

Intermediate Level Introduction to Computing at CARC

1 hour version with SimCov

Matthew Fricke

Version 0.1



Goals

- 1) SLURM scheduler literacy
- 2) Example SimCov
- We wont cover file transfer, storage systems, module system, conda, PBS. (These are all covered in depth in the video tutorials)

Logging into Hopper



First login to the Linux **workstation** in front of you. Your CARC username is on the sign in sheet.

If you have logged in before use your **existing password**

Otherwise, your initial password is **Welcome2CARC**

This is an “important step” so don’t let me move on until you have logged in

Logging into Hopper



```
ssh vanilla@hopper.alliance.unm.edu
```

Should prompt you for a password...

Don't let me move on until you are able to login.

Replace vanilla with your name (unless your last name is Ice)

Logging into Hopper

Welcome to Hopper

Be sure to review the "Acceptable Use" guidelines posted on the CARC website.

For assistance using this system email help@carc.unm.edu.

Tutorial videos can be accessed through the CARC website: Go to <http://carc.unm.edu>, select the "New Users" menu and then click "Introduction to Computing at CARC".

Warning: By default home directories are world readable. Use the chmod command to restrict access.

Don't forget to acknowledge CARC in publications, dissertations, theses and presentations that use CARC computational resources:

"We would like to thank the UNM Center for Advanced Research Computing, supported in part by the National Science Foundation, for providing the research computing resources used in this work."

Please send citations to publications@carc.unm.edu.

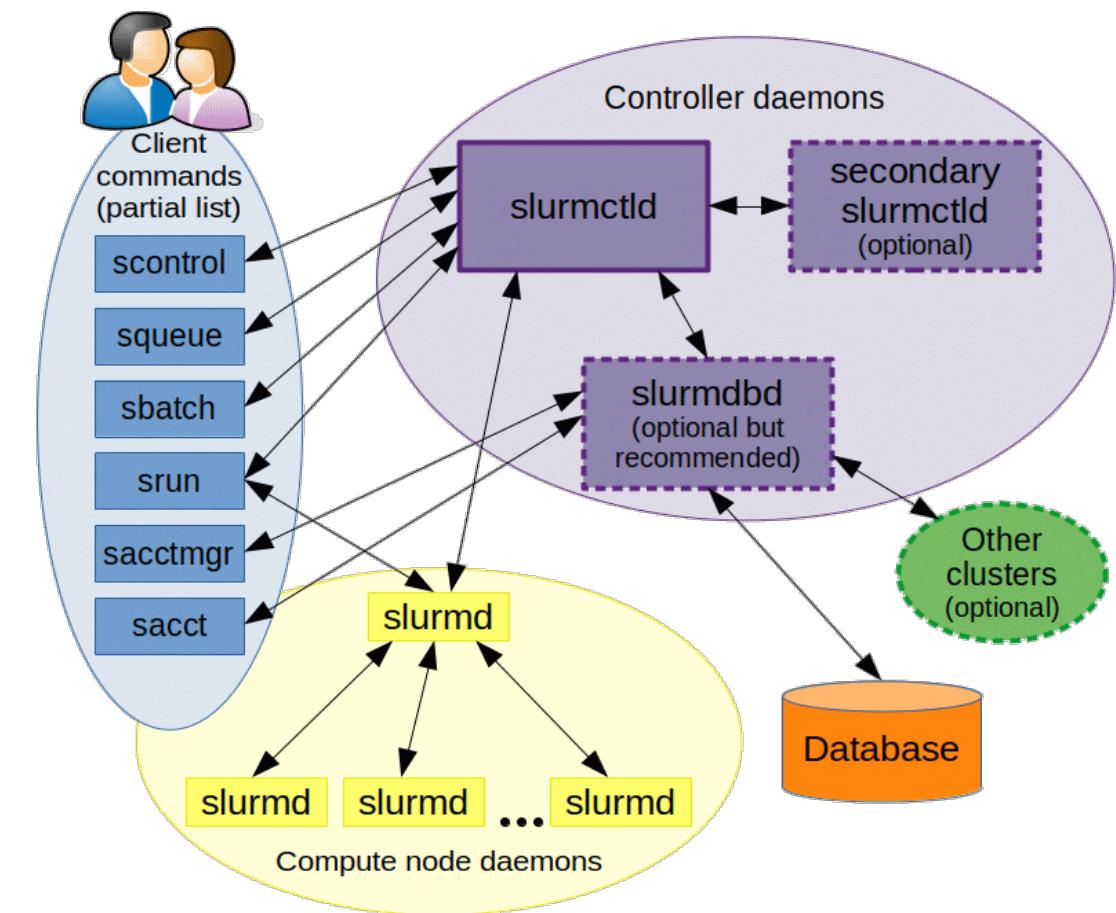
There are three types of slurm partitions on Hopper:

- 1) General - this partition is accessible by all CARC users.
- 2) Condo - preemptable scavenger queue available to all condo users. Your job must use checkpointing to use this queue or you will lose any work you have done if it is preempted by the partition's owner.
- 3) Named partitions - these partitions are available to condo users working under the grant/lab/center that purchased the associated hardware.

Type "qgrok" to get the status of the partitions.

Last login: Wed Jul 27 17:46:13 2022 from 129.24.246.68
mfricke@hopper:~ \$

Simple Linux Utility for Resource Management



ENJOY

Slurm

SODA

IT'S HIGHLY ADDICTIVE!

VOTED #1

SOFT DRINK OF THE 31ST CENTURY!



SELENE MCKENRE



```
[vanilla@hopper ~]$ qgrok
```

queues	free	busy	offline	jobs	nodes	CPUs	GPUs	CPUs/node	GPUs/node	Memory/node	time_limit	CPU_limit
general	4	6	0	4	10	320	0	32	0	93G	2-00:00:00	64
debug	2	0	0	0	2	64	0	32	0	93G	4:00:00	8
condo	22	25	4	7	51	1632	28	32	2	93G-1.5T	2-00:00:00	192
bugs	2	0	0	0	2	64	0	32	0	93G	7-00:00:00	

```
vanilla@hopper:~ $ qgrok
```

queues	free	busy	offline	jobs	nodes	CPUs	GPUs
general	1	9	0	7	10	320	0
debug	2	0	0	0	2	64	0
condo	18	19	1	7	38	1216	8
bugs	0	2	0	0	2	64	0
pcnc	1	1					
pathogen	1	0					
tc	5	5					
gold	2	0					
fishgen	0	1					
neuro-hsc	8	6					
cup-ecs	0	2					
tid	0	1					
biocomp	0	1					
chakra	1	0					
pna	0	0					
totals:	19	28					

Open partitions for use by
everyone with a CARC account.

Purchased by the Office for the
Vice President for Research.

vanilla@hopper:~ \$ qgrok

queues	free	busy	offline	jobs	nodes	CPUs	GPUs
general	1	9	0	7	10	320	0
debug	2	0	0	0	2	64	0
condo	18	19	1	7	38	1216	8
bugs	0	2	0	0	2	64	0
pcnc	1	1	0	0	2	64	0
pathogen	1	0	0	0	1	32	0
tc	5	5	0	3	10	320	0
gold	2	0	0	0	2	64	0
fishgen	0	1	0	0	1	32	0
neuro-hsc	8	6	0	0	14	448	0
cup-ecs	0	2	0	2	2	64	4
tid	0	1	0	0	1	32	2
biocomp	0	1	0	0	1	32	1
chakra	1	0	0	0	1	32	1
pna	0	0	1	0	1	32	0
totals:	19	28	1	14	48	1536	8

```
vanilla@hopper:~ $ qgrok
```

queues	free	busy	offline	jobs	nodes	CPUs	GPUs
general	1	9	0	7	10	320	0
debug	2	0	0	0	2	64	0
condo	18	19	1	7	38	1216	8
bugs	0	2	0	0	2	64	0
pcnc	1	1					
pathogen	1	0					
tc	5	5					
gold	2	0					
fishgen	0	1					
neuro-hsc	8	6					
cup-ecs	0	2					
tid	0	1					
biocomp	0	1					
chakra	1	0					
pna	0	0					
totals:	19	28					

Private partitions

- Reserved for use by the purchaser.
- Request access by emailing support@carc.unm.edu and CC the partition owner.

```
mfricke@hopper:~ $ qgrok
```

queues	free	busy	offline	jobs	nodes	CPUs	GPUs
general	1	9	0	7	10	320	0
debug	2	0	0	0	2	64	0
condo	18	19	1	7	38	1216	8
bugs	0	2	0	0	2	64	0
pcnc	1	1					
pathogen	1	0					
tc	5	5					
gold	2	0					
fishgen	0	1					
neuro-hsc	8	6					
cup-ecs	0	2					
tid	0	1					
biocomp	0	1					
chakra	1	0					
pna	0	0					
totals:	19	28					

Condo “scavenger” partition

- Allows you to use compute nodes purchased by another group that are currently idle.
- May be interrupted at any time if the owners start to use it.

```
[vanilla@hopper ~]$ quotas
```

```
Home Directory (/users/vanilla):
```

```
quota: Cannot resolve mountpoint path /root/.spack: Permission denied
```

```
Disk quotas for user vanilla (uid 659):
```

Filesystem	space	quota	limit	grace	files	quota	limit	grace
chama:/home/homes	1527M	100G	200G		14913	4295m	4295m	

```
Centerwide user scratch (/carc/scratch/users/mfricke)
```

```
Quota information for storage pool Default (ID: 1):
```

user/group	size	chunk	files	
name	used	hard	used	hard
mfricke	592.71 GiB	1024.00 GiB	32784	unlimited

```
Centerwide scratch quota for project mfricke2016174 (/carc/scratch/projects/mfricke2016174)
```

```
Quota information for storage pool Default (ID: 1):
```

user/group	size	chunk	files	
name	used	hard	used	hard
mfricke2016174	190.97 GiB	1024.00 GiB	23704	unlimited

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug        up    4:00:00      2  idle  hopper[011-012]
```

sinfo reports information about
partitions

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug        up    4:00:00      2  idle  hopper[011-012]
```

The debug queues are intended
for testing your programs.

And for interactive jobs.

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug        up    4:00:00      2  idle hopper[011-012]
```



Name

```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST  
debug        up    4:00:00      2  idle  hopper[011-012]
```



You can run a “job” for up to 4 hrs.

```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST  
debug        up    4:00:00      2  idle  hopper[011-012]
```



There are two nodes in this partition.

```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST  
debug        up    4:00:00      2  idle  hopper[011-012]
```



The names of the nodes in the partition

```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST  
debug        up    4:00:00      2  idle  hopper[011-012]
```



The names of the nodes in the partition

```
[vanilla@hopper ~]$ sinfo --partition general
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
general*      up   2-00:00:00      9  alloc  hopper[001-009]
general*      up   2-00:00:00      1  idle   hopper010
```



Hopper001 through 009 are running jobs.

Hopper010 is waiting to be used.

```
[vanilla@hopper ~] $ hostname  
hopper  
[vanilla@hopper ~] $
```



Running on the Head Node.
The head node's name is “hopper”.

```
[vanilla@hopper ~]$ hostname  
hopper
```

```
[vanilla@hopper ~]$ man hostname
```

```
[vanilla@hopper ~]$ hostname  
hopper  
[vanilla@hopper ~]$ man hostname  
(‘q’ to quit)
```

```
[vanilla@hopper ~]$ man man  
(‘q’ to quit)
```

```
[vanilla@hopper ~]$ man sinfo
```

sinfo(1)
Commands

Slurm
sinfo(1)

NAME

sinfo - View information about Slurm nodes and partitions.

SYNOPSIS

`sinfo [OPTIONS...]`

DESCRIPTION

sinfo is used to view partition and node information for a system running Slurm

OPTIONS

`-a, --all`

Display information about all partitions. This causes information to be displayed about partitions that are configured as hidden and partitions that are unavailable to the user's group.

```
[vanilla@hopper ~]$ sinfo --all
```

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
general*	up	2-00:00:00	9	alloc	hopper[001-009]
general*	up	2-00:00:00	1	idle	hopper010
debug	up	4:00:00	2	idle	hopper[011-012]
condo	up	2-00:00:00	1	down*	hopper045
condo	up	2-00:00:00	3	mix	hopper[018-020]
condo	up	2-00:00:00	16	alloc	hopper[013-015,028-036,049-052]
condo	up	2-00:00:00	18	idle	hopper[016-017,021-027,037-044,053]
bugs	up	7-00:00:00	2	alloc	hopper[013-014]
pcnc	up	7-00:00:00	1	alloc	hopper015
pcnc	up	7-00:00:00	1	idle	hopper016
pathogen	up	7-00:00:00	1	idle	hopper017
tc	up	7-00:00:00	3	mix	hopper[018-020]
tc	up	7-00:00:00	2	alloc	hopper[029-030]
tc	up	7-00:00:00	5	idle	hopper[021-025]
gold	up	7-00:00:00	2	idle	hopper[026-027]
fishgen	up	7-00:00:00	1	alloc	hopper028
neuro-hsc	up	7-00:00:00	6	alloc	hopper[031-036]
neuro-hsc	up	7-00:00:00	8	idle	hopper[037-044]
cup-ecs	up	7-00:00:00	2	alloc	hopper[049-050]
tid	up	7-00:00:00	1	alloc	hopper051
biocomp	up	7-00:00:00	1	alloc	hopper052
chakra	up	7-00:00:00	1	idle	hopper053
pna	up	7-00:00:00	1	down*	hopper045

```
[vanilla@hopper ~]$ srun --partition debug hostname
```



Tell slurm to run a program
on a compute node...

```
[vanilla@hopper ~]$ srun --partition debug hostname
```



Run the program on a
compute node in the
debug partition.

```
[vanilla@hopper ~]$ srun --partition debug hostname
```



The program
to run.

```
[vanilla@hopper ~]$ srun --partition debug hostname
srun: Account not specified in script or
~/.default_slurm_account, using latest project
You have not been allocated GPUs. To request GPUs,
use the -G option in your submission script.
hopper011
```

```
[vanilla@hopper ~] $ squeue
```

```
[vanilla@hopper ~]$ squeue
```

JOBID	PARTITION	NAME	USER	ST	TIM
4314	general	PRE	erowland	PD	0:00
4315	general	PRE	erowland	PD	0:00
4317	general	PRE	erowland	PD	0:00
4318	general	PRE	erowland	PD	0:00

**PD means programs
that are waiting their
turn.**

Shows you what the slurm scheduler is doing right now.

Here we can see that user 'erowland' has a lot of programs waiting to run.

```
[vanilla@hopper ~]$ squeue
```

The reason these jobs are not running is that ‘erowland’ is already using the maximum number of CPUs they are allowed.

```
[vanilla@hopper ~]$ squeue -t R --all
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
4405	condo	2ndMA	mfricke	R	1-07:48:30	6	hopper[031-036]
5208	condo	NN	kgu	R	5:48:49	1	hopper015
5210	condo	NN	kgu	R	6:30:13	1	hopper014
5209	condo	NN	kgu	R	6:31:13	1	hopper013
5206	condo	NN	kgu	R	6:32:13	1	hopper051
5207	condo	NN	kgu	R	6:32:13	1	hopper052
5205	condo	NN	kgu	R	6:32:43	1	hopper028
4595	cup-ecs	golConfi	aalasand	R	2-06:51:59	1	hopper050
4594	cup-ecs	golConfi	aalasand	R	2-06:52:03	1	hopper049
5120	general	jupyterh	jacobm	R	11:45:47	1	hopper007
4313	general	PRE	erowland	R	1:17:29	2	hopper[003-004]
5111	general	1stMA	mfricke	R	11:15:28	2	hopper[005-006]
5025	general	c2n	jxzuo	R	1:50	1	hopper001
5024	general	c2n	jxzuo	R	31:28	1	hopper002
5203	general	NN	kgu	R	6:37:50	1	hopper009
5201	general	NN	kgu	R	6:38:14	1	hopper008
4390	tc	UCsTpCyd	lepluart	R	2-15:18:18	3	hopper[018-020]
5198	tc	NN	kgu	R	6:40:19	1	hopper030
5196	tc	NN	kgu	R	6:40:31	1	hopper029

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 hostname
srun: Account not specified in script or
~/.default_slurm_account, using latest project
You have not been allocated GPUs. To request GPUs, use the -G
option in your submission script.
hopper011
hopper011
```

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 hostname
srun: Account not specified in script or
~/.default_slurm_account, using latest project
You have not been allocated GPUs. To request GPUs, use the -G
option in your submission script.
hopper011
hopper011
```

You ran two **copies** of your program.

ntasks is the number of copies to run.

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 8 hostname
srun: Account not specified in script or
~/.default_slurm_account, using latest project
hopper011
hopper011
hopper011
hopper011
You have not been allocated GPUs. To request GPUs, use the -G
option in your submission script.
hopper011
hopper011
hopper011
hopper011
hopper011
```

You ran eight **copies** of your program.

ntasks is the number of copies to run.

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 8 hostname
srun: Account not specified in script or
~/.default_slurm_account, using latest project
hopper011
hopper011
hopper011
hopper011
You have not been allocated GPUs. To request GPUs, use the -G
option in your submission script.
hopper011
hopper011
hopper011
hopper011
hopper011
```

By default, each task (copy of your program) is allowed to use one CPU.

Many programs are able to use more than one CPU at a time.

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest project  
You have not been allocated GPUs. To request GPUs, use the -G option in your submission  
script.  
hopper011  
hopper011
```

Here we are telling SLURM to run 2 copies of our program and let each copy of our program use 2 CPUs.

```
[vanilla@hopper ~]$ srun --partition debug --nodes 2 -ntasks-per-node 4 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using  
latest project  
hopper012  
You have not been allocated GPUs. To request GPUs, use the -G option in  
your submission script.  
hopper012  
hopper011  
hopper011  
hopper012  
hopper012  
hopper011  
hopper011
```

Here we are telling SLURM to run 4 copies of our program on 2 different compute nodes.

This is useful when our programs need a bigger share of the compute node.

```
[vanilla@hopper ~]$ srun --partition debug --nodes 2  
--ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or  
~/.default_slurm_account, using latest project  
hopper011  
You have not been allocated GPUs. To request GPUs, use  
the -G option in your submission script.  
hopper011  
hopper012  
hopper012
```

And we can combine all three.

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2  
hostname
```

```
srun: Account not specified in script or  
~/.default_slurm
```

```
hopper012
```

```
hopper012
```

```
You have not be
```

```
the -G option i
```

```
hopper011
```

```
Hopper011
```

**And we can specify how much
memory we want.**

**--mem 4G means give me 4
gigabytes of memory per node.**

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2  
hostname  
srun: Account not specified in script or  
~/.default_slurm  
hopper012  
hopper012  
You have not been  
the -G option in  
hopper011  
Hopper011
```

Why does all this matter?

The purpose of SLURM is to provide you the hardware your programs need.

So you have to understand what those requirements are really well.

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or  
~/.default_slurm_account using latest project  
hopper012  
hopper012  
You have not been assigned to a project.  
the -G option is required.  
hopper011  
Hopper011
```

- 1) Can my program use multiple CPUs?**
- 2) How much memory does my program need?**
- 3) Can my program use multiple compute nodes (MPI*, GNU Parallel*)?**
- 4) Can my program use GPUs?**

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or  
~/.default_slurm_account using latest project  
hopper012  
hopper012  
You have not been  
the -G option in  
hopper011  
Hopper011
```

This command is getting pretty long.

We can use **shell scripts** to automate
all this in **batch mode**.

Interactive vs Batch Mode

Interactive Mode

- Everything so far has been interactive. You request hardware, run your program, and get the output on your screen right away.

Batch Mode

- Most programs at an HPC center are run in “batch” mode.
- Batch mode means we write a shell script that the SLURM scheduler runs for us. The script requests hardware just like we did with salloc and then runs the commands in the script.
- Whatever would have been written to the screen is saved to a file instead.

```
[vanilla@hopper ~]$ git clone https://lobogit.unm.edu/CARC/workshops.git  
Cloning into 'workshops'...  
remote: Enumerating objects: 132, done.  
remote: Counting objects: 100% (75/75), done.  
remote: Compressing objects: 100% (43/43), done.  
remote: Total 132 (delta 33), reused 74 (delta 32), pack-reused 57  
Receiving objects: 100% (132/132), 57.58 KiB | 3.60 MiB/s, done.  
Resolving deltas: 100% (51/51), done.
```

Rather than make you write shell scripts lets just download some we wrote for this workshop...

```
[vanilla@hopper ~]$ tree workshops
```

```
workshops/
├── intro_workshop
│   ├── code
│   │   ├── calcPiMPI.py
│   │   ├── calcPiSerial.py
│   │   └── vecadd
│   │       ├── Makefile
│   │       ├── vecadd_gpu.cu
│   │       ├── vecadd_mpi_cpu
│   │       ├── vecadd_mpi_cpu.c
│   │       ├── vecaddmpi_cpu.sh
│   │       └── vecadd_mpi_gpu.c
│
└── data
    ├── H2O.gjf
    └── step_sizes.txt
└── slurm
    ├── calc_pi_array.sh
    ├── calc_pi_mpi.sh
    ├── calc_pi_parallel.sh
    ├── calc_pi_serial.sh
    ├── gaussian.sh
    ├── hostname_mpi.sh
    ├── vecadd_hopper.sh
    ├── vecadd_xena.sh
    ├── workshop_example2.sh
    ├── workshop_example3.sh
    └── workshop_example.sh

```

Run tree to see how the workshops directories are organized...

```
[vanilla@hopper ~]$ tree workshops
```

```
workshops/
├── intro_workshop
│   ├── code
│   │   ├── calcPiMPI.py
│   │   ├── calcPiSerial.py
│   │   └── vecadd
│   │       ├── Makefile
│   │       ├── vecadd_gpu.cu
│   │       ├── vecadd_mpi_cpu
│   │       ├── vecadd_mpi_cpu.c
│   │       ├── vecaddmpi_cpu.sh
│   │       └── vecadd_mpi_gpu.c
│
├── data
│   ├── H2O.gjf
│   └── step_sizes.txt
└── slurm
    ├── calc_pi_array.sh
    ├── calc_pi_mpi.sh
    ├── calc_pi_parallel.sh
    ├── calc_pi_serial.sh
    ├── gaussian.sh
    ├── hostname_mpi.sh
    ├── vecadd_hopper.sh
    ├── vecadd_xena.sh
    ├── workshop_example2.sh
    ├── workshop_example3.sh
    └── workshop_example.sh

```

```
README.md
```

Run **tree** to see how the workshops directories are organized...

The workshop files are divided into “code”, “slurm”, and “data” directories.

```
[vanilla@hopper intro_workshop]$ pwd  
/users/vanilla/workshops/intro_workshop  
[vanilla@hopper intro_workshop]$ cat slurm/workshop_example.sh  
#!/bin/bash  
#SBATCH --partition debug  
#SBATCH --ntasks 4  
#SBATCH --time 00:05:00  
#SBATCH --job-name ws_example  
#SBATCH --mail-user your_username@unm.edu  
#SBATCH --mail-type ALL
```

hostname

Let's take a look at the **workshop_example.sh** script in the slurm directory...

```
[vanilla@hopper intro_workshop]$ sbatch slurm/workshop_example.sh  
sbatch: Account not specified in script or ~/.default_slurm_account,  
using latest project  
Submitted batch job 5252  
[vanilla@hopper intro_workshop]$
```

We **submit** our slurm
shell script with the
sbatch command.

```
[vanilla@hopper intro_workshop]$ sbatch slurm/workshop_example.sh  
sbatch: Account not specified in script or ~/.default_slurm_account,  
using latest project  
Submitted batch job 5252  
[vanilla@hopper intro_workshop]$
```

Notice that the only output we get is a job id.

This indicates that the script was successfully sent to the scheduler.

The commands in the script will run as soon as the hardware requested is available.

We submit our slurm shell script with the sbatch command.

Workflow

Head Node

User 1

Program A

Script A

User 2

Program B

Script B

Compute Node 01

Compute Node 02

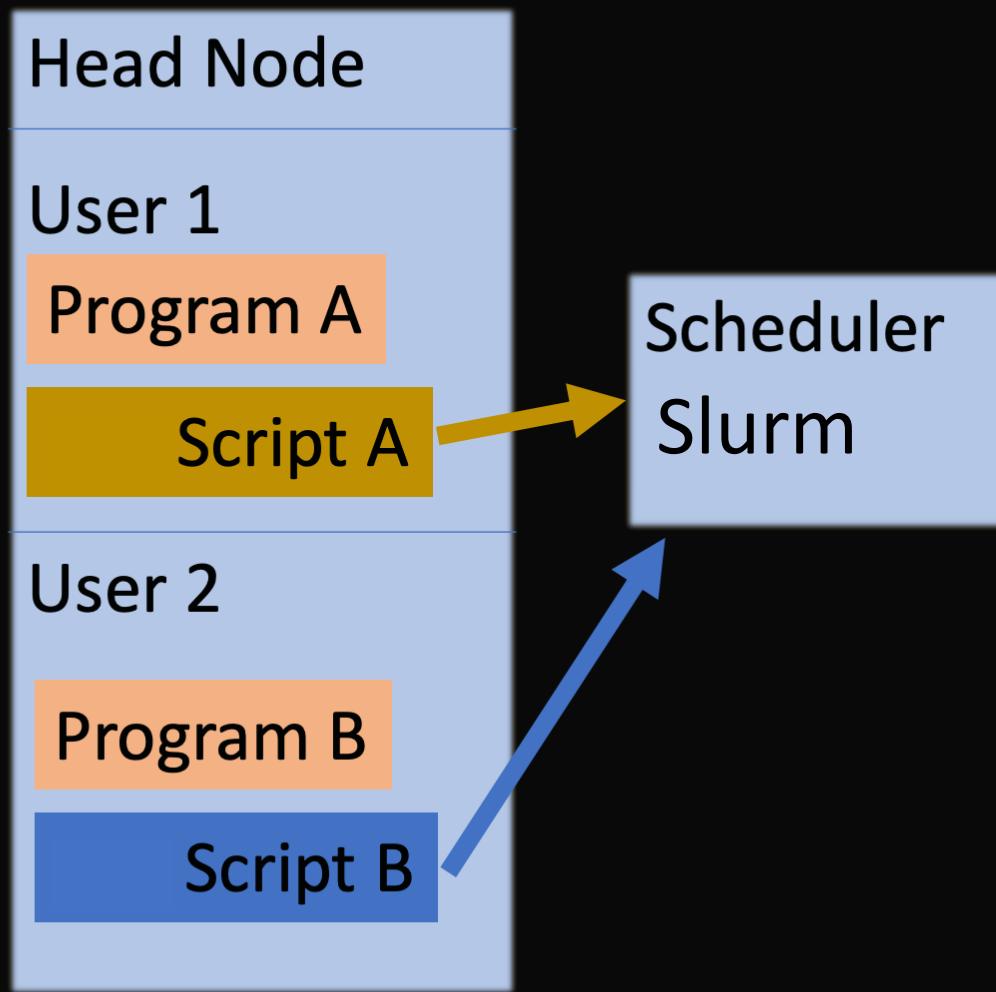
Compute Node 03

Compute Node 04

Compute Node 05

Shared filesystems – All nodes can access the same programs and write output

Workflow



Compute Node 01

Compute Node 02

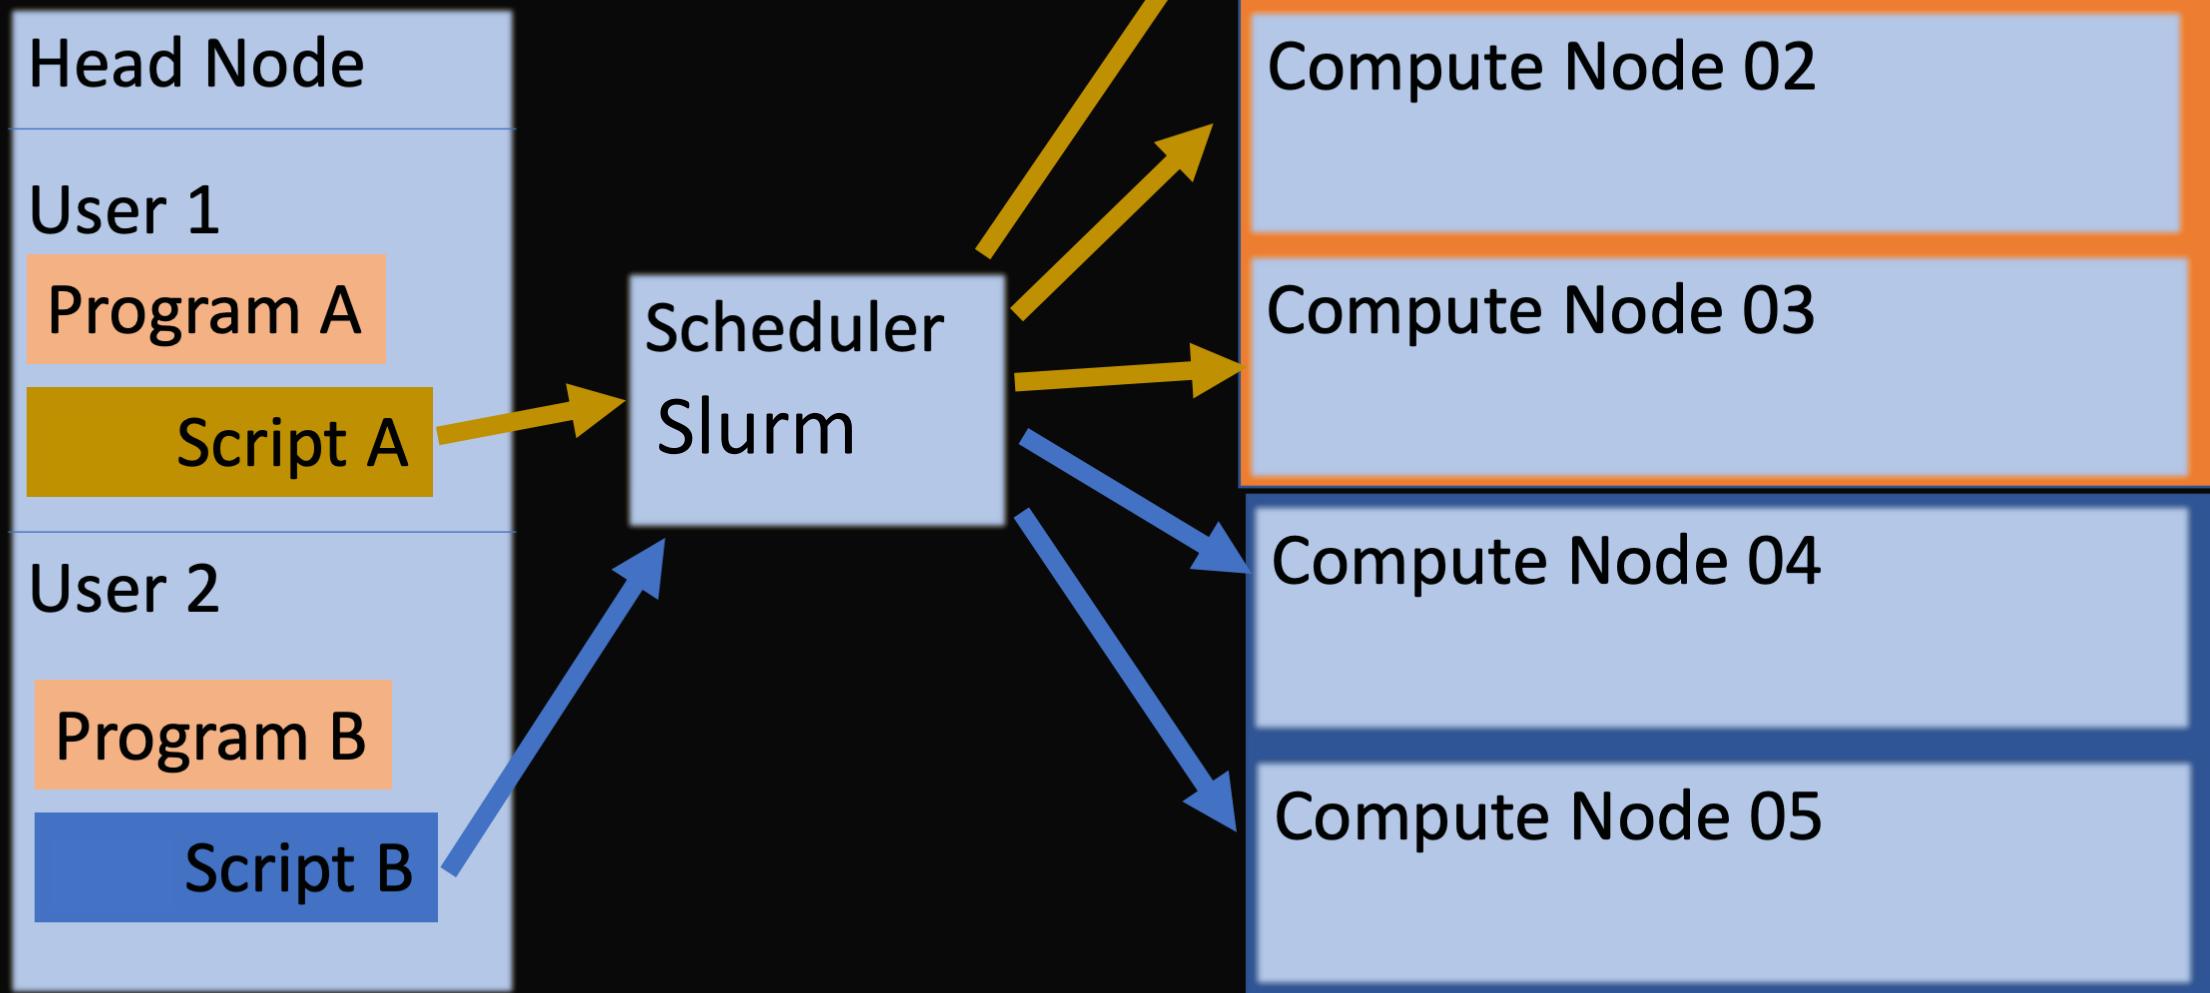
Compute Node 03

Compute Node 04

Compute Node 05

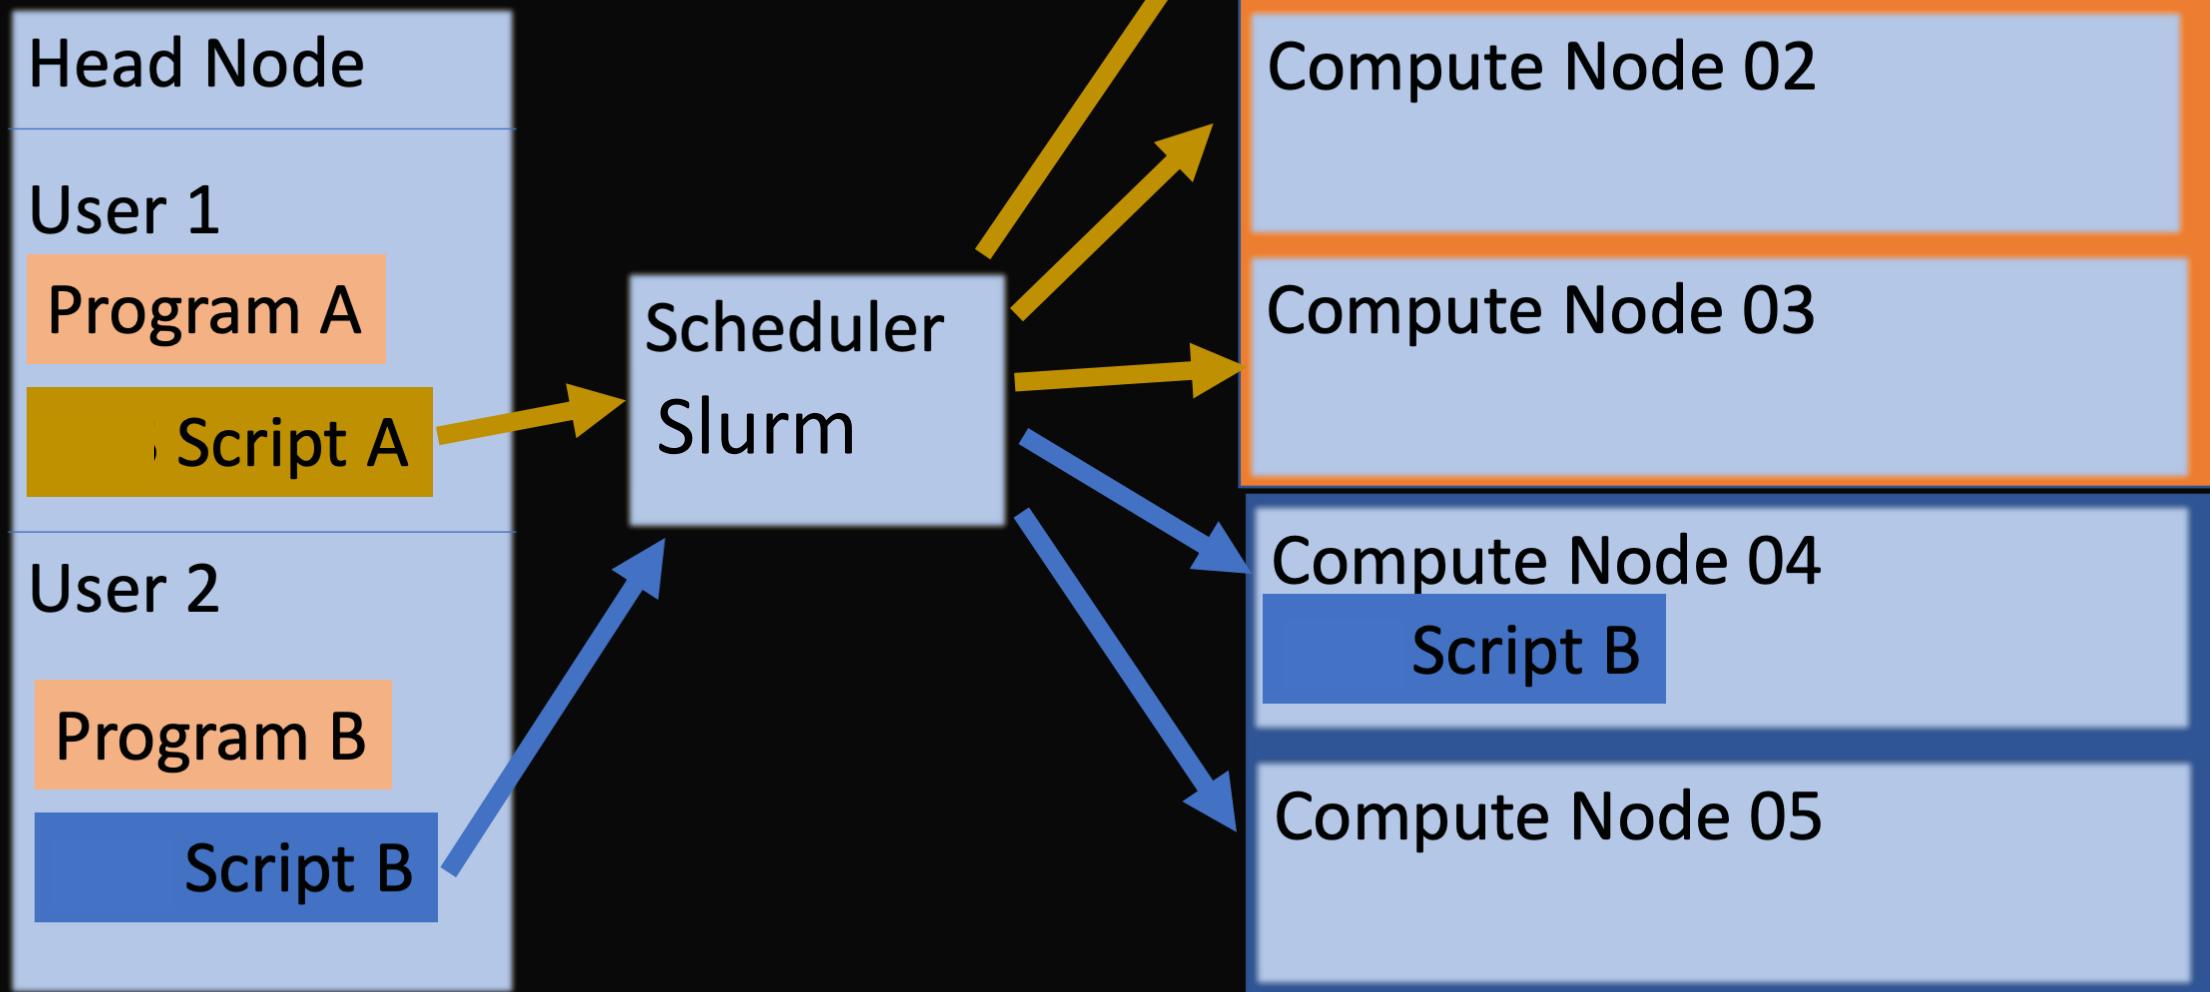
Shared filesystems – All nodes can access the same programs and write output

Workflow



Shared filesystems – All nodes can access the same programs and write output

Workflow



Shared filesystems – All nodes can access the same programs and write output

Workflow

We need something in the script to run the program on all the nodes. E.g. srun.

Head Node

User 1

Program A

Script A

User 2

Program B

Script B

Scheduler
Slurm

Compute Node 01

Script A

Program A

Compute Node 02

Program A

Program A

Compute Node 03

Program A

Program A

Compute Node 04

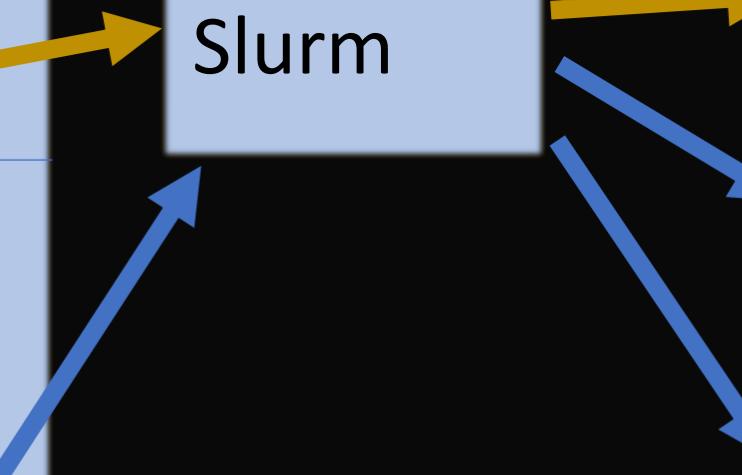
Script B

Program B

Compute Node 05

Program B

Program B



Shared filesystems – All nodes can access the same programs and write output

```
[vanilla@hopper intro_workshop]$ ls  
code  data  pbs  slurm  slurm-5252.out
```

The **hostname** command is very fast so everyone's job should finish in a few seconds.

When it is finished you will have a new file named slurm-{your job id}.out.



```
[vanilla@hopper intro_workshop]$ ls  
code  data  pbs  slurm  slurm-5252.out
```



When it is finished you will have a new file named slurm-{your job id}.out.

```
[vanilla@hopper intro_workshop]$ cat slurm-5252.out  
hopper011
```

```
[vanilla@hopper intro_workshop]$ ls  
code  data  pbs  slurm  slurm-5252.out
```



When it is finished you will have a new file named slurm-{your job id}.out.

```
[vanilla@hopper intro_workshop]$ cat slurm-5252.out  
hopper011
```

Why did it only run the program once instead of 4 times?

```
[vanilla@hopper intro_workshop]$ sbatch slurm/workshop_example1.sh
sbatch: Account not specified in script or ~/.default_slurm_account,
using latest project
Submitted batch job 5252
[vanilla@hopper intro_workshop]$
```

Take a look at the
output file.

```
[vanilla@hopper intro_workshop]$ module load miniconda3
```

```
[vanilla@hopper intro_workshop]$ conda create -n numpy numpy
```

Wait a while – introduce yourselves to your neighbor...

Conda allows you to install software into your home directory. In this case we need the numerical python libraries for calcPiSerial.py

**Let's experiment with a
program that does
slightly more than
print the hostname.**

```
[vanilla@hopper intro_workshop]$source activate numpy  
[vanilla@hopper intro_workshop]$srun --partition debug python  
code/calcPiSerial.py 10  
srun: Using account 2016199 from ~/.default_slurm_account  
You have not been allocated GPUs. To request GPUs, use the -G  
option in your submission script.  
Pi = 3.14242598500109870940, (Diff=0.00083333141130559341)  
(calculated in 0.000005 secs with 10 steps)
```

**Activate the numpy environment and
Run calcPiSerial.py on a compute node.**

**For our example program the more steps it takes the
more accurate it is, but the longer it takes.**

```
[vanilla@hopper intro_workshop]$ sbatch slurm/calc_pi_serial.sh
sbatch: Using account 2016199 from ~/.default_slurm_account
Submitted batch job 5263
vanilla@hopper:~/workshops/intro_workshop$ squeue -me
JOBID PARTITION      NAME      USER ST          TIME   NODES NODELIST(REASON)
5263    debug calc_pi_  vanilla  R          0:44       1 hopper011
```

Edit `slurm/calc_pi_serial.sh`.

Change the email address to your address and submit the script.

Then enter `squeue --me` to see the job status.

Take a look at the job output.

Moses Biological Computation Lab

[Home](#) [Research](#) ▾ [Publications](#) [Views on COVID-19](#) [Lab Page](#) ▾ [Courses](#) ▾ [Swarms](#) [News](#) ▾ [Website Archives](#)

UNM / Home

Home

Research ▶

Publications

Views on COVID-19

Lab Page ▶

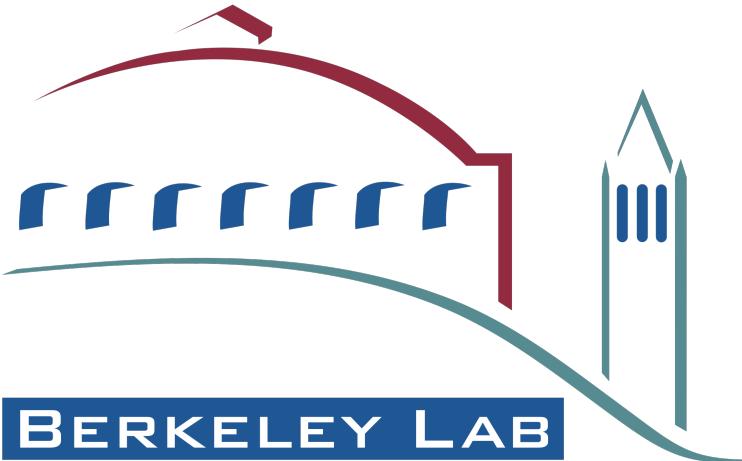
Courses ▶

Swarms



We live in a complex world. Many of the greatest challenges in science and engineering require understanding how complex systems evolve, what makes them efficient, scalable and robust, and why they fail. Our **highly interdisciplinary research** in the Moses Biological Computation lab seeks to understand complexity in natural and engineered systems: we develop computational and mathematical models of biological systems (particularly ant colonies and immune systems) and we apply what we learn in biologically-inspired computation (particularly swarm robotics).

I am the PI of two programs that seek to diversify computer science, **NM CSforAll** and the **NASA Swarmathon**. Combined, these two programs have trained over 3000 students, 75% of them women and/or students of color. These students have been historically underrepresented in STEM, particularly in Computer Science. Our team is changing that NOW!



SimCov Example

SimCov is a C++ program for modelling CoVID virus dynamics in the lungs.
It uses a library called UPCXX to spread the computation over multiple nodes.

RESEARCH ARTICLE

Spatially distributed infection increases viral load in a computational model of SARS-CoV-2 lung infection

Melanie E. Moses^{1,2*}, Steven Hofmeyr³, Judy L. Cannon⁴, Akil Andrews¹, Rebekah Gridley⁴, Monica Hinga¹, Kurtis Leyba⁵, Abigail Pribisova¹, Vanessa Suradjidjaja¹, Humayra Tasnim¹, Stephanie Forrest^{2,5}

1 Department of Computer Science, University of New Mexico, Albuquerque, New Mexico, United States of America, **2** Santa Fe Institute, Santa Fe, New Mexico, United States of America, **3** Lawrence Berkeley National Laboratory, Berkeley, California, United States of America, **4** Department of Molecular Genetics and Microbiology, University of New Mexico School of Medicine, Albuquerque, New Mexico, United States of America, **5** Biodesign Institute, Arizona State University, Tempe, Arizona, United States of America

* melaniem@unm.edu



OPEN ACCESS

Citation: Moses ME, Hofmeyr S, Cannon JL, Andrews A, Gridley R, Hinga M, et al. (2021) Spatially distributed infection increases viral load in a computational model of SARS-CoV-2 lung infection. PLoS Comput Biol 17(12): e1009735. <https://doi.org/10.1371/journal.pcbi.1009735>

Editor: Amber M Smith, University of Tennessee Health Science Center College of Medicine Memphis, UNITED STATES

Received: June 24, 2021

Abstract

A key question in SARS-CoV-2 infection is why viral loads and patient outcomes vary dramatically across individuals. Because spatial-temporal dynamics of viral spread and immune response are challenging to study *in vivo*, we developed Spatial Immune Model of Coronavirus (SIMCoV), a scalable computational model that simulates hundreds of millions of lung cells, including respiratory epithelial cells and T cells. SIMCoV replicates viral growth dynamics observed in patients and shows how spatially dispersed infections can lead to increased viral loads. The model also shows how the timing and strength of the T cell response can affect viral persistence, oscillations, and control. By incorporating spatial interactions, SIMCoV provides a parsimonious explanation for the dramatically different viral load trajectories among patients by varying only the number of initial sites of infection and

```
[vanilla@hopper]$ git clone --recurse-submodules  
https://github.com/AdaptiveComputationLab/simcov.git
```

```
Cloning into 'simcov'...  
remote: Enumerating objects: 3263, done.  
remote: Counting objects: 100% (517/517), done.  
remote: Compressing objects: 100% (131/131), done.  
remote: Total 3263 (delta 446), reused 425 (delta 386), pack-reused 2746  
Receiving objects: 100% (3263/3263), 35.13 MiB | 30.36 MiB/s, done.  
Resolving deltas: 100% (2195/2195), done.  
Updating files: 100% (641/641), done.  
Submodule 'upcxx-utils' (https://bitbucket.org/berkeleylab/upcxx-utils.git) registered for path 'upcxx-utils'  
Cloning into '/users/mfricke/simcov/upcxx-utils'...  
remote: Enumerating objects: 5568, done.  
remote: Counting objects: 100% (4932/4932), done.  
remote: Compressing objects: 100% (2358/2358), done.  
remote: Total 5568 (delta 3415), reused 3655 (delta 2550), pack-reused 636  
Receiving objects: 100% (5568/5568), 1.48 MiB | 2.21 MiB/s, done.  
Resolving deltas: 100% (3771/3771), done.  
Submodule path 'upcxx-utils': checked out '5694db37fa4d5a0e100b58968b280aa4a934f7cd'
```

Git clone the SimCov repository

The recursive argument tells git to also download UPCXX utils

```
[vanilla@hopper]$ cd simcov
```

Move into the simcov git repo directory

```
[vanilla@hopper]$ cd simcov
```

Move into the simcov git repo directory

```
[vanilla@hopper ~/simcov]$ chmod +x ./build_hopper.sh
```

Change the permissions on the build script to make it executable

```
vanilla@hopper ~/simcov]$ module load gcc/12.1.0-crtl  
[vanilla@hopper ~/simcov]$ module load cmake/3.11.4-qkyj  
[vanilla@hopper ~/simcov]$ module load openmpi/4.1.3-j6zb  
[vanilla@hopper ~/simcov]$ module load upcxx/2020.10.0-6eh2
```

Load the modules we need to compile simcov

GCC is the GNU C compiler we will use to compile SimCov

CMAKE is the build manager for building the code

OPENMPI allows programs to use multiple compute nodes

UPCXX uses MPI to combine memory on different compute nodes

(This is in the build script but is good to know what's happening)

```
[vanilla@hopper ~/simcov]$ ./build_hopper.sh Release
Installing to /users/vanilla/simcov/install
-- The C compiler identification is GNU 12.1.0
-- The CXX compiler identification is GNU 12.1.0
-- Check for working C compiler: /opt/spack/opt/spack/linux-rocky8-skylake_avx512/gcc-8.5.0/gcc-12.1.0-crtl4tekll36atmxstxm3m346irylu7t/bin/gcc
-- Check for working C compiler: /opt/spack/opt/spack/linux-rocky8-skylake_avx512/gcc-8.5.0/gcc-12.1.0-crtl4tekll36atmxstxm3m346irylu7t/bin/gcc -- works
-- Detecting C compiler ABI info
-- Detecting C compiler ABI info - done
...
[ 96%] Built target test_aggr_store
[100%] Linking CXX executable simcov
[100%] Built target simcov
Install the project...
-- Install configuration: "Release"
...
-- Installing: /users/vanilla/simcov/install/bin/simcov
Build took 25s
```

Run the compilation shell script. This compiles the code, builds the libraries and executables, and installs simcov into your home directory.

```
[vanilla@hopper ~/simcov]$ cat slurm_hopper.sh
```

```
#!/bin/bash
#SBATCH --partition=general
#SBATCH --nodes=2
#SBATCH --ntasks-per-node=32
#SBATCH --time=2-00:00
#SBATCH --job-name=simcov-tutorial
#SBATCH --mail-user=yournetid@unm.edu
#SBATCH --mail-type=ALL

# Setup environment
export GASNET_MAX_SEGSIZE=128MB/P
export UPCXX_SHARED_HEAP_SIZE='128 MB'
export UPCXX_THREADMODE=seq
export UPCXX_CODEMODE=opt
export UPCXX_NETWORK=ibv

# Load modules
module purge
module load gcc/12.1.0-crtl
module load cmake/3.11.4-qkyj
module load openmpi/4.1.3-j6zb
module load upcxx/2020.10.0-6eh2
srun --mpi=pmi2 install/bin/simcov --config covid_default.config
```

```
[vanilla@hopper ~simcov]$ sbatch slurm_hopper.sh  
sbatch: Account not specified in script or ~/.default_slurm_account, using latest  
project  
Submitted batch job 1790591
```

Submit the SLURM shell script

```
[vanilla@hopper ~simcov]$ sbatch slurm_hopper.sh  
sbatch: Account not specified in script or ~/.default_slurm_account, using latest  
project  
Submitted batch job 1790591
```

Submit the SLURM shell script

```
[vanilla@hopper ~simcov]$ squeue --me  
JOBID PARTITION      NAME      USER ST      TIME   NODES NODELIST(REASON)  
1790591  general simcov-t  vanilla R      0:11      2 hopper[008-009]
```

Get the status of the job

```
[vanilla@hopper ~/simcov]$ sbatch slurm_hopper.sh
sbatch: Account not specified in script or ~/.default_slurm_account, using latest
project
Submitted batch job 1790591
```

Submit the SLURM shell script

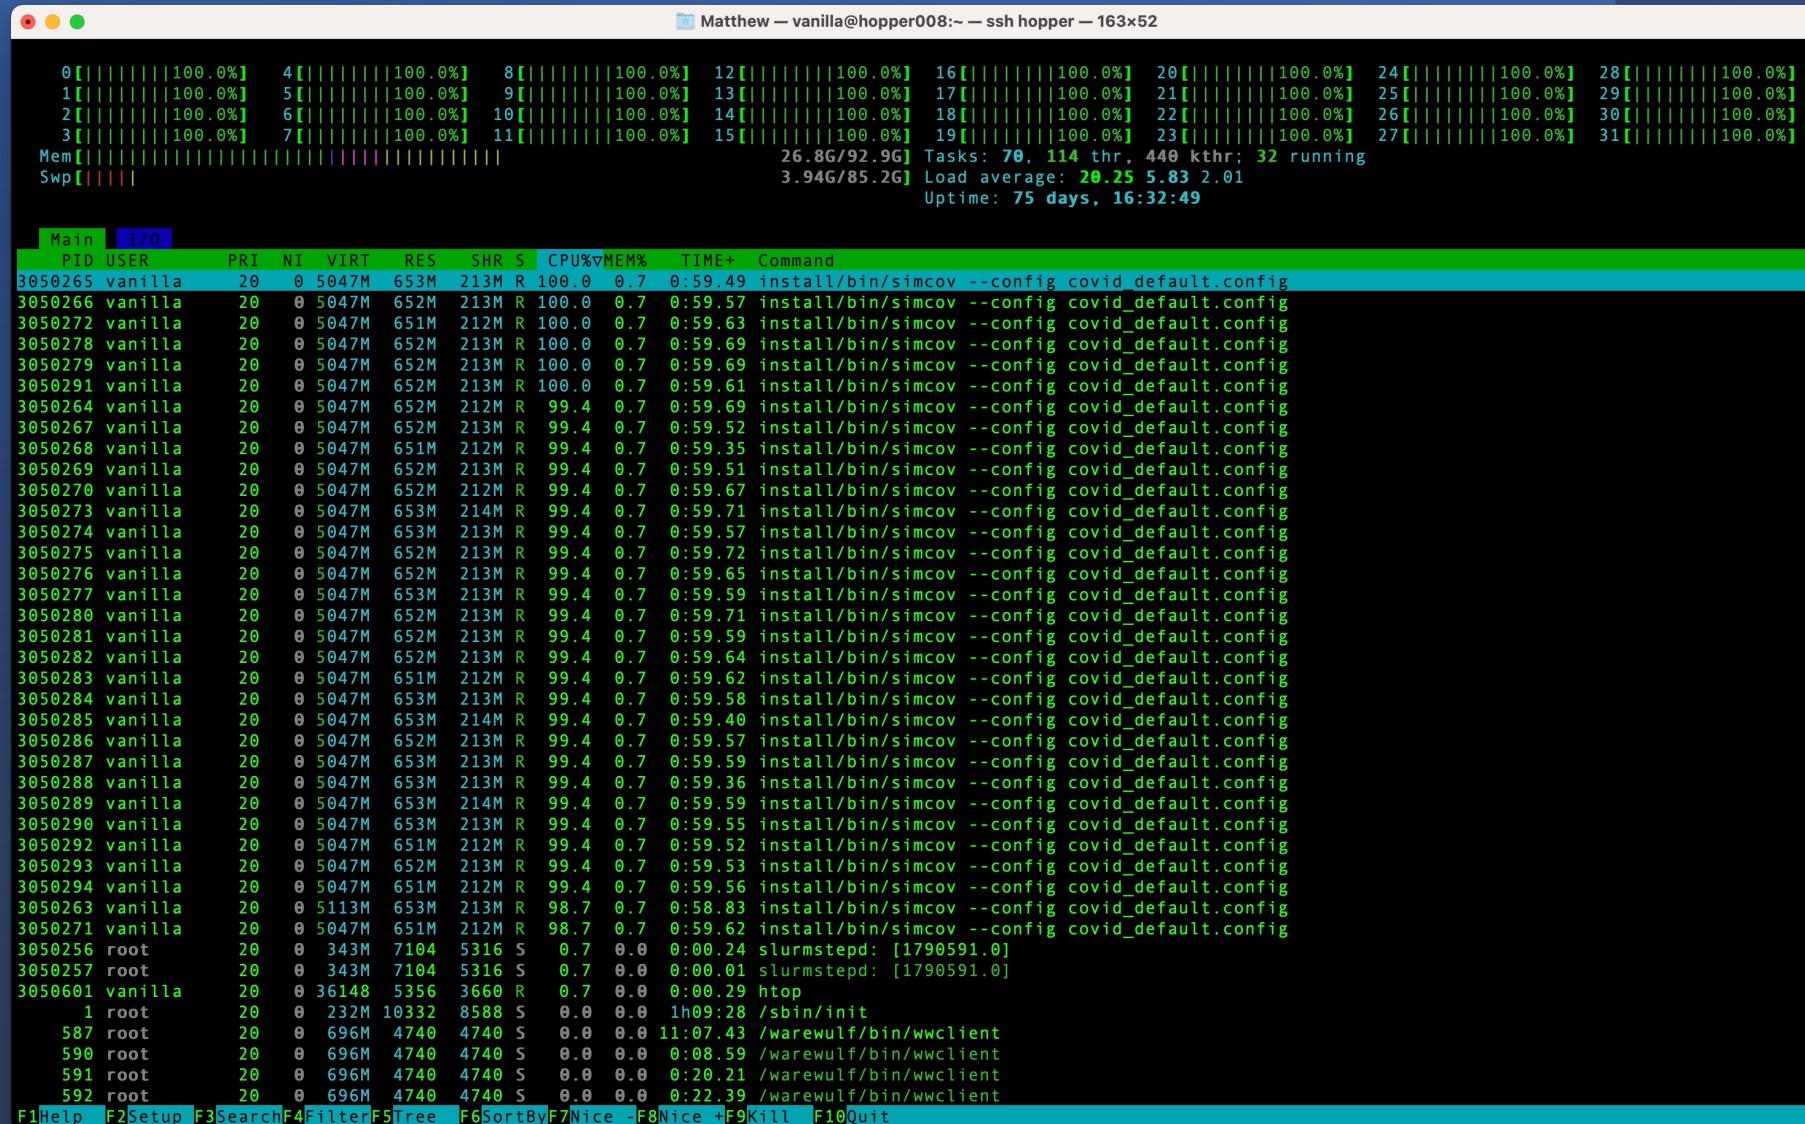
```
[vanilla@hopper ~/simcov]$ squeue --me
      JOBID PARTITION      NAME      USER ST      TIME   NODES NODELIST(REASON)
 1790591    general simcov-t  vanilla  R  0:11      2 hopper[008-009]
```

Get the status of the job

```
[vanilla@hopper ~/simcov]$ ssh hopper008
Last login: Tue Jan 31 07:47:35 2023 from 172.17.0.2
[vanilla@hopper008 ~]$
```

Login to one of the compute nodes we were assigned

[vanilla@hopper008 ~]\$ htop



Check to see that our job is using the node properly

```
[vanilla@hopper008 ~]$ exit  
logout  
Connection to hopper008 closed.  
[vanilla@hopper ~/simcov]$
```

Logout of the compute node

```
[vanilla@hopper008 ~]$ exit  
logout  
Connection to hopper008 closed.  
[vanilla@hopper ~/simcov]$
```

Logout of the compute node

```
[vanilla@hopper ~/simcov]$ ls -lt  
total 148  
-rw-r--r-- 1 vanilla users 33460 Jan 31 08:22 slurm-1790596.out  
drwx----- 3 vanilla users 4096 Jan 31 08:19 simcov-results-n64-N2  
-rw-r--r-- 1 vanilla users 597 Jan 31 08:18 slurm_hopper.sh
```

Inspect the results of your computation

ls –rt sorts the output by modification time. The most recent slurm-{jobid}.out file will be at the top.

```
[vanilla@hopper ~/simcov]$ cat slurm-{jobid}.out
```

After a bunch of warnings about InfiniBand conduits not being configured you should see...

```
Matthew — vanilla@hopper:~/simcov — ssh hopper — 165x55

SimCov version SimCov version 0.1.1.8024051-dirty-master built on 20230131_073221
Options:
  dim = 15000 15000 1
  whole-lung-dim = 48000 40000 20000
  timesteps = 33120
  infection-coords = uniform:1
  initial-infection = 1000
  incubation-period = 480
  apoptosis-period = 180
  expressing-period = 900
  infectivity = 0.001
  infectivity-multiplier = 1.0
  virion-production = 1.1
  virion-production-multiplier = 1.0
  virion-clearance = 0.004
  virion-diffusion = 0.15
  chemokine-production = 1.0
  chemokine-decay = 0.01
  chemokine-diffusion = 1.0
  min-chemokine = 1e-6
  antibody-factor = 1
  antibody-period = 5760
  tcell-generation-rate = 105000
  tcell-initial-delay = 10080
  tcell-vascular-period = 5760
  tcell-tissue-period = 1440
  tcell-binding-period = 10
  max-binding-prob = 1
  tcells-follow-gradient = false
  seed = 29
  sample-period = 0
  sample-resolution = 1
  max-block-dim = 10
  output = simcov-results-n64-N2

Starting run with 64 processes on 2 nodes at 01/31/23 07:47:40
Initial free memory across all nodes: 152.82GB
Starting with 76.54GB free on node 0
Dividing 225000000 grid points into 3515625 squares of size 64 (8^2), with 54932 per process
Total initial memory required per process is at least 281.63MB with each grid point requiring on average 84 bytes
Memory used on node 0 after initialization is 10.91GB
Fraction of circulating T cells extravasating is 5.85938e-06
# datetime step incb expr apop dead tvas ttis chem virs chempts %infct <%active lln>
[0] <main.cpp:141> WARNING: Time step 0: SUCCESSFUL initial infection at (7500, 7500, 0) after 0 tries
[01/31/23 07:47:44 0.00s]: 0 1 0 0 0 0 0 0.00e+00 0.00e+00 1.00e+00 0.00 < 0.000 0.016 >
[01/31/23 07:47:44 0.15s]: 662 154 71 0 0 0 0 1.79e-06 5.99e+03 6.58e+03 0.00 < 0.004 0.649 >
[01/31/23 07:47:44 0.29s]: 1324 745 490 0 2 0 0 5.76e-06 6.86e+04 1.69e+04 0.00 < 0.007 0.811 >
[01/31/23 07:47:44 0.36s]: 1986 1406 1422 0 179 0 0 1.37e-05 2.87e+05 2.23e+04 0.00 < 0.010 0.820 >
[01/31/23 07:47:45 0.45s]: 2648 1922 2450 0 1138 0 0 2.34e-05 5.59e+05 2.89e+04 0.00 < 0.013 0.807 >
[01/31/23 07:47:45 0.54s]: 3310 2529 3441 0 2854 0 0 3.36e-05 8.49e+05 3.66e+04 0.00 < 0.016 0.892 >
[01/31/23 07:47:46 0.65s]: 3972 3062 4591 0 5267 0 0 4.38e-05 1.15e+06 4.42e+04 0.01 < 0.020 0.831 >
[01/31/23 07:47:47 0.75s]: 4634 3652 5640 0 8530 0 0 5.45e-05 1.45e+06 5.34e+04 0.01 < 0.024 0.921 >
[01/31/23 07:47:48 0.89s]: 5296 4331 6789 0 12545 0 0 6.54e-05 1.75e+06 6.37e+04 0.01 < 0.028 0.915 >
[01/31/23 07:47:49 1.00s]: 5958 5011 8051 0 17389 0 0 7.69e-05 2.08e+06 7.46e+04 0.01 < 0.033 0.895 >
[01/31/23 07:47:50 1.25s]: 6620 5588 9252 0 23135 0 0 8.76e-05 2.41e+06 8.63e+04 0.02 < 0.038 0.928 >
```

Useful Slurm Commands

squeue --me --long	shows information about jobs you submitted
squeue --me --start	shows when slurm expects your job to start
scancel jobid	cancels a job
scancel --u \$USER	cancels all your jobs
sacct	shows your job history
seff jobid	shows how efficiently the hardware was used