

Beginner's Introduction to Computing at CARC with JupyterHub

Version 0.1

Goals

- 1) Basic Linux Literacy
- 2) Using Compute Nodes with Slurm
- 3) Jupyterhub

I hope to spend about an hour on each with a 15 minute break.

Outline

- High Performance Computing Overview
- Logging in
- The BASH Shell
- The Slurm Job Scheduler
- Transferring data to and from CARC
- Accessing software and the module system
- Conda environments
- Jupyterhub

High Performance Computing

- What is high performance computing?
 - Really just means something that is a lot more powerful than your desktop or laptop.
 - Hardware:
 - That might mean more and faster processors to do the calculations more quickly (eg 400 CPUs instead of 4)
 - More RAM so you can work on bigger problems (3,000 GB instead of 8)
 - Bigger file systems so you can process larger datasets
 - More and bigger GPUs to accelerate your computations (12 GPUs at a time instead of 1)
 - People:
 - Someone else to manage the systems and keep them running and secure
 - Someone to answer your questions and help with problems

Logging into Wheeler



First login to the Linux **workstation** in front of you.

Use your CARC username and password.

Keven, Tannor, Viacheslav (Slava), and Jose can help you login if you have trouble.

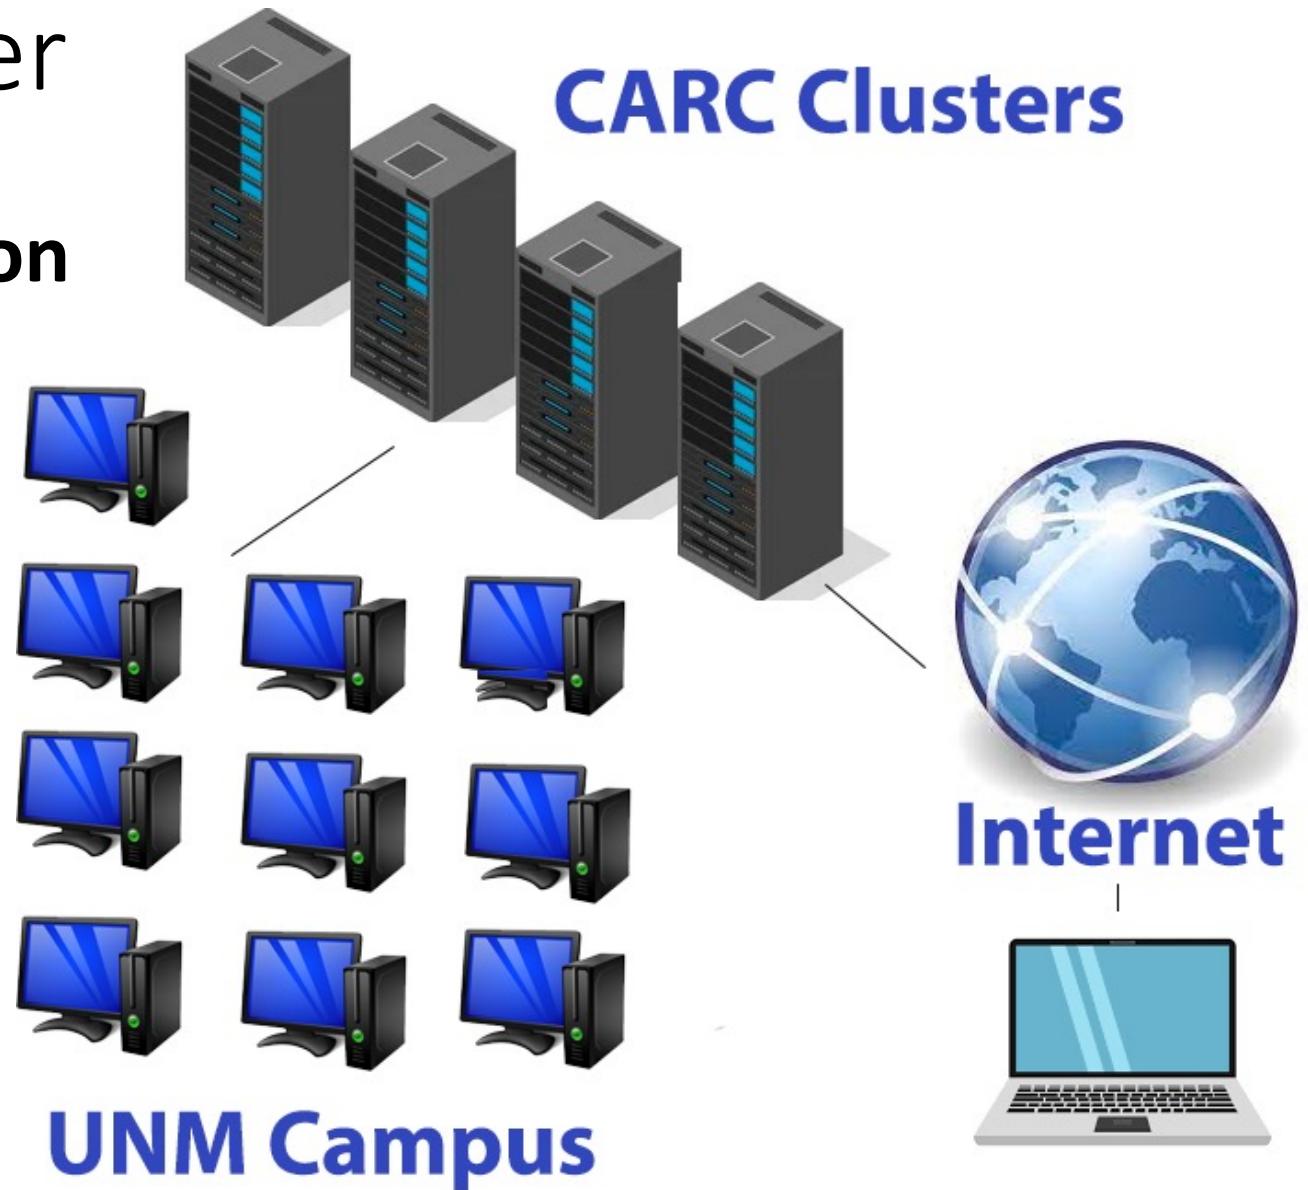
This is an “important step” so don’t let me move on until you have logged in

Logging into Wheeler

First login to the **workstation** in front of you.

You will always login to CARC cluster remotely.

These clusters don't even have monitors.



Logging into Wheeler

We are going to use a program called **secure shell**.

Secure shell (ssh) is now built into every major operating system (Windows, OSX, and Linux).

You don't need third party programs like putty anymore.

Logging into Wheeler



```
ssh vanilla@wheeler.alliance.unm.edu
```

Should prompt you for a password...

Don't let me move on until you are able to login.

Logging into Wheeler

```
ast login: Tue Jun 14 14:47:24 2022 from fricke.co.uk
-----
Welcome to Wheeler

Be sure to review the "Acceptable Use" guidelines posted on the CARC website.

For assistance using this system email help@carc.unm.edu.

Tutorial videos can be accessed through the CARC website: Go to
http://carc.unm.edu, select the "New Users" menu and then click
"Introduction to Computing at CARC".

Warning: By default home directories are world readable. Use the chmod command
to restrict access.

Don't forget to acknowledge CARC in publications, dissertations, theses and
presentations that use CARC computational resources:

"We would like to thank the UNM Center for Advanced Research Computing,
supported in part by the National Science Foundation, for providing the
research computing resources used in this work."

Please send citations to publications@carc.unm.edu.
-----
tarting SSH Key Agent...
gent pid 19486
fricke@wheeler:~ $
```

Please enter the following command



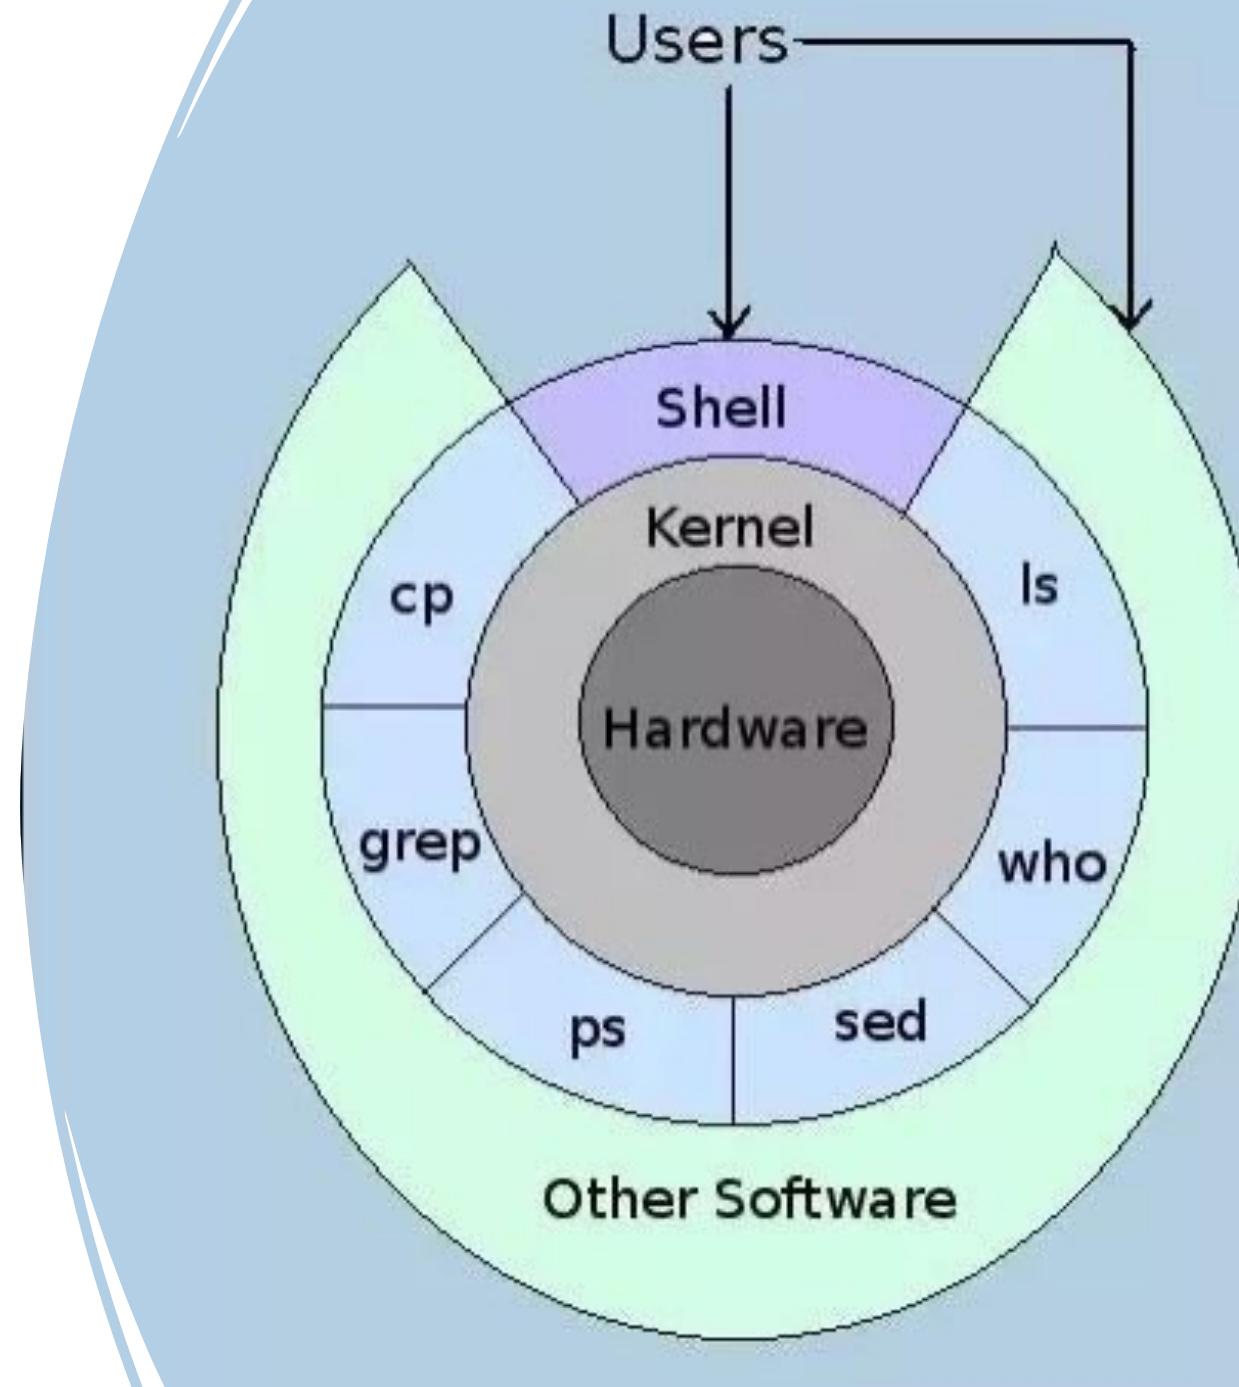
```
cp -r /projects/shared/workshops/beginner/mystuff ~/
```

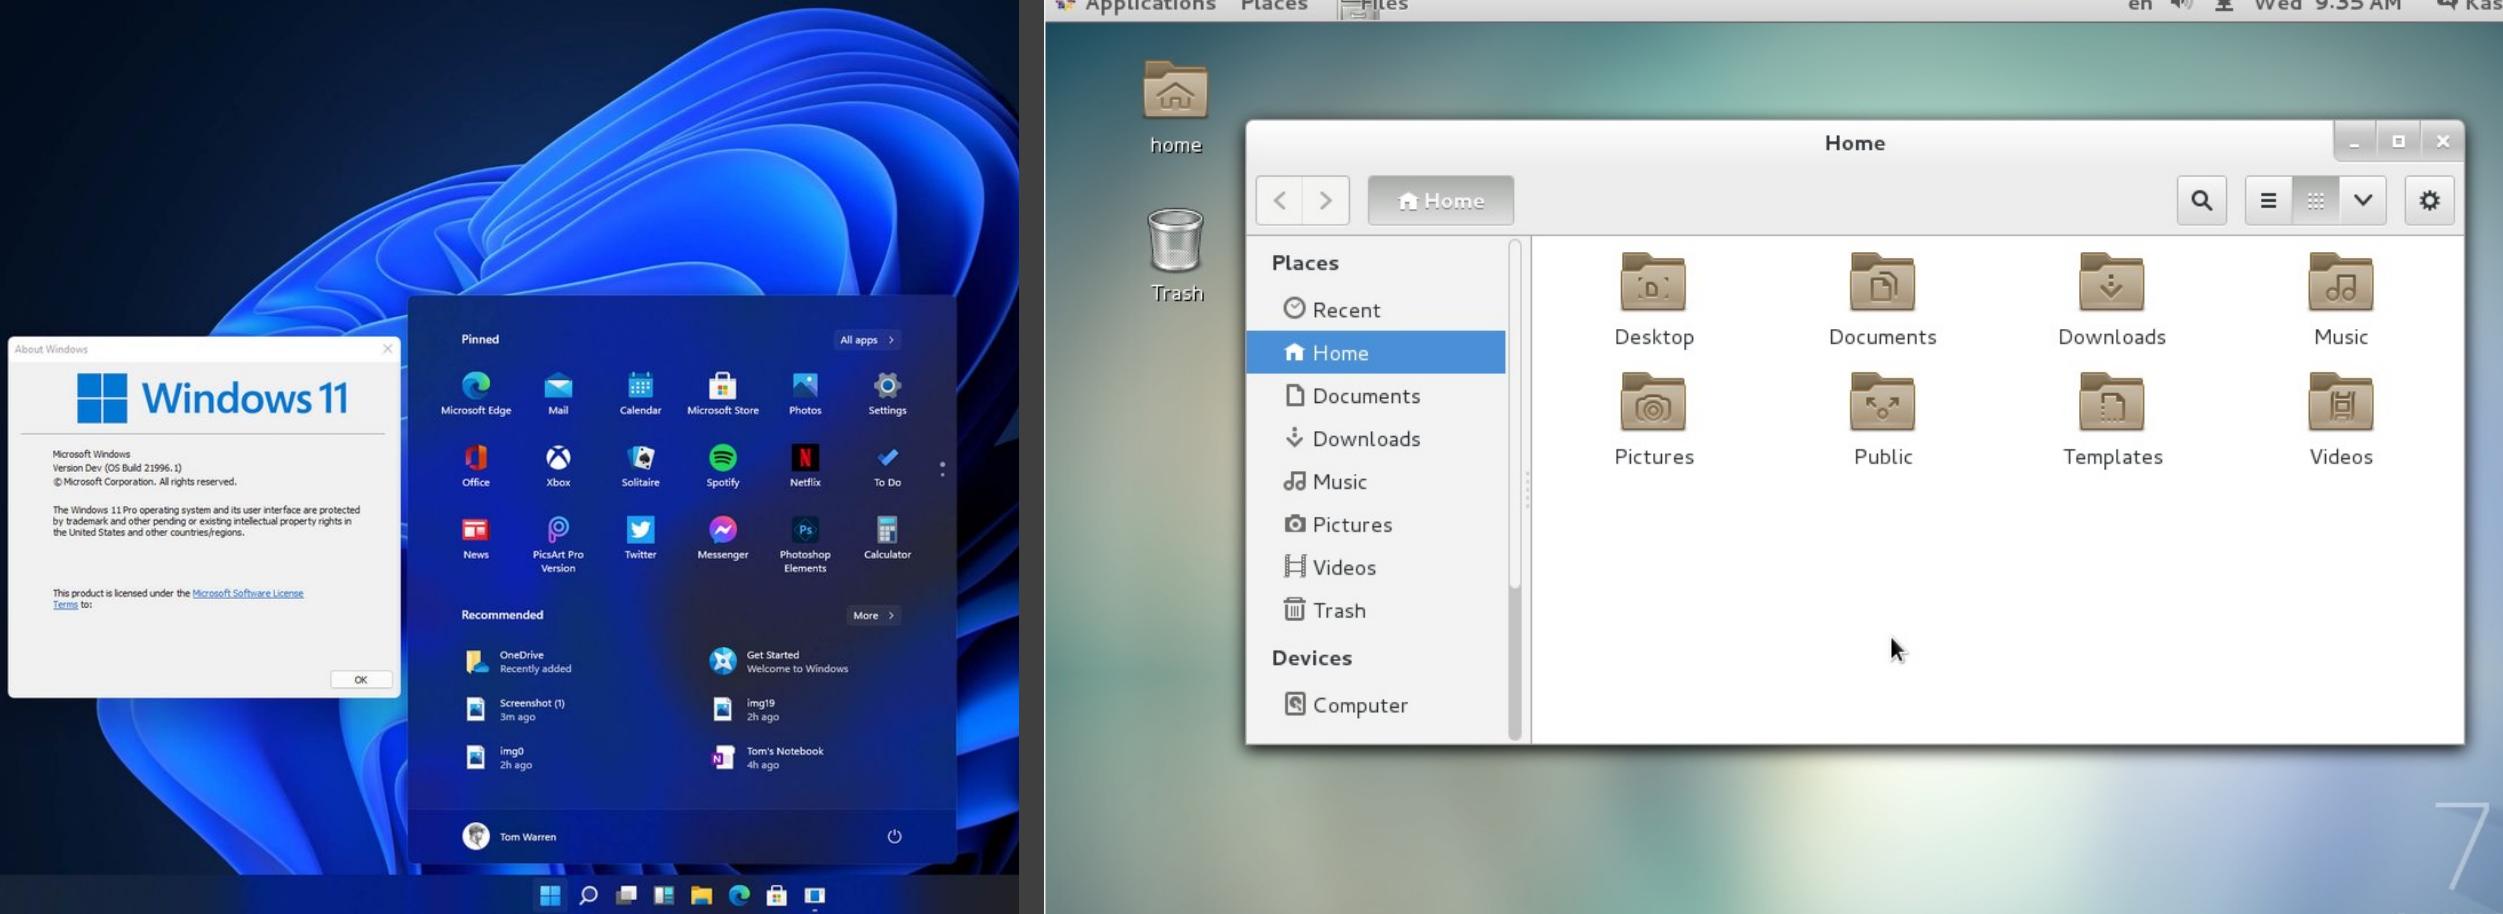
We will come help you if you have any trouble.

(Later I will go over what this command does)

Linux and the BASH Shell

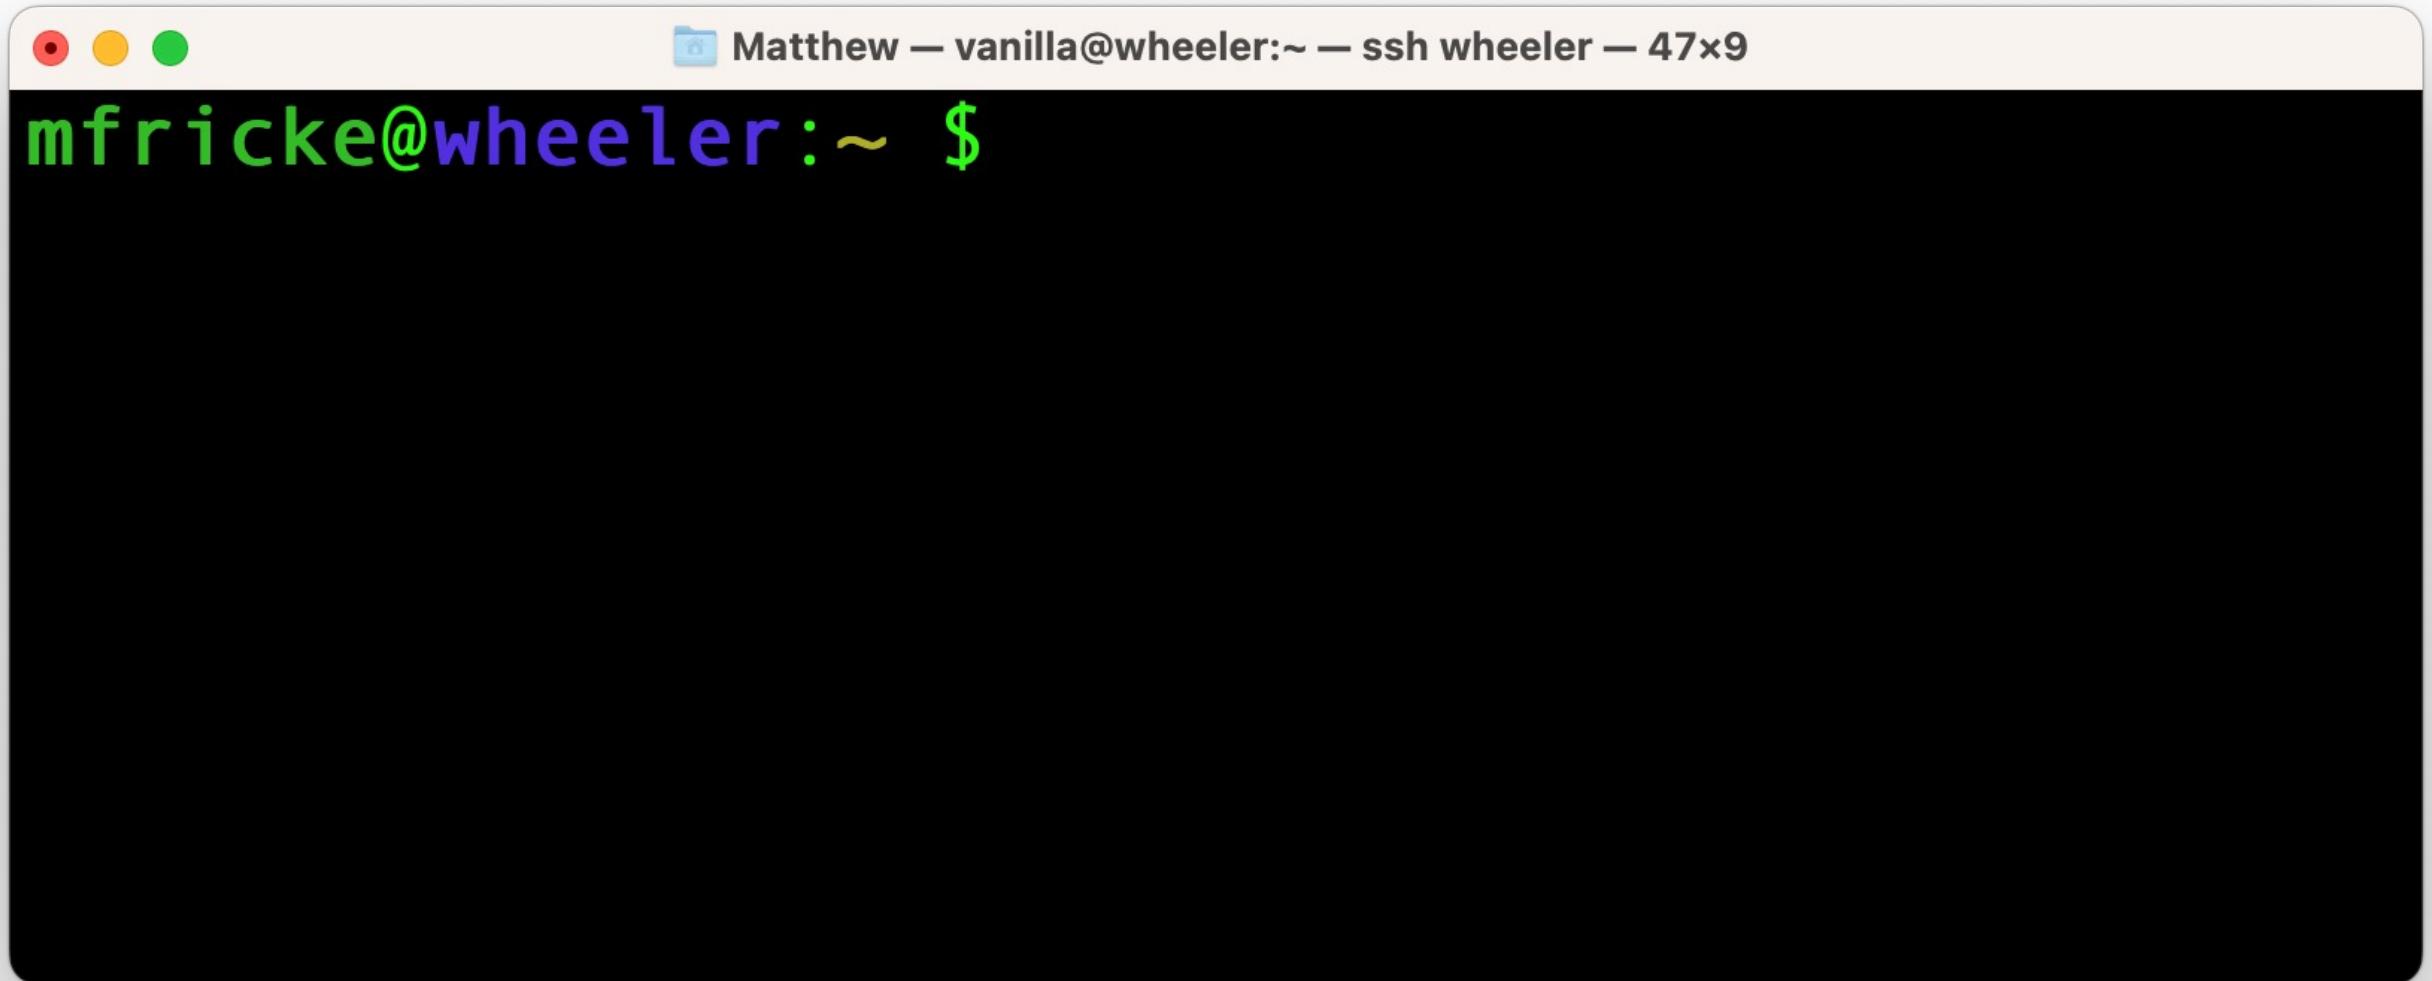
- The Kernel manages access to the hardware in a computer.
- An Operating System (OS) is the Kernel plus useful programs provided by the OS.
- The “shell” is the outermost layer of the OS.
- It is where the user interacts with the OS.





Graphical Shells (GUIs)

Logging into Wheeler



A screenshot of a macOS terminal window. The window has a white header bar with three colored circular icons (red, yellow, green) on the left. The title bar in the center contains a blue folder icon followed by the text "Matthew — vanilla@wheeler:~ — ssh wheeler — 47x9". The main area of the terminal is black and shows the command prompt "mfricke@wheeler:~ \$".

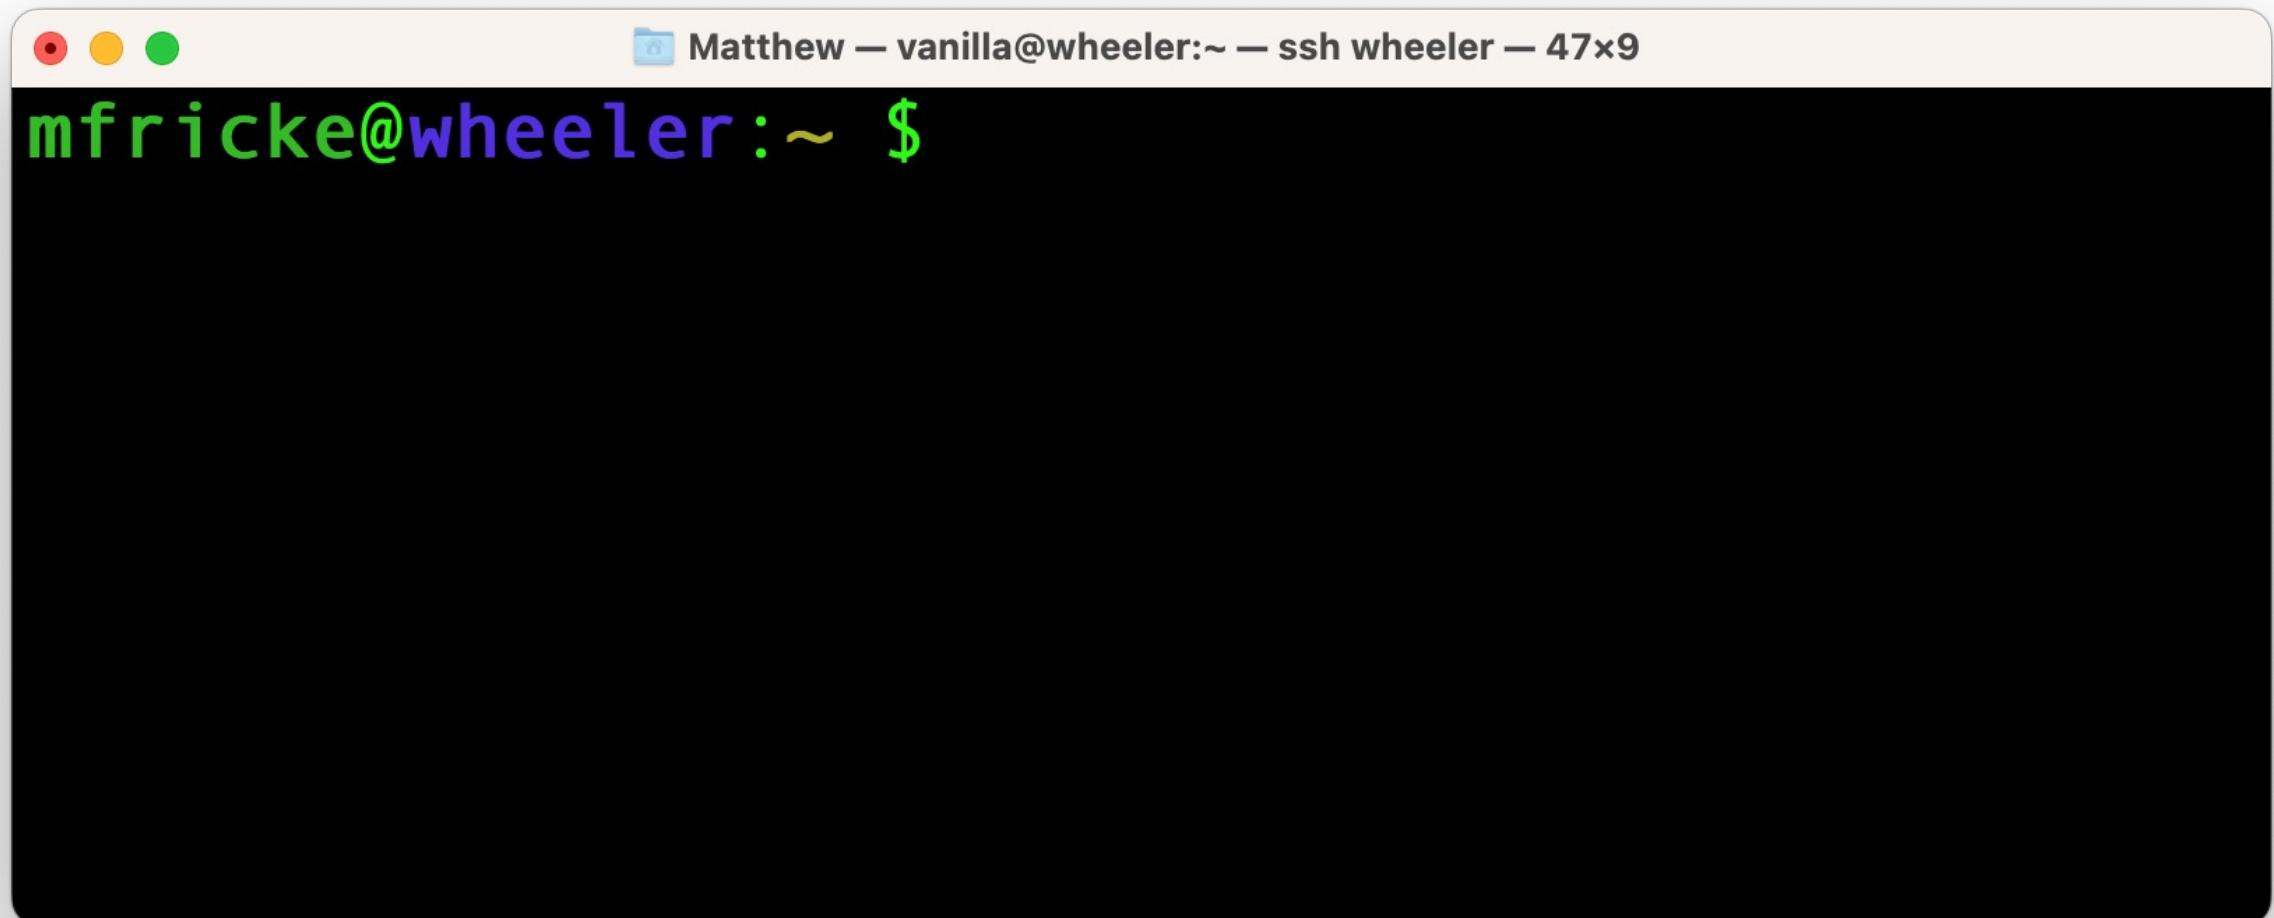
Linux and the BASH Shell



WHERE THERE IS A SHELL, THERE IS A WAY!

The Borne-Again Shell (BASH)

Written in 1976 by Stephen Bourne for UNIX version 7.



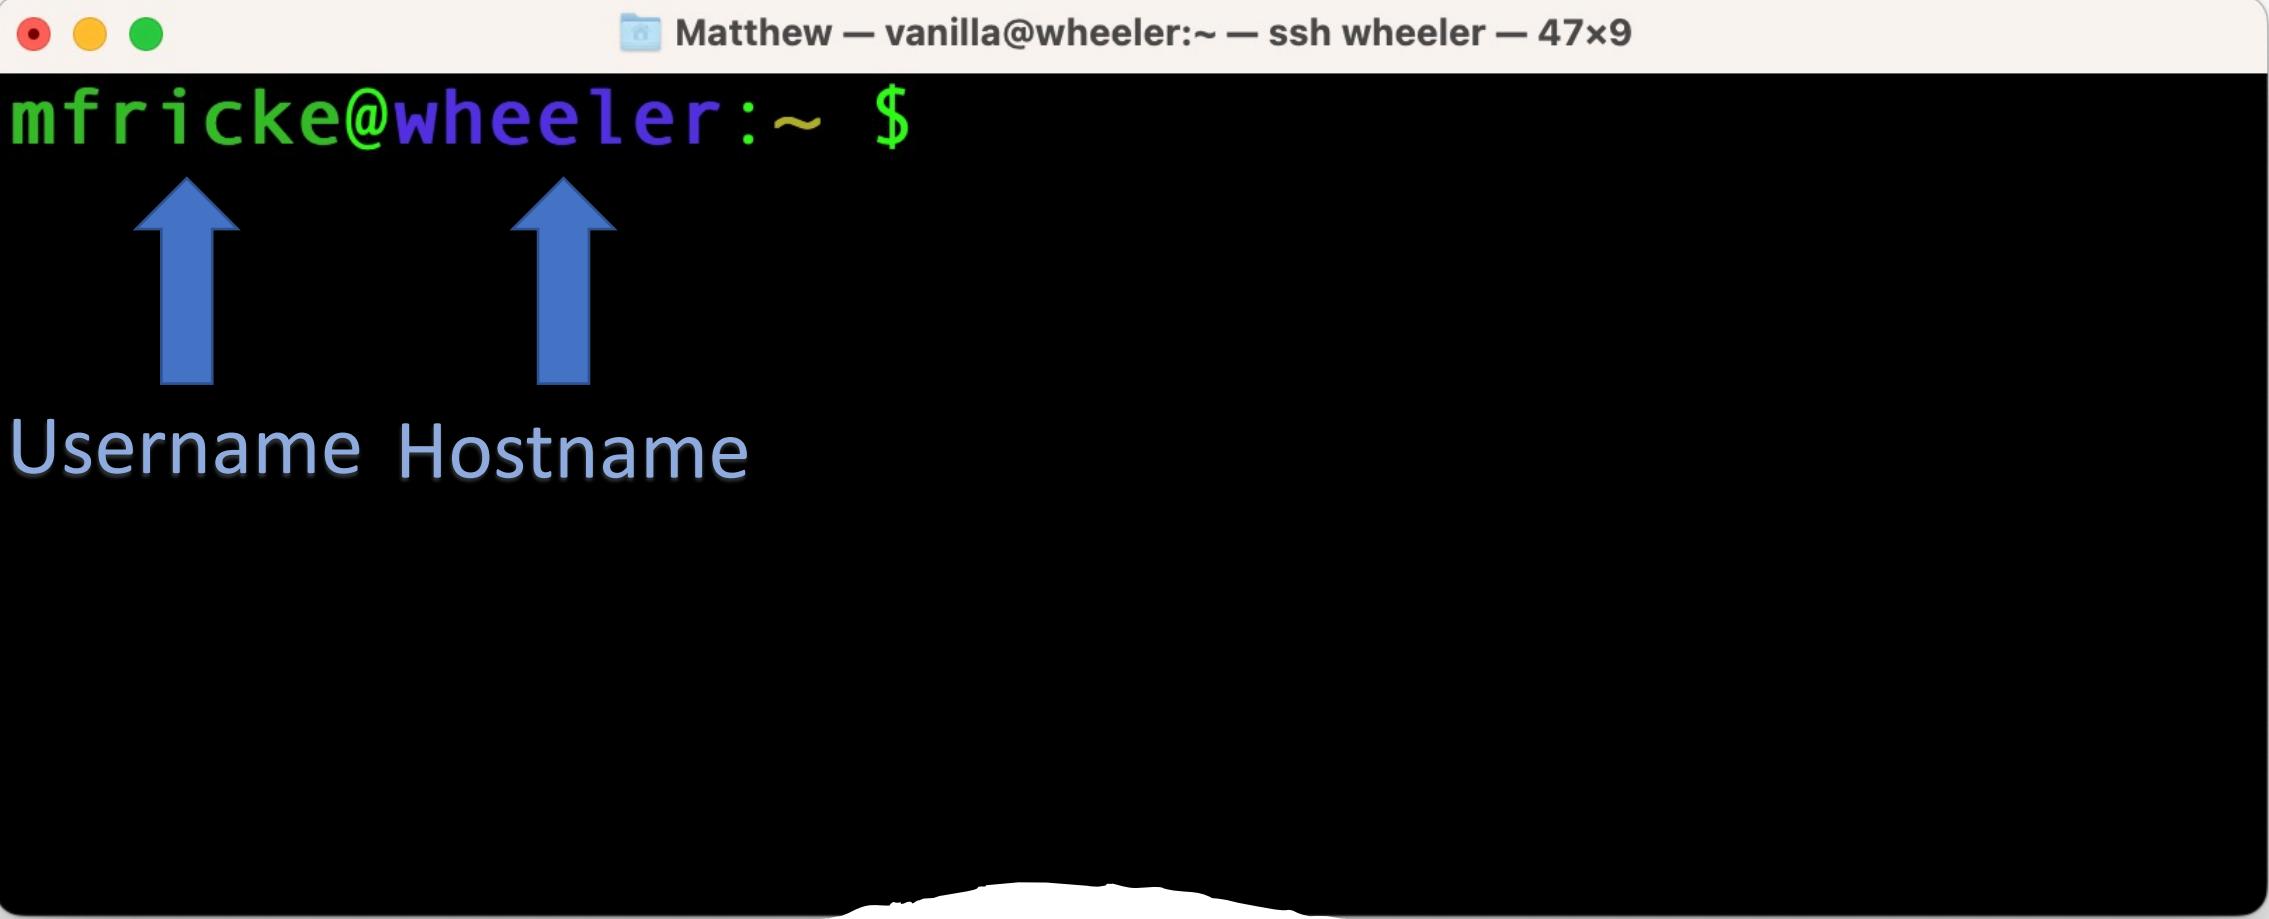
A screenshot of a macOS terminal window. The window title bar shows three colored circular icons (red, yellow, green) on the left, followed by a folder icon and the text "Matthew — vanilla@wheeler:~ — ssh wheeler — 47x9". The main terminal area is black and contains the text "mfricke@wheeler:~ \$".

mfricke@wheeler:~ \$

↑

Username

Understanding the BASH prompt...



mfricke@wheeler:~ \$

↑ ↑

Username Hostname

Understanding the BASH prompt...



Matthew — vanilla@wheeler:~ — ssh wheeler — 47x9

mfricke@wheeler:~ \$



This is the current working directory.
“~” is short for home directory

Understanding the BASH prompt...



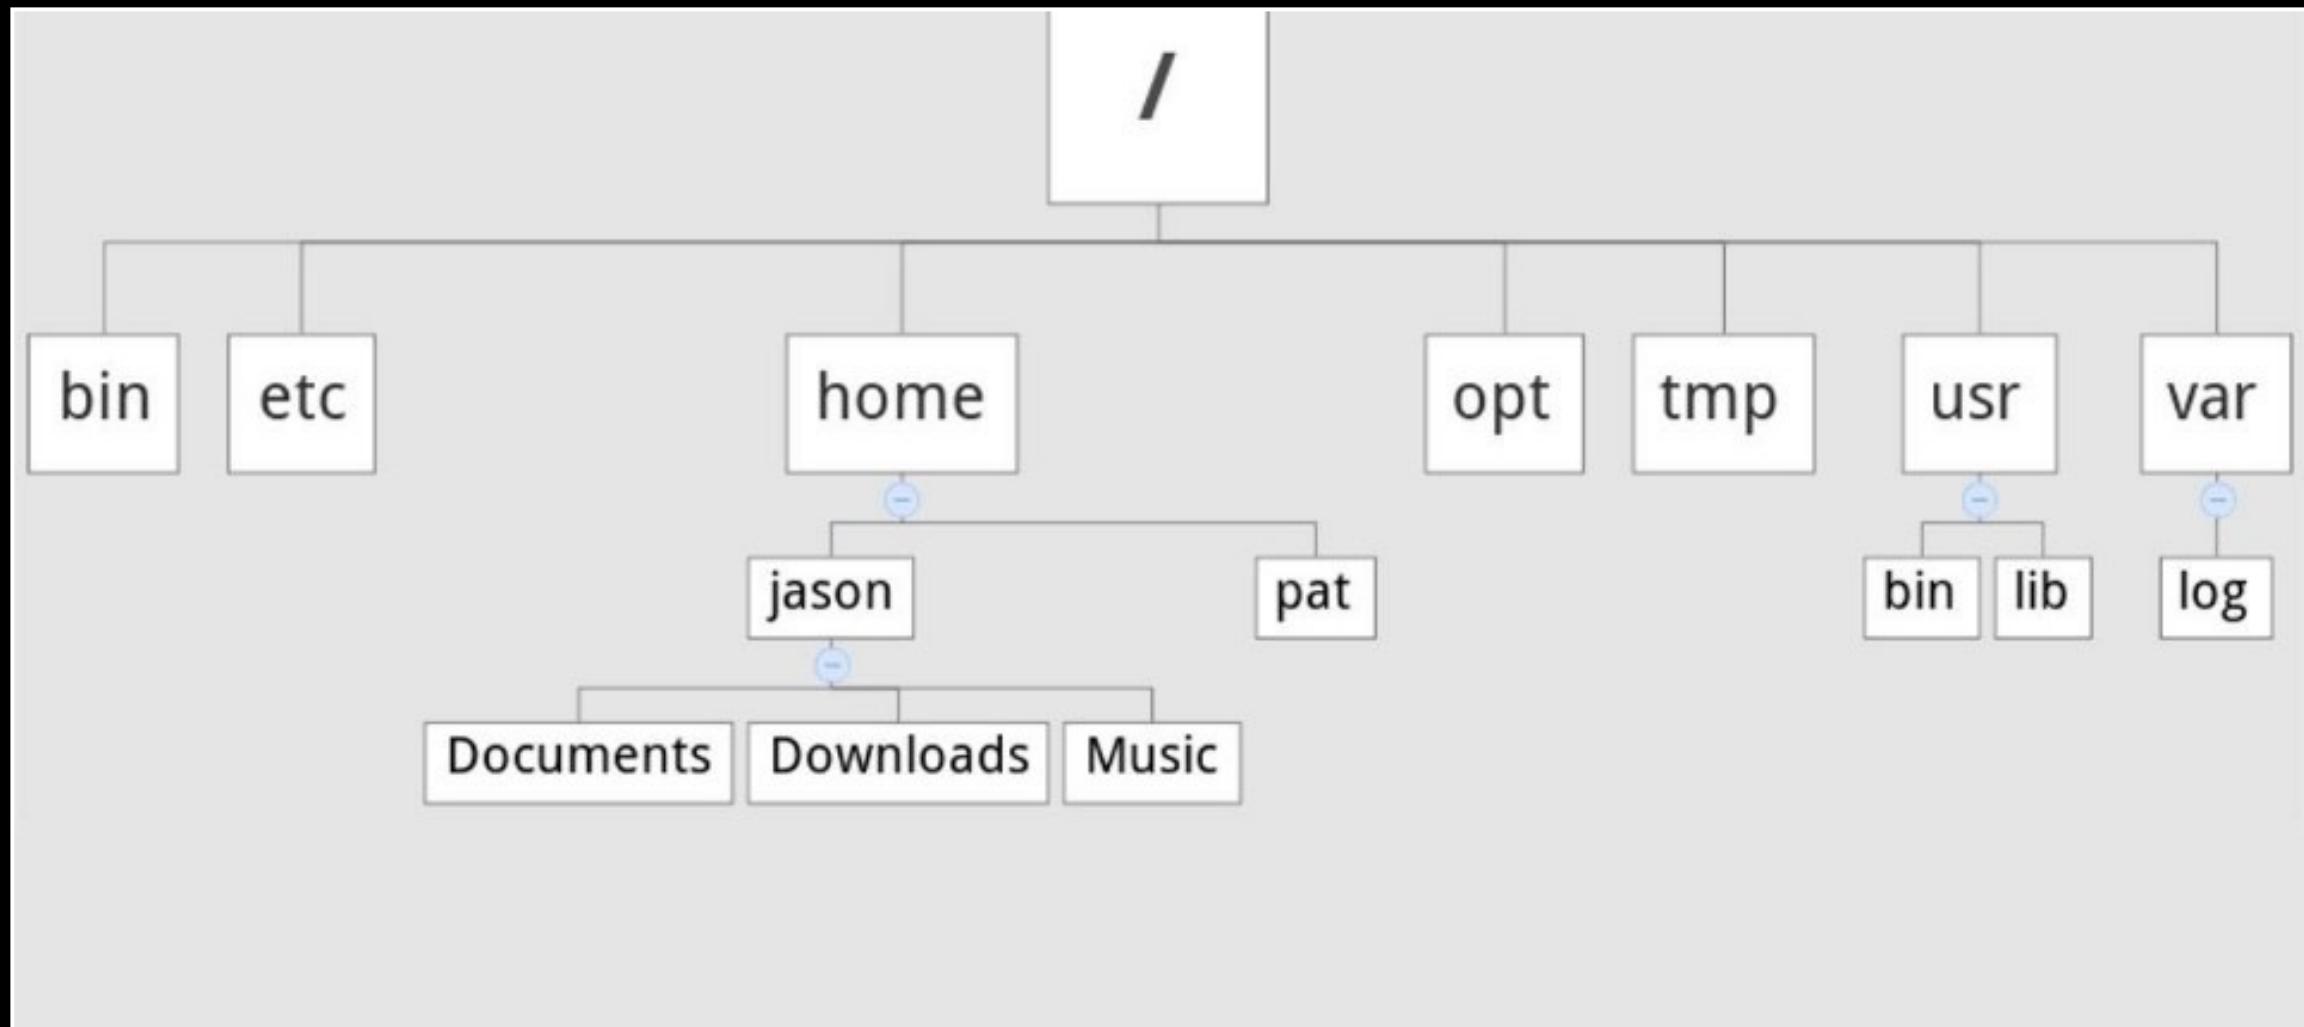
Matthew — vanilla@wheeler:~ — ssh wheeler — 47x9

mfricke@wheeler:~ \$



“\$” means this user is standard user
(i.e. not a system administrator)

Understanding the BASH prompt...



Example Filesystem Tree



Matthew — vanilla@wheeler:~ — ssh wheeler — 42x13

```
[vanilla@wheeler ~]$ pwd  
/users/vanilla  
[vanilla@wheeler ~]$ █
```

Figuring out where you are in the
filesystem... |



Matthew — vanilla@wheeler:~ — ssh wheeler — 42x13

```
[vanilla@wheeler ~]$ ls  
mystuff  wheeler-scratch  
[vanilla@wheeler ~]$
```

Figuring out where you are in the
filesystem...



Matthew — vanilla@wheeler:~ — ssh wheeler — 49x13

```
[vanilla@wheeler ~]$ tree
```

```
.
├── mystuff
│   ├── myfile1
│   └── myfile2
└── wheeler-scratch -> /wheeler/scratch/vanilla
```

```
2 directories, 2 files
```

```
[vanilla@wheeler ~]$
```

Figuring out where you are in
the filesystem...



```
[vanilla@wheeler ~]$ tree
```



```
└── mystuff
    ├── myfile1
    └── myfile2
└── wheeler-scratch -> /wheeler/scratch/vanilla
```

This . means the current directory

2 directories, 2 files

```
[vanilla@wheeler ~]$
```

Figuring out where you are in
the filesystem...

“Absolute” paths vs “relative” paths

- A path is a list of directories and/or files. It is a path through the directory tree that tells one how to get somewhere in the filesystem.
- An absolute path tells one how to get to the destination from starting from the root of the filesystem. E.g “/users/vanilla/mystuff/”
- A relative path specifies how to get there *starting from the current working directory*. E.g vanilla/mystuff/

 Matthew — vanilla@wheeler:~ — ssh wheeler — 49x13

```
[vanilla@wheeler ~]$ ls mystuff/
myfile1 myfile2
[vanilla@wheeler ~]$ █
```

Figuring out where you going...



Matthew — vanilla@wheeler:~ — ssh wheeler — 49x13

```
[vanilla@wheeler ~]$ ls /users/vanilla/mystuff
myfile1 myfile2
[vanilla@wheeler ~]$
```

Figuring out where you going...

```
[vanilla@wheeler ~]$ ls ./mystuff/  
myfile1 myfile2  
[vanilla@wheeler ~]$ ls ~/mystuff/  
myfile1 myfile2  
[vanilla@wheeler ~]$
```

Figuring out where you going...

```
[vanilla@wheeler ~]$ ls -a
.
..
.
.addressbook
.addressbook.lu
.bashrc
.cache
.comsol
.config
.flexlmrc
.modulesbeginenv
mystuff
.oracle_jre_usage
.pinerc
.pki
.rhosts
.shosts
.spack
.ssh
```

Figuring out where you going...

```
[vanilla@wheeler ~]$ ls -l
total 4
drwxr-xr-x 2 vanilla users 4096 Jun 14 22:05 mystuff
lrwxrwxrwx 1 vanilla users    24 Jun 14 21:20 wheeler-scratch ->
/wheeler/scratch/vanilla
[vanilla@wheeler ~]$
```

Figuring out where you going...

```
[vanilla@wheeler ~]$ ls -l mystuff/
total 473704
-rw-r--r-- 1 vanilla users 483165473 Jun 14 23:20 myfile1
-rw-r--r-- 1 vanilla users          0 Jun 14 22:05 myfile2
[vanilla@wheeler ~]$
```

Figuring out where you going...

```
[vanilla@wheeler ~]$ ls -lh mystuff/
total 463M
-rw-r--r-- 1 vanilla users 461M Jun 14 23:20 myfile1
-rw-r--r-- 1 vanilla users     0 Jun 14 22:05 myfile2
[vanilla@wheeler ~]$
```

Figuring out where you going...

```
[vanilla@wheeler ~]$ du -s
```

```
499704 .
```

```
[vanilla@wheeler ~]$ du -sh
```

```
488M .
```

```
[vanilla@wheeler ~]$
```

Figuring out where you going...

```
[vanilla@wheeler ~]$ df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
devtmpfs	24G	0	24G	0%	/dev
tmpfs	24G	64K	24G	1%	/dev/shm
tmpfs	24G	968M	23G	5%	/run
tmpfs	24G	0	24G	0%	/sys/fs/cgroup
/dev/mapper/centos-root	930G	567G	363G	61%	/
/dev/sdc2	836G	72G	764G	9%	/tmp
/dev/md126p1	2.0G	333M	1.7G	17%	/boot
172.17.2.254:/mnt/wheeler-scratch	37T	28T	8.7T	77%	/wheeler/scratch
172.17.2.255:/mnt/wheeler-scratch2	37T	28T	9.0T	76%	/wheeler/scratch2
beegfs_nodev	110T	51T	60T	46%	/carc/scratch
chama:/home/homes	65T	36T	30T	55%	/users
chama:/home/carc_projects	65T	36T	30T	55%	/projects

Figuring out where you going...

```
[vanilla@wheeler ~]$ quota -S
Disk quotas for user vanilla (uid 659):
  Filesystem    space   quota   limit   grace   files   quota   limit   grace
chama:/home/homes
                  488M    100G    200G
[vanilla@wheeler ~]$
```

Figuring out where you going...

```
[vanilla@wheeler ~]$ stat mystuff/myfile1
  File: 'mystuff/myfile1'
  Size: 483165473 Blocks: 947408      IO Block: 65536   regular file
Device: 28h/40d Inode: 9232782834205560540 Links: 1
Access: (0644/-rw-r--r--) Uid: ( 659/ vanilla)  Gid: ( 100/ users)
Access: 2022-06-14 22:05:27.503289000 -0600
Modify: 2022-06-14 23:20:26.945918000 -0600
Change: 2022-06-14 23:20:48.754917000 -0600
 Birth: -
```

Figuring out what you've got...

```
[vanilla@wheeler ~]$ find -name myfile2  
./mystuff/myfile2
```

```
[vanilla@wheeler ~]$ find -name "myfile*"  
./mystuff/myfile1  
./mystuff/myfile2  
./mystuff/myfile3  
./mystuff/myfile0
```



Wildcard

Figuring out what you've got...

```
[vanilla@wheeler ~]$ cd mystuff/  
[vanilla@wheeler ~/mystuff]$
```

Use the tab key to autocomplete

Going somewhere new...

Now it is your turn...



- For this path:

`/projects/shared/workshops/beginner/vecadd`

- What are the names of the files in that directory?
- When were they last modified?
- How large are the files?

You can find this information with the `ls` command.

Now it is your turn...



- For this path:

/projects/shared/workshops/beginner/vecadd

Now “cd” into that directory using <tab> autocomplete.

- Now you know how to find your way around filesystems using bash
- Let's see how to modify the filesystem.
 - In bash to move a file we use the `mv` command.
 - To copy a file it is `cp`.
 - To copy files from CARC to a personal computer use `scp` or `rsync`.

```
[vanilla@wheeler beginner]$ pwd  
/projects/shared/workshops/beginner  
[vanilla@wheeler beginner]$ cd ~  
[vanilla@wheeler ~]$ pwd  
/users/vanilla  
[vanilla@wheeler ~]$
```

First return to your home
directory...

```
[vanilla@wheeler ~]$ cd mystuff  
[vanilla@wheeler ~/mystuff]$ mv myfile1 myfile0  
[vanilla@wheeler ~/mystuff]$ ls  
myfile0 myfile2 myfile3  
[vanilla@wheeler ~/mystuff]
```

Modifying the filesystem...
moving a file.

```
[vanilla@wheeler ~/mystuff]$ cp myfile0 myfile1  
[vanilla@wheeler ~/mystuff]$
```



Source Destination

```
[vanilla@wheeler ~/mystuff]$ ls  
myfile0 myfile1 myfile2 myfile3  
[vanilla@wheeler ~/mystuff]$
```

Modifying the filesystem...
copying a file.

```
[vanilla@wheeler ~]$ cp -r mystuff mystuff2  
[vanilla@wheeler ~]$
```



Source



Destination

```
[vanilla@wheeler ~]$ ls  
mystuff  mystuff2  wheeler-scratch
```

Copying a whole directory tree...

```
Lycaon:~ matthew$ scp vanilla@wheeler.alliance.unm.edu:~/mystuff/myfile3 Desktop/
```



Source

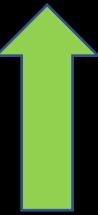


Destination

```
(vanilla@wheeler.alliance.unm.edu) Password:  
myfile3          100%    40      2.0KB/s   00:00
```

Copying data to a personal
computer from CARC...

```
Lycaon:~ matthew$ scp -r vanilla@wheeler.alliance.unm.edu:~/mystuff Desktop/
```



Source



Destination

```
(vanilla@wheeler.alliance.unm.edu) Password:
```

myfile1	100%	1024KB	6.5MB/s	00:00
myfile2	100%	2048KB	382.5KB/s	00:05
myfile3	100%	40	3.2KB/s	00:00
myfile0	100%	1024KB	8.8MB/s	00:00

Copying data to a personal
computer from CARC...

```
Lycaon:~ matthew$ scp -r Desktop/mystuff vanilla@wheeler.alliance.unm.edu:~/
```



Source



Destination

```
(vanilla@wheeler.alliance.unm.edu) Password:
```

myfile1	100%	1024KB	591.5KB/s	00:01
myfile0	100%	1024KB	2.0MB/s	00:00
myfile2	100%	2048KB	2.1MB/s	00:00
myfile3	100%	40	2.1KB/s	00:00

To copy from a personal
computer to CARC...

```
ssh vanilla@wheeler.alliance.unm.edu
```

Log back into wheeler...

```
[vanilla@wheeler ~]$ file mystuff/myfile0  
mystuff/myfile0: data
```

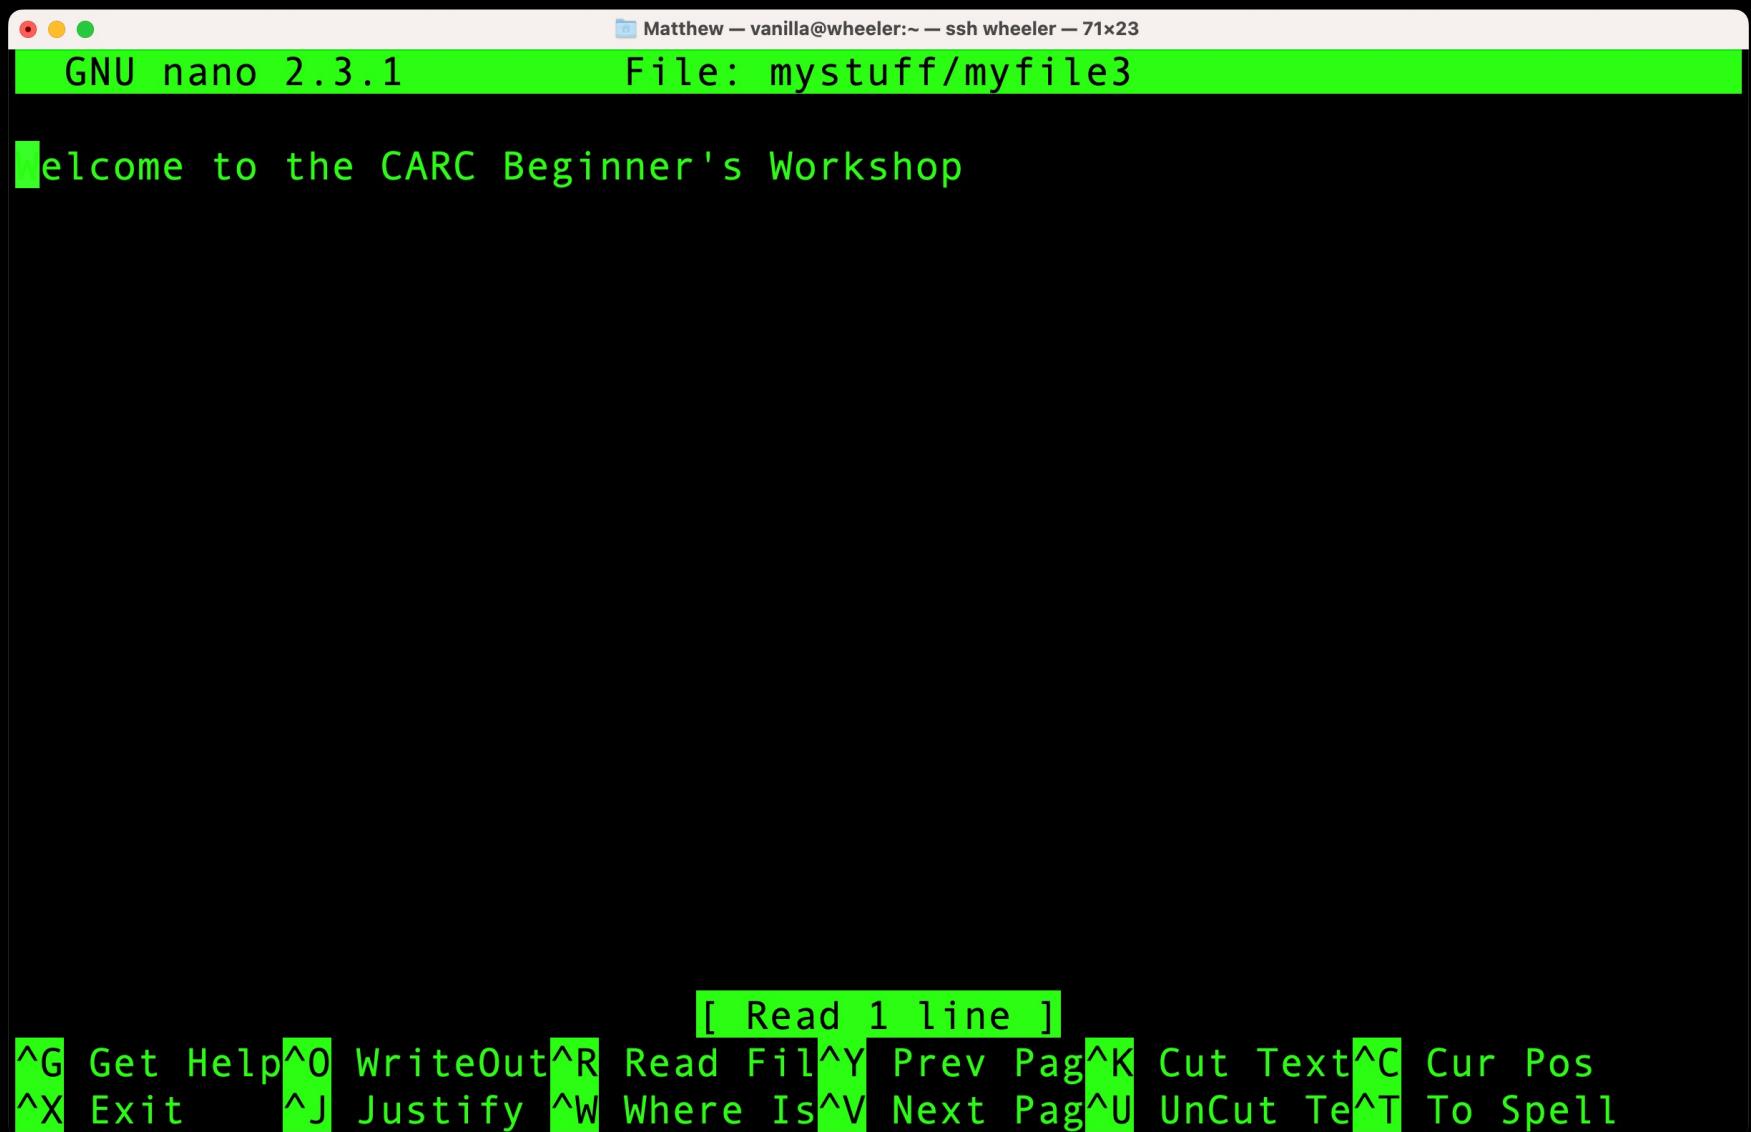
```
[vanilla@wheeler ~]$ file mystuff/myfile3  
mystuff/myfile3: ASCII text
```

Figuring out file types ...

```
[vanilla@wheeler ~]$ cat mystuff/myfile3  
Welcome to the CARC Beginner's Workshop
```

Text files ...

```
[vanilla@wheeler ~]$ nano mystuff/myfile3
```



The screenshot shows a terminal window titled "Matthew — vanilla@wheeler:~ — ssh wheeler — 71x23". Inside the terminal, the command "nano mystuff/myfile3" is run, opening the file "myfile3". The file contains the text "Welcome to the CARC Beginner's Workshop". At the bottom of the screen, the nano command-line interface is visible, showing various keyboard shortcuts for editing.

```
GNU nano 2.3.1          File: mystuff/myfile3
Welcome to the CARC Beginner's Workshop

[ Read 1 line ]
^G Get Help ^O WriteOut ^R Read Fil^Y Prev Pag ^K Cut Text ^C Cur Pos
^X Exit      ^J Justify ^W Where Is ^V Next Pag ^U UnCut Te^T To Spell
```

15 Minute Break

Software Access

Lmod
Modules

Conda

```
[vanilla@wheeler ~]$ module spider matlab
```

```
- matlab:
```

```
- Versions:
```

```
  matlab/R2017a
  matlab/R2018b
  matlab/R2019a
  matlab/R2020a
  matlab/R2021a
```

Getting access to software...

```
[vanilla@wheeler ~]$ module load matlab/R2021a
Lmod has detected the following error: Matlab may only be run on compute
nodes. wheeler is not a compute node. Exiting...
While processing the following module(s):
  Module fullname      Module Filename
  -----      -----
matlab/R2021a      /opt/local/modules/matlab/R2021a.lua
```

Getting access to software...

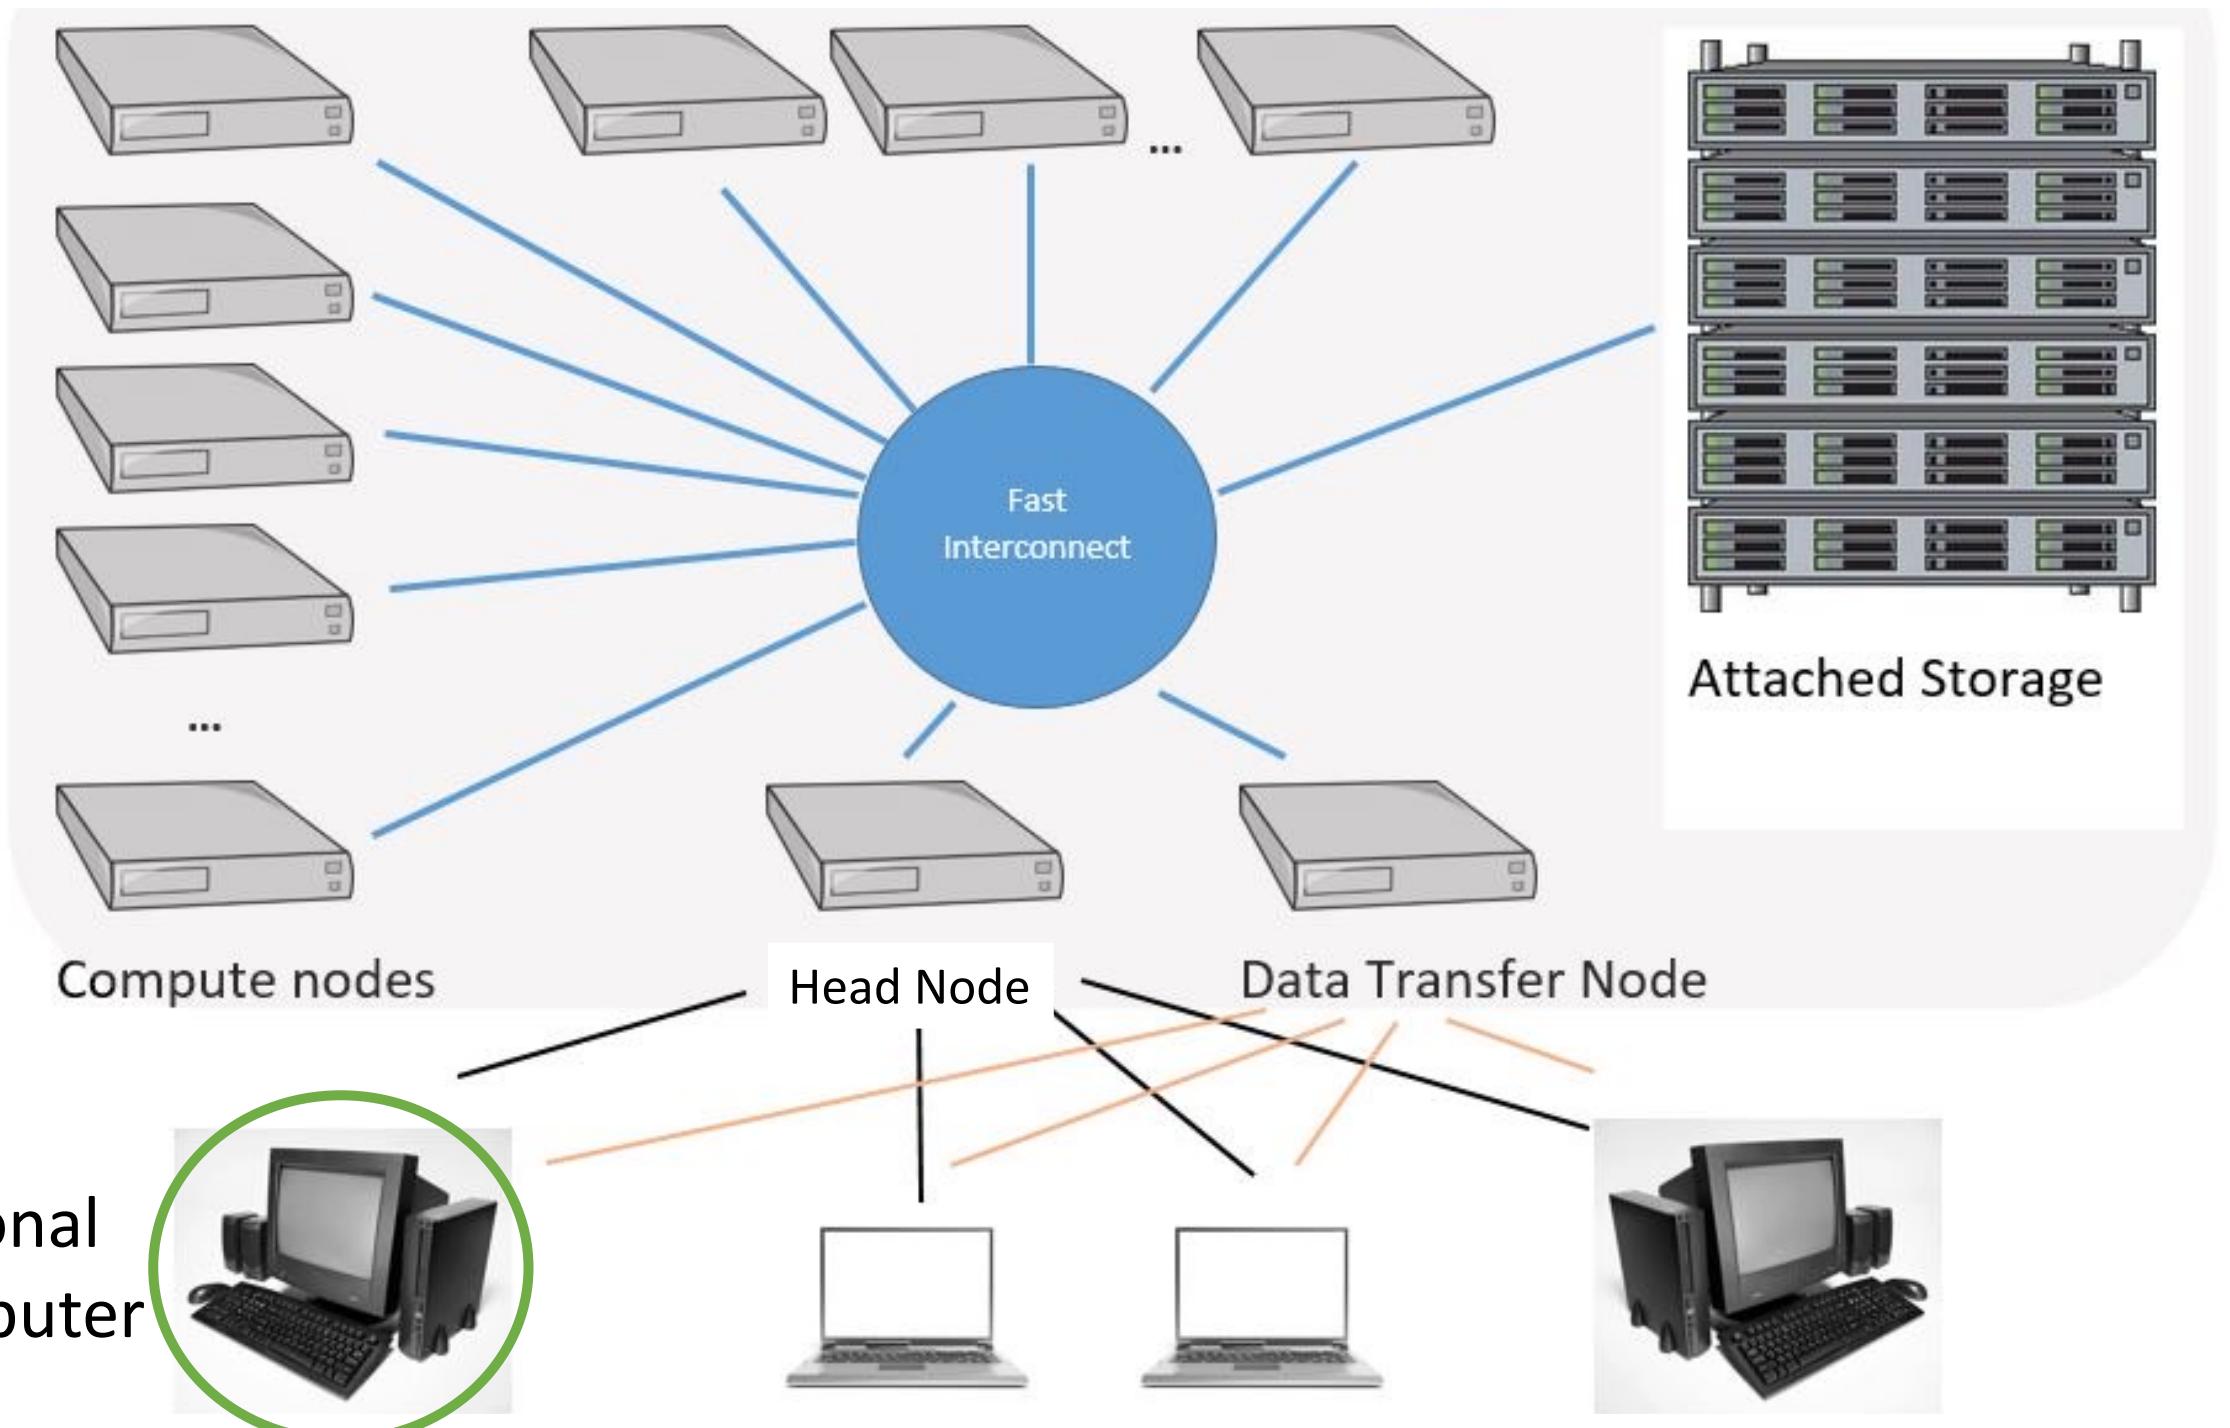
```
[vanilla@wheeler ~]$ module load matlab/R2021a
Lmod has detected the following error: Matlab may only be run on compute
nodes. wheeler is not a compute node. Exiting...
While processing the following module(s):
  Module fullname      Module Filename
  -----      -----
matlab/R2021a      /opt/local/modules/matlab/R2021a.lua
```

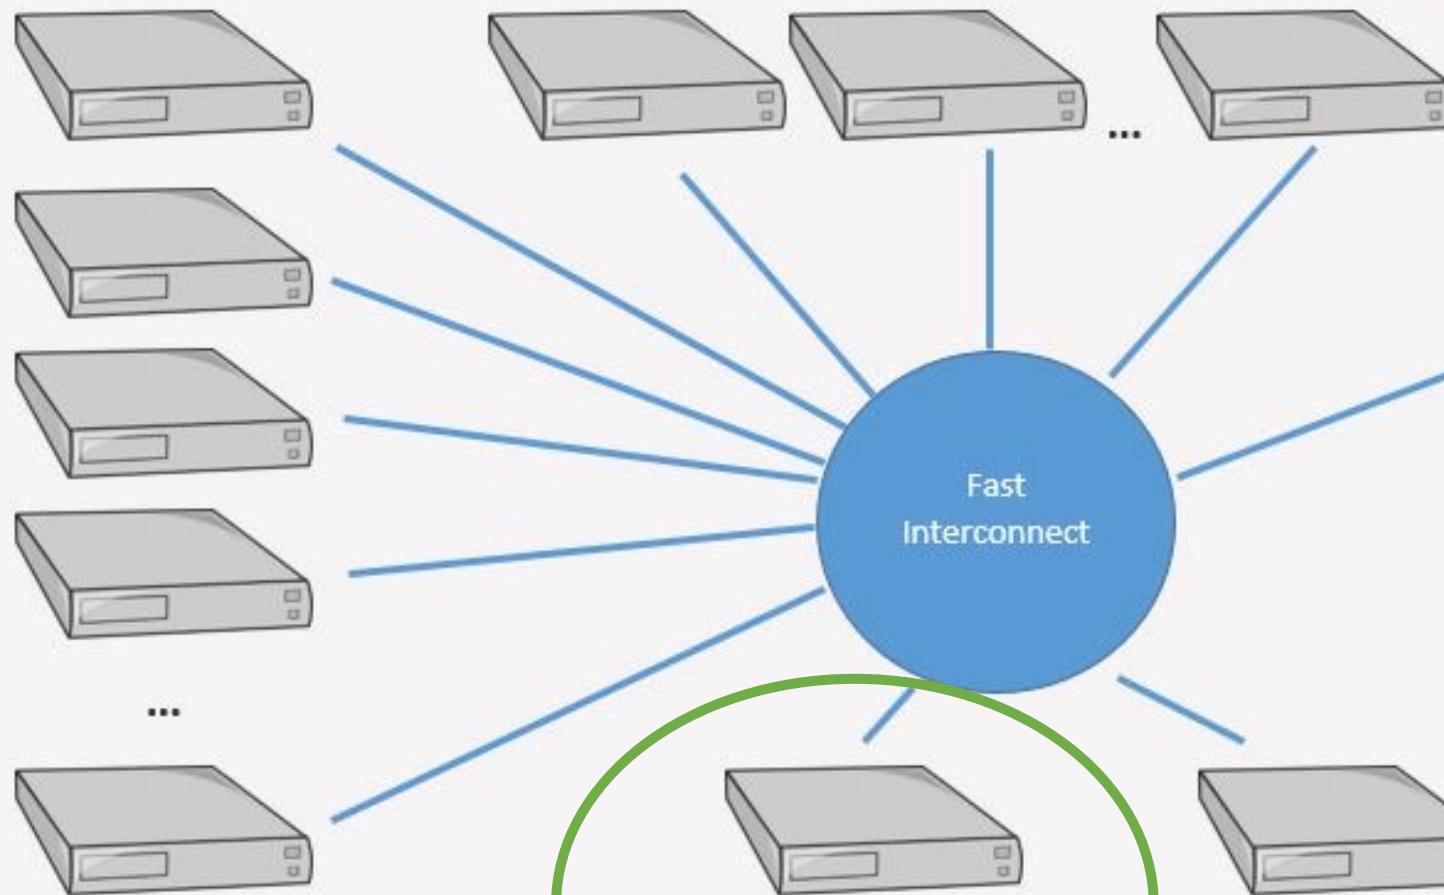
What is a compute node?

Getting access to software...

HPC Cluster





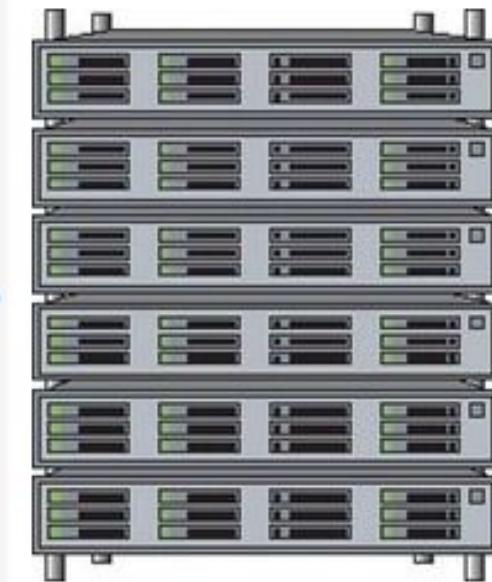


Compute nodes

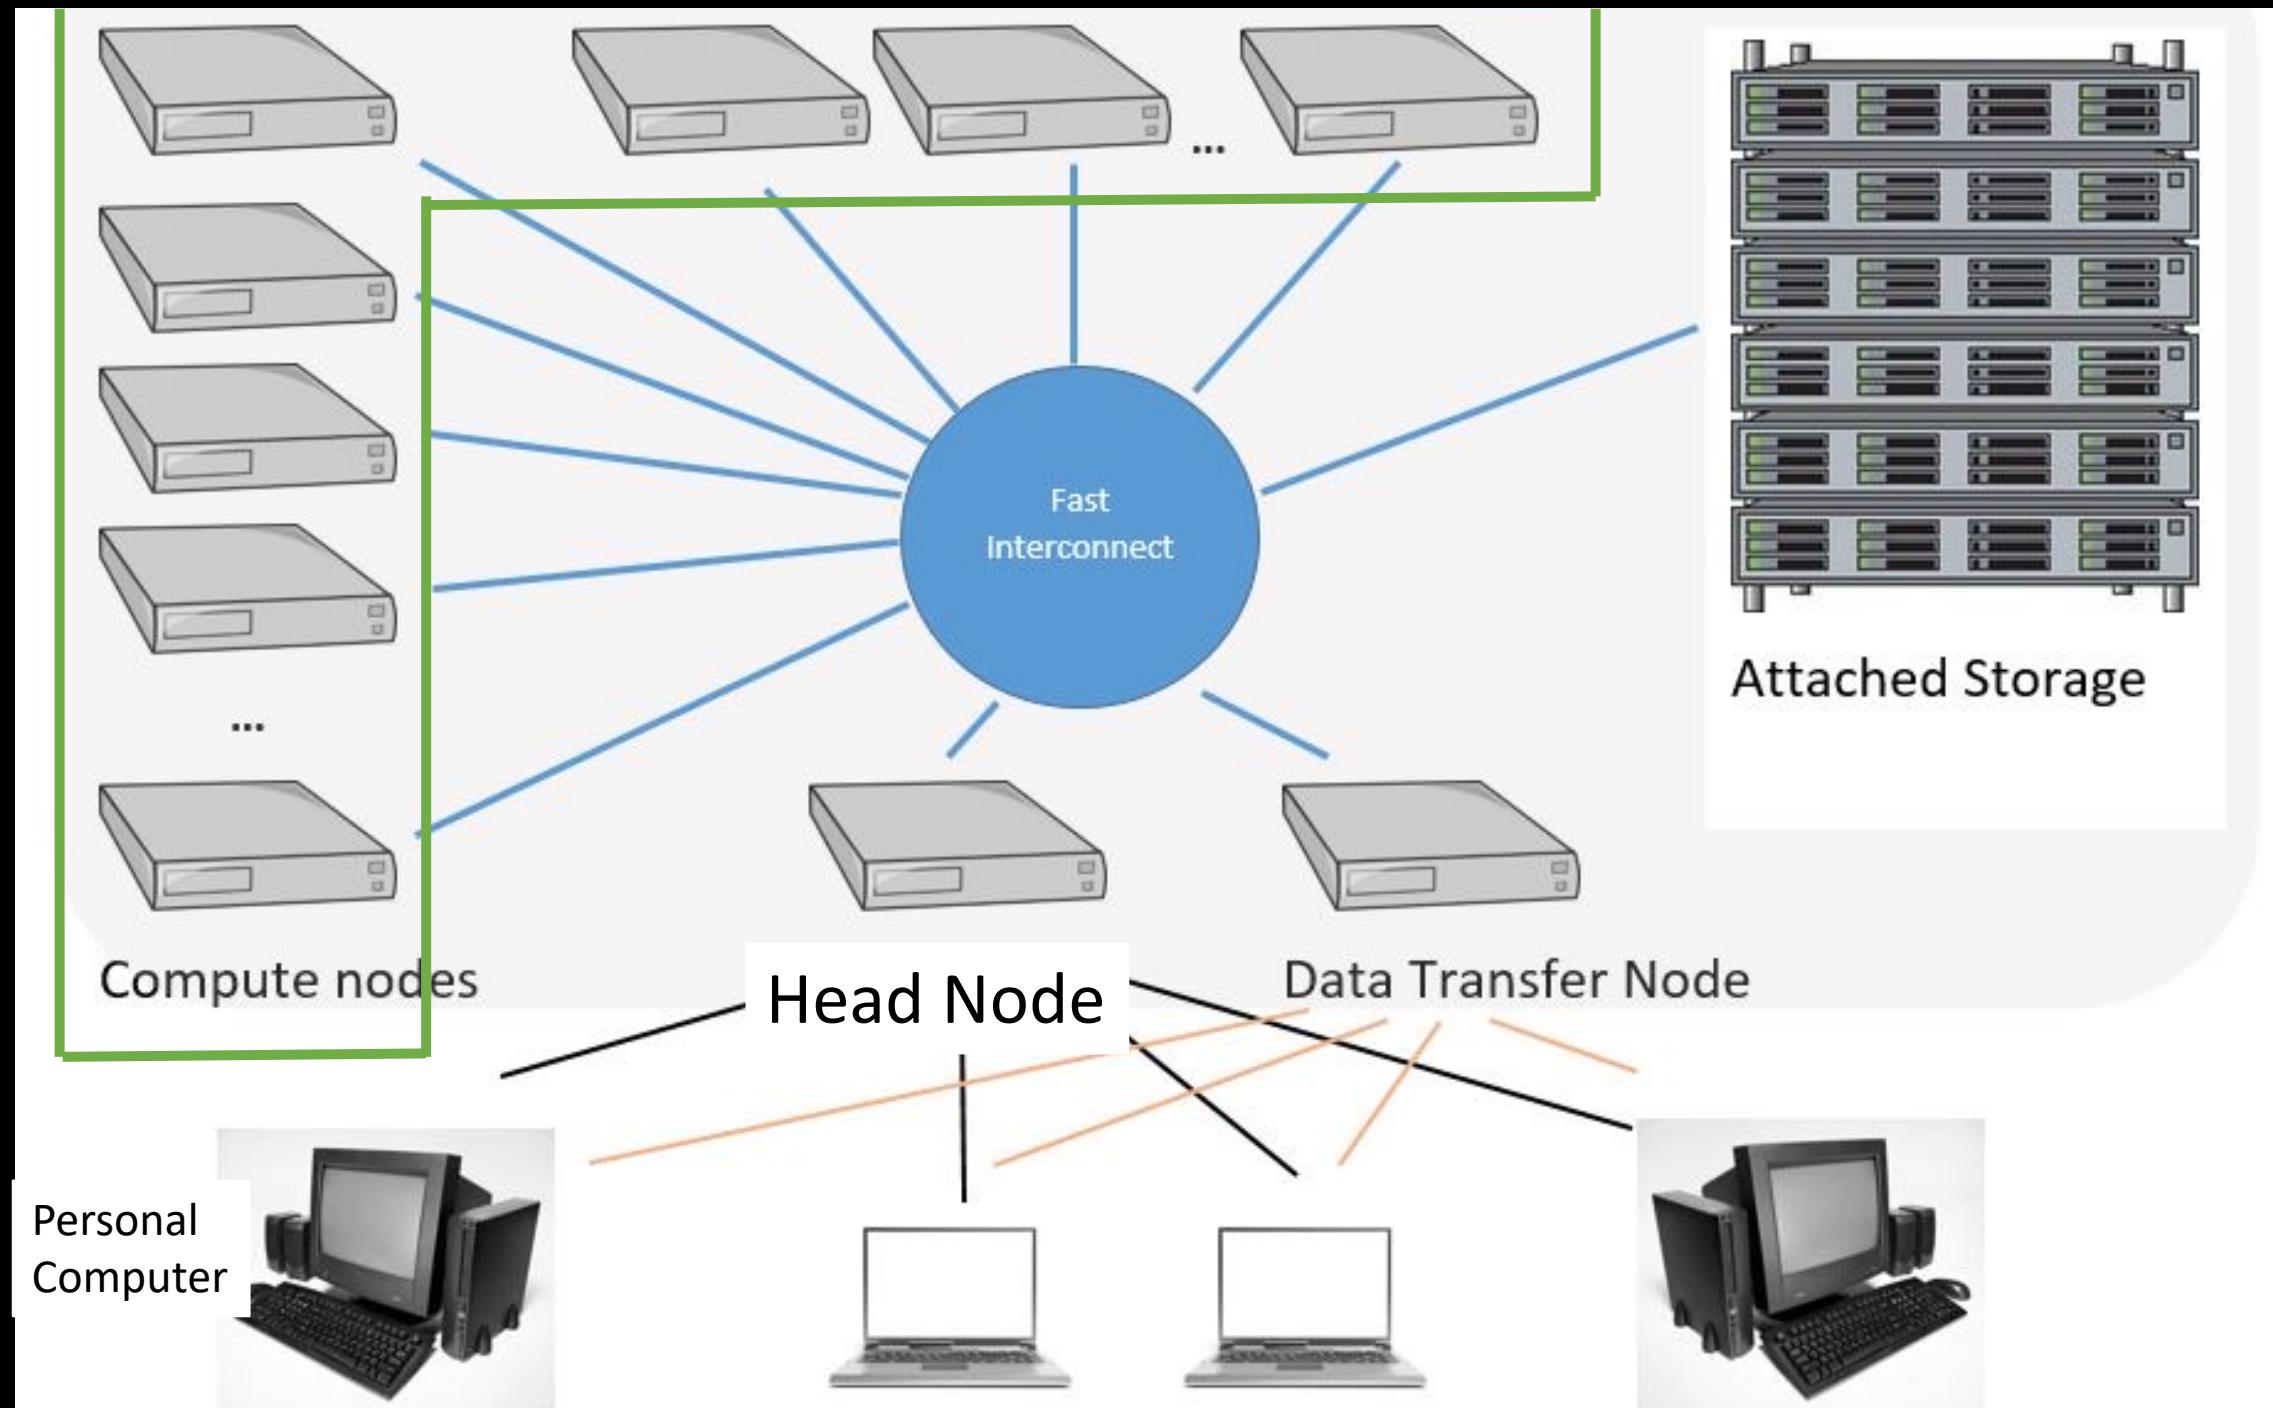
Head Node

Data Transfer Node

Personal Computer



Attached Storage





WARNING !!

Never run computations on the head node

Always use compute nodes

```
[vanilla@wheeler ~]$ qgrok
```

queues	free	busy	offline	jobs	nodes	CPUs
-----	-----	-----	-----	-----	-----	-----
normal	0	299	1	97	300	2400
debug	4	0	0	0	4	32
totals:	4	299	1	97	304	2432

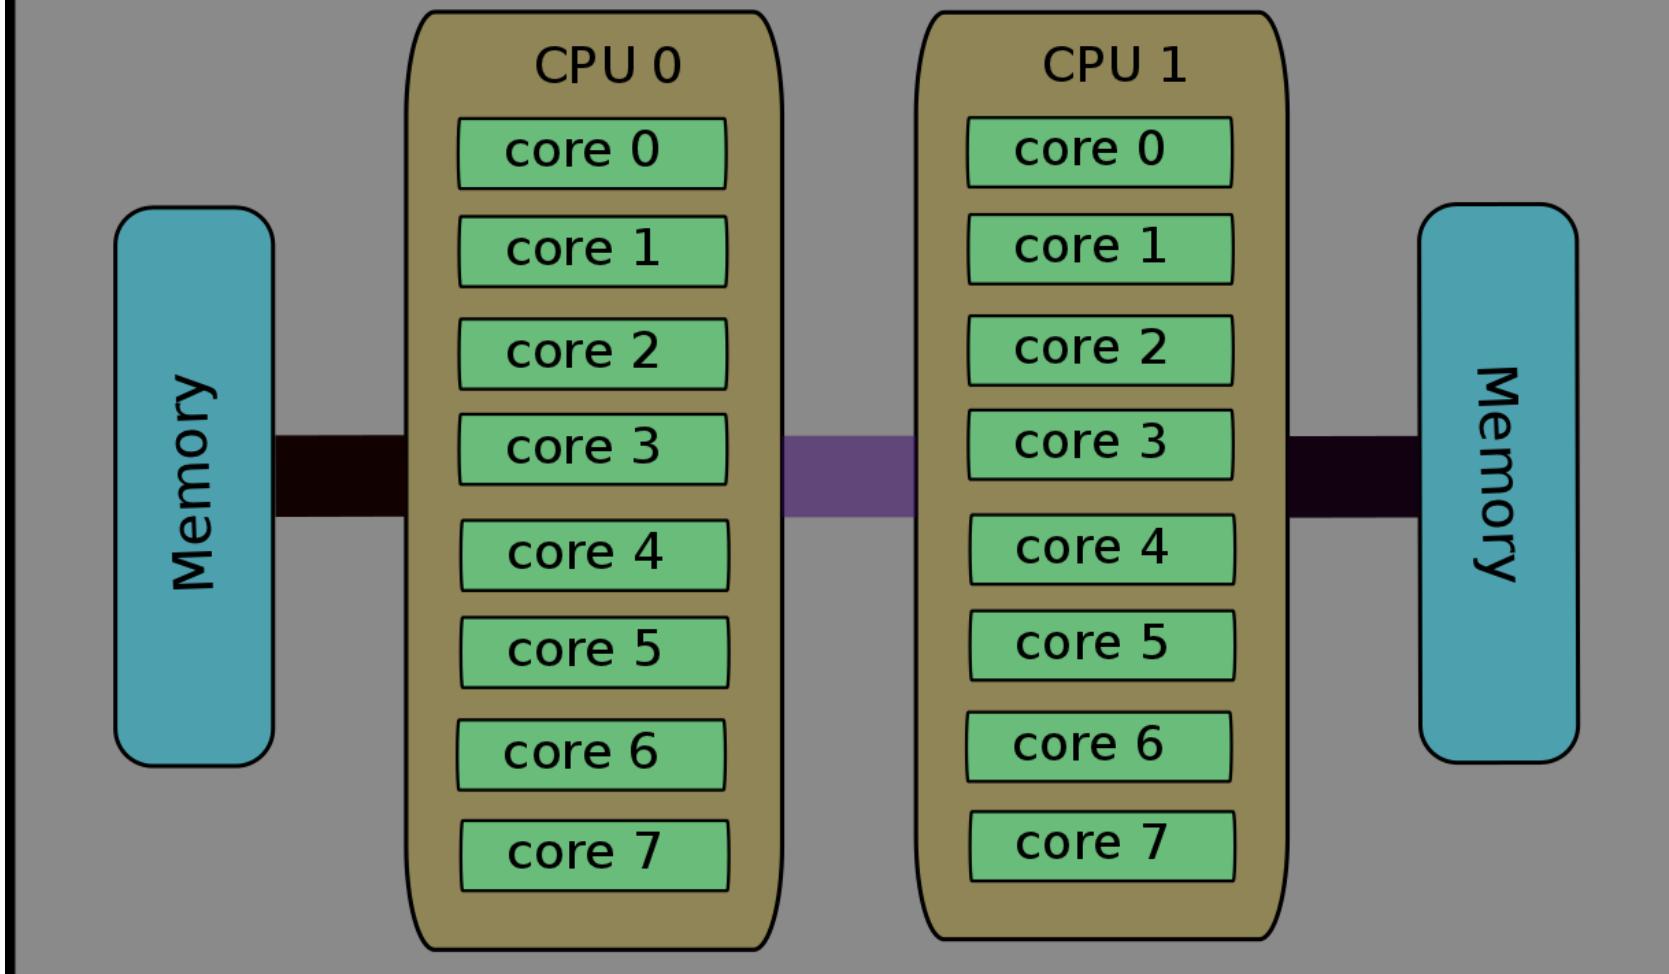
Compute nodes and partitions...

```
qmfricke@hopper:~ $ qgrok
```

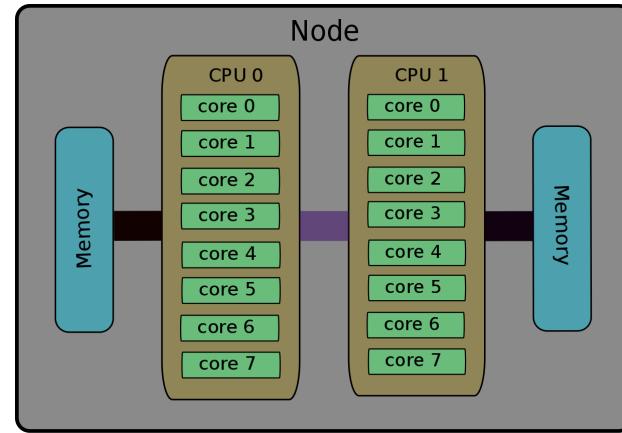
queues	free	busy	offline	jobs	nodes	CPUs
general	12	0	0	0	12	384
condo	29	1	7	0	37	1184
bugs	2	0	0	0	2	64
pcnc	2	0	0	0	2	64
pathogen	0	1	0	1	1	32
tc	7	0	3	0	10	320
gold	2	0	0	0	2	64
fishgen	1	0	0	0	1	32
neuro-hsc	13	0	1	0	14	448
cup-ecs	2	0	0	0	2	64
tid	0	0	1	0	1	32
biocomp	0	0	1	0	1	32
chakra	0	0	1	0	1	32
totals:	70	2	14	1	86	2752

Compute nodes and partitions...

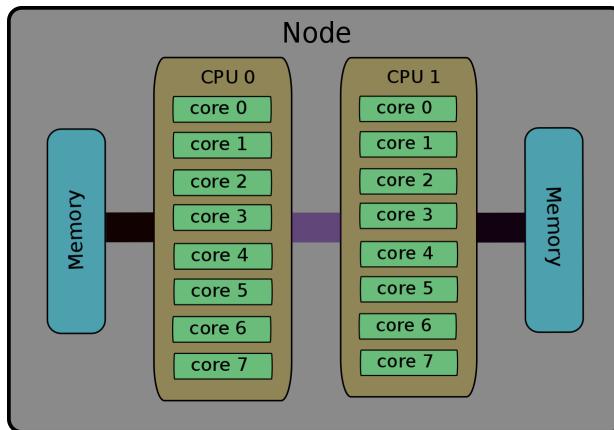
Node



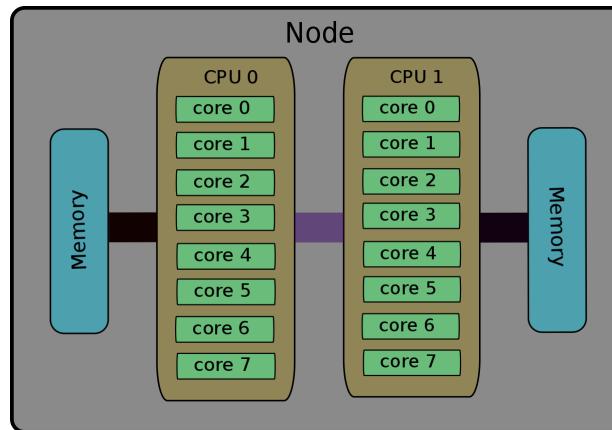
Head Node (wheeler)



Compute Nodes

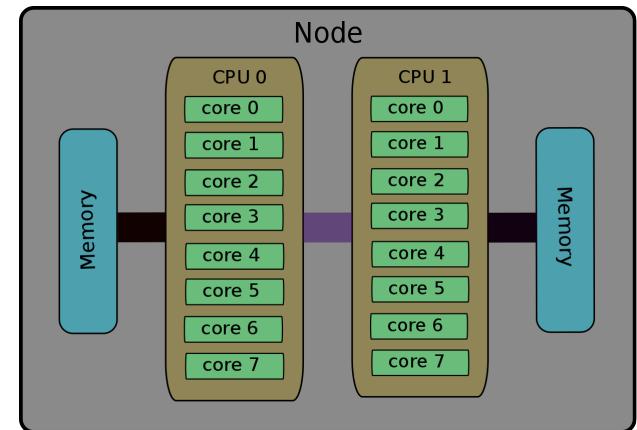


wheeler001



wheeler002

...



wheeler304

```
[vanilla@wheeler ~]$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                8
On-line CPU(s) list:  0-7
Thread(s) per core:   1
Core(s) per socket:   4
Socket(s):             2
NUMA node(s):          2
Vendor ID:             GenuineIntel
CPU family:            6
Model:                 26
Model name:            Intel(R) Xeon(R) CPU X5550 @ 2.67GHz
```

Wheeler has 8 cores per node....

```
mfricke@hopper:~ $ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                64
On-line CPU(s) list:  0-63
Thread(s) per core:   2
Core(s) per socket:   16
Socket(s):             2
NUMA node(s):          2
Vendor ID:             GenuineIntel
CPU family:            6
Model:                 85
Model name:            Intel(R) Xeon(R) Gold 6226R CPU @ 2.90GHz
```

Hopper has 32 real cores (64 virtual cores) per node....

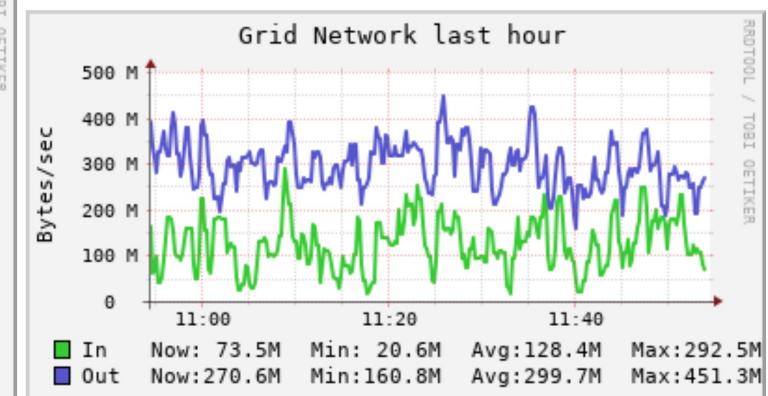
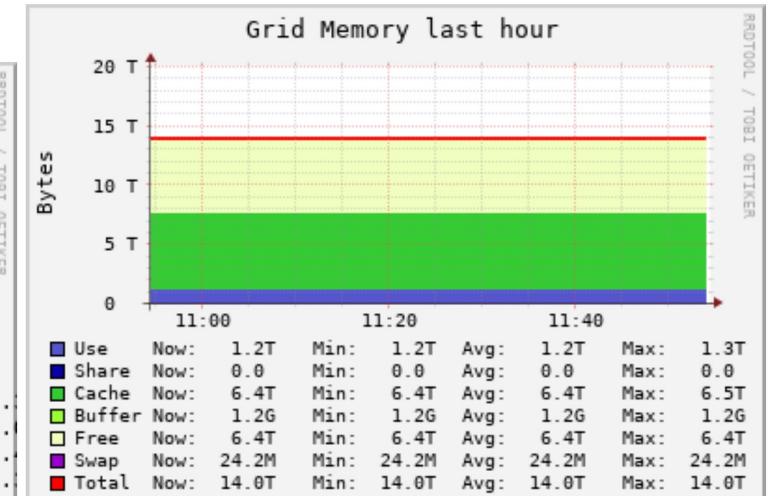
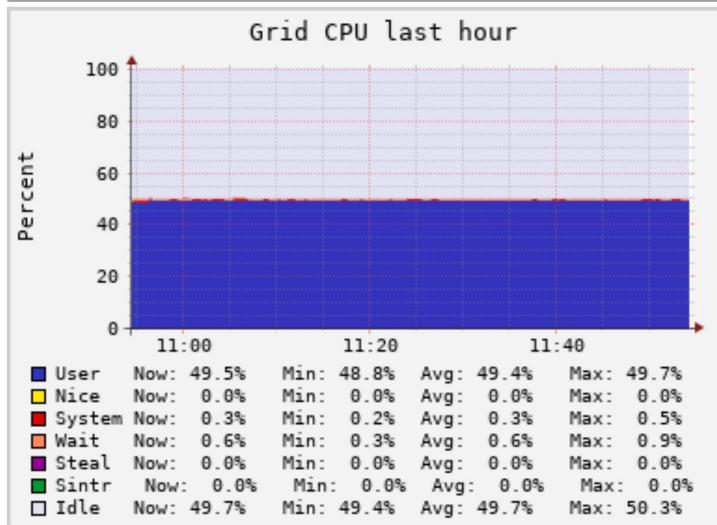
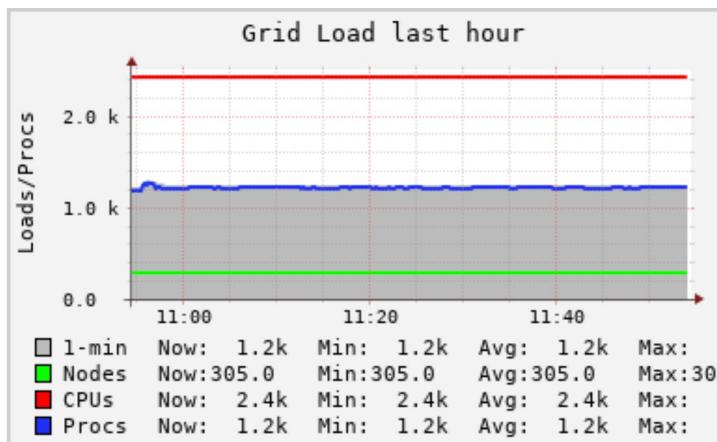
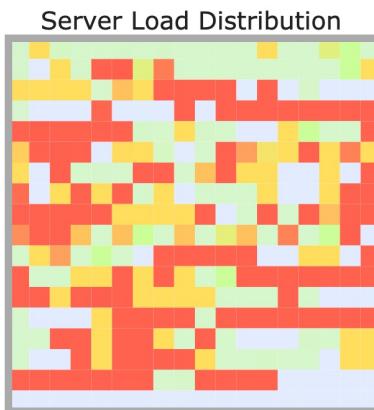
Architecture

Wheeler Cluster Grid > Wheeler Cluster > --Choose a Node

Overview of @ 2022-03-03 18:54

CPUs Total: **2440**
Hosts up: **305**
Hosts down: **0**

Current Load Avg (15, 5, 1m):
51%, 51%, 51%
Avg Utilization (last hour):
51%





Technology, IT etc.

SLURM

means

Simple Linux Utility for Resource Management

ENJOY

Slurm

SODA

IT'S HIGHLY ADDICTIVE!

VOTED #1

SOFT DRINK OF THE 31ST CENTURY!



SLURM MCKENZIE

```
[vanilla@wheeler ~]$ Squeue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
159914	normal	co-mcpdf	nsharma2	CG	2-00:00:26	1	wheeler257
159915	normal	co-mcpdf	nsharma2	CG	2-00:00:26	1	wheeler257
159916	normal	co-mcpdf	nsharma2	CG	2-00:00:26	1	wheeler257
159917	normal	co-mcpdf	nsharma2	CG	2-00:00:26	1	wheeler257
159918	normal	co-mcpdf	nsharma2	CG	2-00:00:26	1	wheeler257
159919	normal	co-mcpdf	nsharma2	CG	2-00:00:26	1	wheeler257
159912	normal	co-mcpdf	nsharma2	CG	2-00:00:28	1	wheeler257
159913	normal	co-mcpdf	nsharma2	CG	2-00:00:28	1	wheeler257
166800_[21-100%10]	normal	Jannat	jannat	PD	0:00	1	(JobArrayTaskLimit)
167067	normal	WINDENER	rubeldas	PD	0:00	36	
(QOSMaxCpuPerUserLimit)							
167068	normal	WINDENER	rubeldas	PD	0:00	24	

Slurm....

```
[vanilla@wheeler ~]$ srun --partition debug --nodes 2 hostname
srun: Account not specified in script or
~/default_slurm_account, using latest project
wheeler302.alliance.unm.edu
You have not been allocated GPUs. To request GPUs, use the -G
option in your submission script.
wheeler301.alliance.unm.edu
[vanilla@wheeler ~]$ srun --partition debug --nodes 2 hostname
```

The srun command...

```
[vanilla@wheeler ~]$ srun --partition debug --ntasks 8 hostname
srun: Account not specified in script or ~/.default_slurm_account, using
latest project
wheeler302.alliance.unm.edu
wheeler302.alliance.unm.edu
wheeler302.alliance.unm.edu
wheeler302.alliance.unm.edu
wheeler302.alliance.unm.edu
wheeler302.alliance.unm.edu
You have not been allocated GPUs. To request GPUs, use the -G option in your
submission script.
wheeler302.alliance.unm.edu
wheeler302.alliance.unm.edu
```

The srun command...

JupyterHub



python™

julia

R

The screenshot shows a Jupyter Notebook interface. On the left, there's a sidebar with the Jupyter logo and a 'Welcome to P...' message. Below it, a yellow box contains a 'WARNING' message: 'Don't rely on this server'. Further down, there's a section titled 'Run some Python code' with instructions on how to run code cells.

```
In [ ]: %matplotlib inline  
import pandas as pd  
import numpy as np  
import matplotlib
```

Exploring the Lorenz System

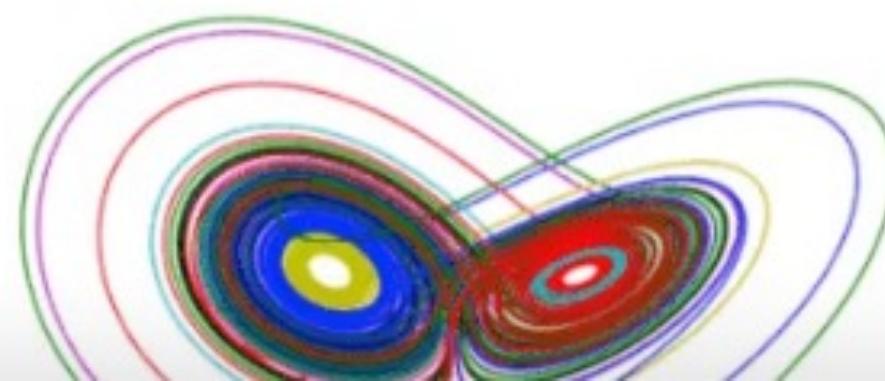
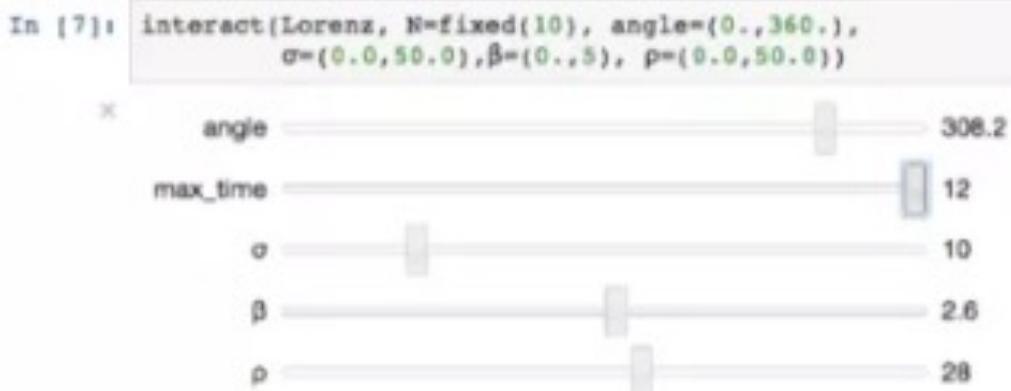
In this Notebook we explore the Lorenz system of differential equations:

$$\dot{x} = \sigma(y - x)$$

$$\dot{y} = \rho x - y - xz$$

$$\dot{z} = -\beta z + xy$$

This is one of the classic systems in non-linear differential equations. It exhibits a range of complex behaviors as the parameters (σ , β , ρ) are varied, including what are known as chaotic solutions. The system was originally developed as a simplified mathematical model for atmospheric convection in 1963.



<https://courses.seas.harvard.edu/climate/eli/Courses/EPS101/index.html>



Description

Global Warming Science: A quantitative introduction to the science of climate change and its consequences, meant to assist students to process issues that often appear in the news and public debates. Topics include: the greenhouse effect and consequences of the rise of greenhouse gasses, including sea level rise, ocean acidification, heat waves, droughts, glacier melting, forest fires, expected changes to hurricanes and more. The scientific basis for each subject will be covered, and every class will involve a hands-on analysis of observations, climate models, and climate feedbacks, using python Jupyter notebooks. Throughout, an ability to critically evaluate observations, predictions and risk will be developed.

Sea level rise in the news

Global Warming Science 101, Sea level, Xiaoting Yang and Eli Tziperman



St Mark's square in Venice on Nov 13, 2019. sea level rise of 1.87 m, highest in more than 50 yrs, flooding over 85% of city

Current sea level rise is ~3.5 mm/year...

a submerged burial ground in Fiji's Togoru village. Five seawalls built to protect it have been knocked down by rising waters.

<https://www.nytimes.com/interactive/2014/03/27/world/climate-rising-seas.html>



<https://www.theguardian.com/environment/2019/dec/10/venice-floods-sea-level-rise-mose-project>

```
vanilla@hopper:~ $ cp -r /projects/shared/workshops/climate_change_jupyterhub ~  
vanilla@hopper:~ $
```

Dependencies

- The example code depends on these libraries:
- cartopy
- gsw
- Using it with jupyterhub requires ipykernel

```
[kfotso@hopper:~$ module avail miniconda
```

```
----- /opt/spack/share/spack/modules/linux-rocky8-skylake_avx512 -----
miniconda3-4.10.3-gcc-8.5.0-yikvdeb
```

```
----- /opt/spack/share/spack/modules/linux-rocky8-cascadelake -----
miniconda3-4.10.3-gcc-11.2.0-tl4rbd6
```

```
----- /opt/spack/share/spack/lmod/linux-rocky8-x86_64/Core -----
miniconda3/4.10.3-yikv
```

If the avail list is too long consider trying:

"module --default avail" or "ml -d av" to just list the default modules.

"module overview" or "ml ov" to display the number of modules for each name.

Use "module spider" to find all possible modules and extensions.

Use "module keyword key1 key2 ..." to search for all possible modules matching any of the "keys".

```
[kfotso@hopper:~$ module load miniconda3
```

```
[kfotso@hopper:~$ module list
```

Currently Loaded Modules:

- 1) miniconda3/4.10.3-yikv

```
Collecting package metadata (current_repodata.json): done
Solving environment: done
```

```
--> WARNING: A newer version of conda exists. <==
    current version: 4.10.3
    latest version: 4.14.0
```

```
Please update conda by running
```

```
$ conda update -n base -c defaults conda
```

```
## Package Plan ##
```

```
environment location: /users/kfotso/.conda/envs/climate_change
```

```
[Proceed ([y]/n)? y
```

```
Preparing transaction: done
Verifying transaction: done
Executing transaction: done
"
```

```
[kfotso@hopper:~$ source activate climate_change  
(climate_change) kfotso@hopper:~$ ]
```

```
(climate_change) kfotso@hopper:~$ conda install cartopy ipykernel  
Collecting package metadata (current_repodata.json): done  
Solving environment: done
```

```
==> WARNING: A newer version of conda exists. <==  
  current version: 4.10.3  
  latest version: 4.14.0
```

Please update conda by running

```
$ conda update -n base -c defaults conda
```

```
## Package Plan ##
```

```
environment location: /users/kfotso/.conda/envs/climate_change
```

```
added / updated specs:  
- cartopy  
- ipykernel
```

The following NEW packages will be INSTALLED:

_libgcc_mutex	conda-forge/linux-64::__libgcc_mutex-0.1-conda_forge
_openmp_mutex	conda-forge/linux-64::__openmp_mutex-4.5-2_gnu

```
(climate_change) kfotso@hopper:~$ conda install -c conda-forge gsw
Collecting package metadata (current_repodata.json): done
Solving environment: done
```

```
==> WARNING: A newer version of conda exists. <==
  current version: 4.10.3
  latest version: 4.14.0
```

Please update conda by running

```
$ conda update -n base -c defaults conda
```

```
## Package Plan ##
```

```
environment location: /users/kfotso/.conda/envs/climate_change
```

```
added / updated specs:
- gsw
```

The following NEW packages will be INSTALLED:

```
gsw           conda-forge/linux-64::gsw-3.4.0-py310hde88566_3
```

Proceed ([y]/n)? █

carc.unm.edu

Xena JupyterHub CARC Helpdesk ParentVUE CARC AIRS fricke.co.uk Fricke Email Spam Manager CARC Systems Systems - Share... CARC Asana My UNM Chrome River UNM Directory UNM OneDrive XDMoD Portal

Paused

Home About Research Education & training News & events New users Systems User support Contact us Donate

Welcome to the Center for Advanced Research Computing

The UNM Center for Advanced Research Computing is the hub of computational research at UNM and one of the largest computing centers in the State of New Mexico. It is an interdisciplinary community that uses computational resources to create new research insights. The goal is to lead and grow the computational research community at UNM.

CARC provides not just the computing resources but also the expertise and support to help the university's researchers. This service is available to faculty, staff, and student researchers free of charge through support from the UNM Office of the Vice President for Research.

Get started

Create a help ticket

Downtime notices

CARC systems information
CARC infrastructure
Resource limits
Storage and backup
Export control
JupyterHub cluster links
System status and downtime
System Usage by Principal Investigator

Students go on strike over greenhouse gas emissions

Hopper Gibbs Taos Wheeler Xena

JupyterHub cluster links

UNM Earth Science Professor to study the health of groundwater aquifers across the state

Explore or exploit: How our brains make choices



Sign in

Username

vanilla

Passwords

• • • • •

Sign in

Server Options

Select a job profile:

- ✓ Debug Queue, 1 hours, 1 core, 4GB RAM
- Debug Queue, 4 hours, 8 cores, 24GB RAM
- General Access Queue, 48 hours, 32 cores, 90GB RAM
- Condo Queue (Preemtable), 48 hours, 32 cores, 90GB RAM
- Condo Queue (Preemtable), 48 hours, 32 cores, 90GB RAM, 1 GPU
- Chakra (restricted), 48 hours, 32 cores, 180GB RAM, 1 GPU
- Geodef (restricted), 1 hours, 1 core, 14GB RAM**
- Geodef (restricted), 48 hours, 32 cores, 500GB RAM





Home

Token

Admin

mfricke

 Logout

Your server is starting up.

You will be redirected automatically when it's ready for you.

Cluster job running... waiting to connect

Event log

jupyterhub

- jupyterhub_hopper_slurm_cmd.err
- jupyterhub_hopper_slurm_cmd.out
- jupyterhub_hopper_batchspawner_28984.log
- jupyterhub_hopper_batchspawner_28982.log
- climate_change_jupyterhub
- Climate_course





Files

Running

Clusters

Select items to perform actions on them.

0



/ Climate_course / Sea-level / code-for-students



sea_level_workshop.ipynb

sea_level_variables.pickle



jupyterhub sea_level_workshop Last Checkpoint: an hour ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Trusted Python [conc]

Interrupt [I, I] ▾

Restart [0, 0]

Restart & Clear Output

Restart & Run All

Reconnect

Shutdown

https://courses.seas...

Global War

Sea Level

Please use the template below to answer your name:

Required installations first:

from a terminal try:
conda install cartopy
pip install gsw

Another option for cartopy: in Anaconda Navigator, click Environments, then change in the pull-down menu "Installed" to "All"; find cartopy installation and click Apply.

If both of these options do not work, try "pip install cartopy" instead of "conda install cartopy" for cartopy.

Change kernel ▾

- Python 3
- Python [conda env:.conda-DarkMatter]
- Python [conda env:.conda-EQTransformer_final]
- Python [conda env:.conda-climate]
- Python [conda env:.conda-climate_change]**
- Python [conda env:.conda-climatechange]
- Python [conda env:.conda-compat_gpu]
- Python [conda env:.conda-eqt]
- Python [conda env:.conda-mpi_numpy]
- Python [conda env:.conda-taos_rgdal_import]
- R [conda env:.conda-taos_rgdal_import]



If both of these options do not work, try "pip install cartopy" instead of "conda install cartopy" for cartopy.



If both of these options do not work, try "pip install cartopy" instead of "conda install cartopy" for cartopy.

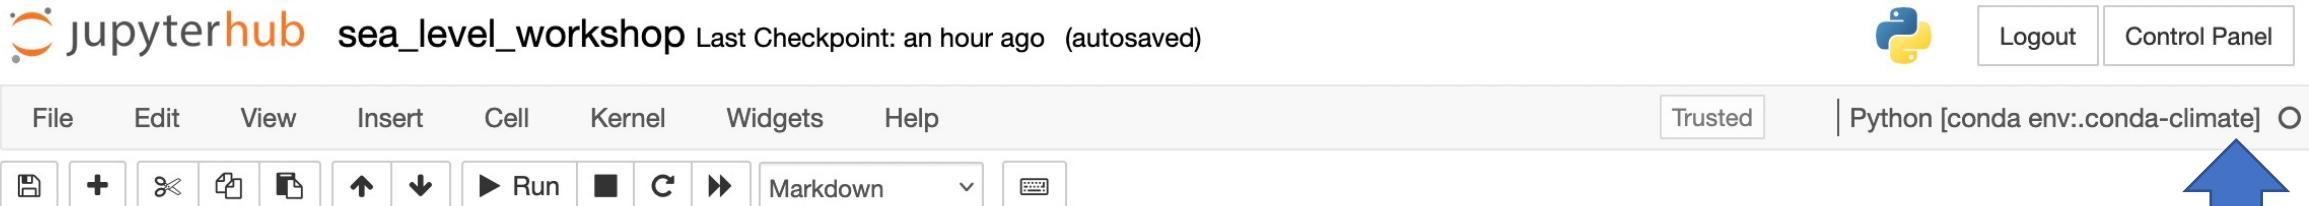
In [1]:

```
# import needed libraries and load data:  
import numpy as np  
import numpy.matlib as matlib  
import pickle  
import matplotlib.pyplot as plt  
import scipy.integrate as integrate  
import gsw  
import cartopy.crs as ccrs  
import cartopy.feature as cfeature  
from cartopy.feature import NaturalEarthFeature  
from cartopy import config  
from matplotlib.colors import BoundaryNorm  
from matplotlib import ticker  
%matplotlib inline
```

```
# Read 2-d sea surface data for historical and rcp85 experiment  
with open('./sea_level_variables.pickle', 'rb') as file:  
    d = pickle.load(file)  
    # print information about each extracted variable:  
    for key in list(d.keys()):  
        print("extracting pickled variable: name=", key, "; size=", d[key].shape)  
        #print("type=", type(d[key]))  
    globals().update(d)
```

```
extracting pickled variable: name= sealevel_lon ; size= (180, 360)  
extracting pickled variable: name= sealevel_lat ; size= (180, 360)  
extracting pickled variable: name= sealevel_historical ; size= (156, 180, 360)  
extracting pickled variable: name= sealevel_historical_years ; size= (156,)  
extracting pickled variable: name= sealevel_rcp85 ; size= (95, 180, 360)
```

Shut down your jupyterhub server



Global Warming Science

<https://courses.seas.harvard.edu/climate/eli/Courses/EPS101/>

Sea Level

Please use the template below to answer the *complete workshop questions from the course notes*.

your name:



Home

Token

Admin



Stop My Server

My Server

CARC Resources

- Tutorial Videos
- Written Tutorials



QuickBytes

Quickbytes are tutorials designed to help CARC users.

- Linux-Intro
- Running jobs
 - Logging in
 - SSH keys and Config file
 - Transferring data
 - PBS/TORQUE
 - Sample PBS script
 - Submitting jobs
 - Check running jobs
 - Managing modules
 - Intro to Slurm
 - Converting PBS to Slurm
 - Intro to Slurm accounting at CARC

A screenshot of a YouTube channel page for 'UNMCARC'. The channel name is 'QuickBytes'. It has 16 videos and 2,870 views, last updated on Jan 5, 2022. A note on the page states: 'Short Tutorials on CARC Systems - DUE TO THE RECENT WHEELER UPGRADE THESE VIDEOS WILL BE REPLACED SOON.' The channel has a logo for 'UNM CENTER FOR ADVANCED RESEARCH COMPUTING'. There are 6 video thumbnails listed:

1. Intro to the UNM Center for Advanced Research Computing (5:53)
2. Projects and Accounts (8:47)
3. Logging into CARC Systems (12:48)
4. Storage Systems (13:11)
5. Transferring data (21:56)
6. Environment Modules (6:04)

A red 'SUBSCRIBE' button is located below the video list.

Getting Help

help@carc.unm.edu

Office hours