

Intermediate Level Introduction to Computing at CARC

1 hour version with SimCov

Matthew Fricke

Version 0.1



Goals

- 1) SLURM scheduler literacy
- 2) Example SimCov
- We wont cover file transfer, storage systems, module system, conda, PBS. (These are all covered in depth in the video tutorials)

Logging into Hopper



First login to the Linux **workstation** in front of you. Your CARC username is on the sign in sheet.

If you have logged in before use your **existing password**

Otherwise, your initial password is **Welcome2CARC**

This is an “important step” so don’t let me move on until you have logged in

Logging into Hopper



```
ssh vanilla@hopper.alliance.unm.edu
```

Should prompt you for a password...

Don't let me move on until you are able to login.

Replace vanilla with your name (unless your last name is Ice)

Logging into Hopper

Welcome to Hopper

Be sure to review the "Acceptable Use" guidelines posted on the CARC website.

For assistance using this system email help@carc.unm.edu.

Tutorial videos can be accessed through the CARC website: Go to <http://carc.unm.edu>, select the "New Users" menu and then click "Introduction to Computing at CARC".

Warning: By default home directories are world readable. Use the chmod command to restrict access.

Don't forget to acknowledge CARC in publications, dissertations, theses and presentations that use CARC computational resources:

"We would like to thank the UNM Center for Advanced Research Computing, supported in part by the National Science Foundation, for providing the research computing resources used in this work."

Please send citations to publications@carc.unm.edu.

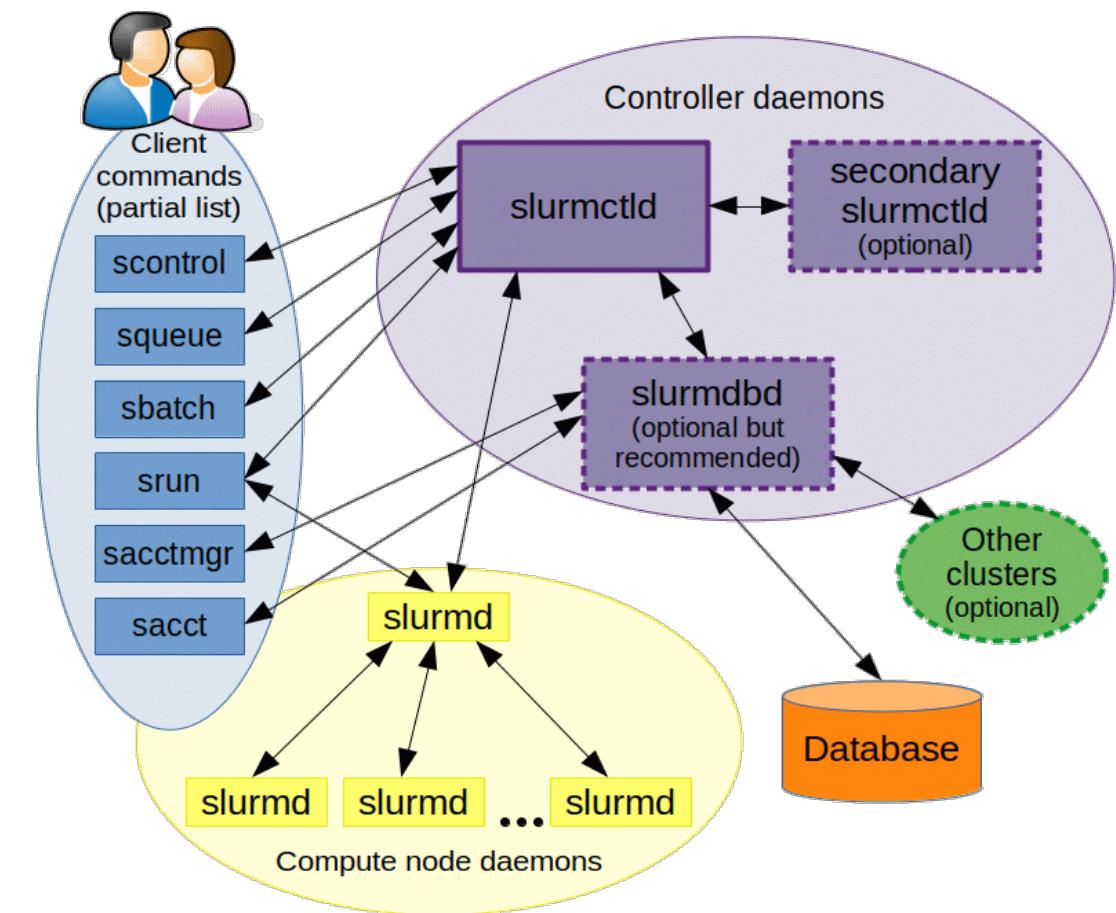
There are three types of slurm partitions on Hopper:

- 1) General - this partition is accessible by all CARC users.
- 2) Condo - preemptable scavenger queue available to all condo users. Your job must use checkpointing to use this queue or you will lose any work you have done if it is preempted by the partition's owner.
- 3) Named partitions - these partitions are available to condo users working under the grant/lab/center that purchased the associated hardware.

Type "qgrok" to get the status of the partitions.

Last login: Wed Jul 27 17:46:13 2022 from 129.24.246.68
mfricke@hopper:~ \$

Simple Linux Utility for Resource Management



ENJOY

Slurm

SODA

IT'S HIGHLY ADDICTIVE!

VOTED #1 SOFT DRINK OF THE 31ST CENTURY!



SELENE MCKENRE

```
[vanilla@hopper ~]$ qgrok
```

queues	free	busy	offline	jobs	nodes	CPUs	GPUs	CPUs/node	GPUs/node	Memory/node	time_limit	CPU_limit
general	4	6	0	4	10	320	0	32	0	93G	2-00:00:00	64
debug	2	0	0	0	2	64	0	32	0	93G	4:00:00	8
condo	22	25	4	7	51	1632	28	32	2	93G-1.5T	2-00:00:00	192
bugs	2	0	0	0	2	64	0	32	0	93G	7-00:00:00	

```
vanilla@hopper:~ $ qgrok
```

queues	free	busy	offline	jobs	nodes	CPUs	GPUs
general	1	9	0	7	10	320	0
debug	2	0	0	0	2	64	0
condo	18	19	1	7	38	1216	8
bugs	0	2	0	0	2	64	0
pcnc	1	1					
pathogen	1	0					
tc	5	5					
gold	2	0					
fishgen	0	1					
neuro-hsc	8	6					
cup-ecs	0	2					
tid	0	1					
biocomp	0	1					
chakra	1	0					
pna	0	0					
totals:	19	28					

Open partitions for use by
everyone with a CARC account.

Purchased by the Office for the
Vice President for Research.

vanilla@hopper:~ \$ qgrok

queues	free	busy	offline	jobs	nodes	CPUs	GPUs
general	1	9	0	7	10	320	0
debug	2	0	0	0	2	64	0
condo	18	19	1	7	38	1216	8
bugs	0	2	0	0	2	64	0
pcnc	1	1	0	0	2	64	0
pathogen	1	0	0	0	1	32	0
tc	5	5	0	3	10	320	0
gold	2	0	0	0	2	64	0
fishgen	0	1	0	0	1	32	0
neuro-hsc	8	6	0	0	14	448	0
cup-ecs	0	2	0	2	2	64	4
tid	0	1	0	0	1	32	2
biocomp	0	1	0	0	1	32	1
chakra	1	0	0	0	1	32	1
pna	0	0	1	0	1	32	0
totals:	19	28	1	14	48	1536	8

```
vanilla@hopper:~ $ qgrok
```

queues	free	busy	offline	jobs	nodes	CPUs	GPUs
general	1	9	0	7	10	320	0
debug	2	0	0	0	2	64	0
condo	18	19	1	7	38	1216	8
bugs	0	2	0	0	2	64	0
pcnc	1	1					
pathogen	1	0					
tc	5	5					
gold	2	0					
fishgen	0	1					
neuro-hsc	8	6					
cup-ecs	0	2					
tid	0	1					
biocomp	0	1					
chakra	1	0					
pna	0	0					
totals:	19	28					

Private partitions

- Reserved for use by the purchaser.
- Request access by emailing support@carc.unm.edu and CC the partition owner.

```
mfricke@hopper:~ $ qgrok
```

queues	free	busy	offline	jobs	nodes	CPUs	GPUs
general	1	9	0	7	10	320	0
debug	2	0	0	0	2	64	0
condo	18	19	1	7	38	1216	8
bugs	0	2	0	0	2	64	0
pcnc	1	1					
pathogen	1	0					
tc	5	5					
gold	2	0					
fishgen	0	1					
neuro-hsc	8	6					
cup-ecs	0	2					
tid	0	1					
biocomp	0	1					
chakra	1	0					
pna	0	0					
totals:	19	28					

Condo “scavenger” partition

- Allows you to use compute nodes purchased by another group that are currently idle.
- May be interrupted at any time if the owners start to use it.

```
[vanilla@hopper ~]$ quotas
```

```
Home Directory (/users/vanilla):
```

```
quota: Cannot resolve mountpoint path /root/.spack: Permission denied
```

```
Disk quotas for user vanilla (uid 659):
```

Filesystem	space	quota	limit	grace	files	quota	limit	grace
chama:/home/homes	1527M	100G	200G		14913	4295m	4295m	

```
Centerwide user scratch (/carc/scratch/users/mfricke)
```

```
Quota information for storage pool Default (ID: 1):
```

user/group	size	chunk	files	
name	used	hard	used	hard
mfricke	592.71 GiB	1024.00 GiB	32784	unlimited

```
Centerwide scratch quota for project mfricke2016174 (/carc/scratch/projects/mfricke2016174)
```

```
Quota information for storage pool Default (ID: 1):
```

user/group	size	chunk	files	
name	used	hard	used	hard
mfricke2016174	190.97 GiB	1024.00 GiB	23704	unlimited

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug        up    4:00:00      2  idle  hopper[011-012]
```

sinfo reports information about
partitions

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug        up    4:00:00      2  idle  hopper[011-012]
```

The debug queues are intended
for testing your programs.

And for interactive jobs.

```
[vanilla@hopper ~]$ sinfo --partition debug
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug        up    4:00:00      2  idle hopper[011-012]
```



Name

```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST  
debug        up    4:00:00      2  idle  hopper[011-012]
```



You can run a “job” for up to 4 hrs.

```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST  
debug        up    4:00:00      2  idle  hopper[011-012]
```



There are two nodes in this partition.

```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST  
debug        up    4:00:00      2  idle  hopper[011-012]
```



The names of the nodes in the partition

```
[vanilla@hopper ~]$ sinfo --partition debug  
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST  
debug        up    4:00:00      2  idle  hopper[011-012]
```



The names of the nodes in the partition

```
[vanilla@hopper ~]$ sinfo --partition general
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
general*      up   2-00:00:00      9  alloc  hopper[001-009]
general*      up   2-00:00:00      1  idle   hopper010
```



Hopper001 through 009 are running jobs.

Hopper010 is waiting to be used.

```
[vanilla@hopper ~] $ hostname  
hopper  
[vanilla@hopper ~] $
```



Running on the Head Node.
The head node's name is “hopper”.

```
[vanilla@hopper ~]$ hostname  
hopper
```

```
[vanilla@hopper ~]$ man hostname
```

```
[vanilla@hopper ~]$ hostname  
hopper
```

```
[vanilla@hopper ~]$ man hostname  
(‘q’ to quit)
```

```
[vanilla@hopper ~]$ man man  
(‘q’ to quit)
```

```
[vanilla@hopper ~]$ man sinfo
```

sinfo(1)
Commands

Slurm
sinfo(1)

NAME

sinfo - View information about Slurm nodes and partitions.

SYNOPSIS

`sinfo [OPTIONS...]`

DESCRIPTION

sinfo is used to view partition and node information for a system running Slurm

OPTIONS

`-a, --all`

Display information about all partitions. This causes information to be displayed about partitions that are configured as hidden and partitions that are unavailable to the user's group.

```
[vanilla@hopper ~]$ sinfo --all
```

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
general*	up	2-00:00:00	9	alloc	hopper[001-009]
general*	up	2-00:00:00	1	idle	hopper010
debug	up	4:00:00	2	idle	hopper[011-012]
condo	up	2-00:00:00	1	down*	hopper045
condo	up	2-00:00:00	3	mix	hopper[018-020]
condo	up	2-00:00:00	16	alloc	hopper[013-015,028-036,049-052]
condo	up	2-00:00:00	18	idle	hopper[016-017,021-027,037-044,053]
bugs	up	7-00:00:00	2	alloc	hopper[013-014]
pcnc	up	7-00:00:00	1	alloc	hopper015
pcnc	up	7-00:00:00	1	idle	hopper016
pathogen	up	7-00:00:00	1	idle	hopper017
tc	up	7-00:00:00	3	mix	hopper[018-020]
tc	up	7-00:00:00	2	alloc	hopper[029-030]
tc	up	7-00:00:00	5	idle	hopper[021-025]
gold	up	7-00:00:00	2	idle	hopper[026-027]
fishgen	up	7-00:00:00	1	alloc	hopper028
neuro-hsc	up	7-00:00:00	6	alloc	hopper[031-036]
neuro-hsc	up	7-00:00:00	8	idle	hopper[037-044]
cup-ecs	up	7-00:00:00	2	alloc	hopper[049-050]
tid	up	7-00:00:00	1	alloc	hopper051
biocomp	up	7-00:00:00	1	alloc	hopper052
chakra	up	7-00:00:00	1	idle	hopper053
pna	up	7-00:00:00	1	down*	hopper045

```
[vanilla@hopper ~]$ srun --partition debug hostname
```



Tell slurm to run a program
on a compute node...

```
[vanilla@hopper ~]$ srun --partition debug hostname
```



Run the program on a
compute node in the
debug partition.

```
[vanilla@hopper ~]$ srun --partition debug hostname
```



The program
to run.

```
[vanilla@hopper ~]$ srun --partition debug hostname
srun: Account not specified in script or
~/.default_slurm_account, using latest project
You have not been allocated GPUs. To request GPUs,
use the -G option in your submission script.
hopper011
```

```
[vanilla@hopper ~] $ squeue
```

```
[vanilla@hopper ~]$ squeue
```

JOBID	PARTITION	NAME	USER	ST	TIM	2 (QOSMaxCpuPerUserLimit)
4314	general	PRE	erowland	PD	0:00	2 (QOSMaxCpuPerUserLimit)
4315	general	PRE	erowland	PD	0:00	2 (QOSMaxCpuPerUserLimit)
4317	general	PRE	erowland	PD	0:00	2 (QOSMaxCpuPerUserLimit)
4318	general	PRE	erowland	PD	0:00	2 (QOSMaxCpuPerUserLimit)

PD means programs
that are waiting their
turn.

Shows you what the slurm
scheduler is doing right now.

Here we can see that user
'erowland' has a lot of programs
waiting to run.

```
[vanilla@hopper ~]$ squeue
```

The reason these jobs are not running is that ‘erowland’ is already using the maximum number of CPUs they are allowed.

```
[vanilla@hopper ~]$ squeue -t R --all
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
4405	condo	2ndMA	mfricke	R	1-07:48:30	6	hopper[031-036]
5208	condo	NN	kgu	R	5:48:49	1	hopper015
5210	condo	NN	kgu	R	6:30:13	1	hopper014
5209	condo	NN	kgu	R	6:31:13	1	hopper013
5206	condo	NN	kgu	R	6:32:13	1	hopper051
5207	condo	NN	kgu	R	6:32:13	1	hopper052
5205	condo	NN	kgu	R	6:32:43	1	hopper028
4595	cup-ecs	golConfi	aalasand	R	2-06:51:59	1	hopper050
4594	cup-ecs	golConfi	aalasand	R	2-06:52:03	1	hopper049
5120	general	jupyterh	jacobm	R	11:45:47	1	hopper007
4313	general	PRE	erowland	R	1:17:29	2	hopper[003-004]
5111	general	1stMA	mfricke	R	11:15:28	2	hopper[005-006]
5025	general	c2n	jxzuo	R	1:50	1	hopper001
5024	general	c2n	jxzuo	R	31:28	1	hopper002
5203	general	NN	kgu	R	6:37:50	1	hopper009
5201	general	NN	kgu	R	6:38:14	1	hopper008
4390	tc	UCsTpCyd	lepluart	R	2-15:18:18	3	hopper[018-020]
5198	tc	NN	kgu	R	6:40:19	1	hopper030
5196	tc	NN	kgu	R	6:40:31	1	hopper029

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 hostname
srun: Account not specified in script or
~/.default_slurm_account, using latest project
You have not been allocated GPUs. To request GPUs, use the -G
option in your submission script.
hopper011
hopper011
```

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 hostname
srun: Account not specified in script or
~/.default_slurm_account, using latest project
You have not been allocated GPUs. To request GPUs, use the -G
option in your submission script.
hopper011
hopper011
```

You ran two **copies** of your program.

ntasks is the number of copies to run.

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 8 hostname
srun: Account not specified in script or
~/.default_slurm_account, using latest project
hopper011
hopper011
hopper011
hopper011
You have not been allocated GPUs. To request GPUs, use the -G
option in your submission script.
hopper011
hopper011
hopper011
hopper011
hopper011
```

You ran eight **copies** of your program.

ntasks is the number of copies to run.

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 8 hostname
srun: Account not specified in script or
~/.default_slurm_account, using latest project
hopper011
hopper011
hopper011
hopper011
You have not been allocated GPUs. To request GPUs, use the -G
option in your submission script.
hopper011
hopper011
hopper011
hopper011
hopper011
```

By default, each task (copy of your program) is allowed to use one CPU.

Many programs are able to use more than one CPU at a time.

```
[vanilla@hopper ~]$ srun --partition debug --ntasks 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using latest project  
You have not been allocated GPUs. To request GPUs, use the -G option in your submission  
script.  
hopper011  
hopper011
```

Here we are telling SLURM to run 2 copies of our program and let each copy of our program use 2 CPUs.

```
[vanilla@hopper ~]$ srun --partition debug --nodes 2 -ntasks-per-node 4 hostname  
srun: Account not specified in script or ~/.default_slurm_account, using  
latest project  
hopper012  
You have not been allocated GPUs. To request GPUs, use the -G option in  
your submission script.  
hopper012  
hopper011  
hopper011  
hopper012  
hopper012  
hopper011  
hopper011
```

Here we are telling SLURM to run 4 copies of our program on 2 different compute nodes.

This is useful when our programs need a bigger share of the compute node.

```
[vanilla@hopper ~]$ srun --partition debug --nodes 2  
--ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or  
~/.default_slurm_account, using latest project  
hopper011  
You have not been allocated GPUs. To request GPUs, use  
the -G option in your submission script.  
hopper011  
hopper012  
hopper012
```

And we can combine all three.

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2  
hostname
```

```
srun: Account not specified in script or  
~/.default_slurm
```

```
hopper012
```

```
hopper012
```

```
You have not be
```

```
the -G option i
```

```
hopper011
```

```
Hopper011
```

**And we can specify how much
memory we want.**

**--mem 4G means give me 4
gigabytes of memory per node.**

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2  
hostname  
srun: Account not specified in script or  
~/.default_slurm  
hopper012  
hopper012  
You have not been assigned to a partition.  
the -G option is required.  
hopper011  
Hopper011
```

Why does all this matter?

The purpose of SLURM is to provide you the hardware your programs need.

So you have to understand what those requirements are really well.

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or  
~/.default_slurm_account using latest project  
hopper012  
hopper012  
You have not been assigned to a project.  
the -G option is required.  
hopper011  
Hopper011
```

- 1) Can my program use multiple CPUs?**
- 2) How much memory does my program need?**
- 3) Can my program use multiple compute nodes (MPI*, GNU Parallel*)?**
- 4) Can my program use GPUs?**

```
[vanilla@hopper ~]$ srun --partition debug --mem 4G  
--nodes 2 --ntasks-per-node 2 --cpus-per-task 2 hostname  
srun: Account not specified in script or  
~/.default_slurm_account using latest project  
hopper012  
hopper012  
You have not been  
the -G option in  
hopper011  
Hopper011
```

This command is getting pretty long.

We can use **shell scripts** to automate
all this in **batch mode**.

Interactive vs Batch Mode

Interactive Mode

- Everything so far has been interactive. You request hardware, run your program, and get the output on your screen right away.

Batch Mode

- Most programs at an HPC center are run in “batch” mode.
- Batch mode means we write a shell script that the SLURM scheduler runs for us. The script requests hardware just like we did with salloc and then runs the commands in the script.
- Whatever would have been written to the screen is saved to a file instead.

```
[vanilla@hopper ~]$ git clone https://lobogit.unm.edu/CARC/workshops.git  
Cloning into 'workshops'...  
remote: Enumerating objects: 132, done.  
remote: Counting objects: 100% (75/75), done.  
remote: Compressing objects: 100% (43/43), done.  
remote: Total 132 (delta 33), reused 74 (delta 32), pack-reused 57  
Receiving objects: 100% (132/132), 57.58 KiB | 3.60 MiB/s, done.  
Resolving deltas: 100% (51/51), done.
```

Rather than make you write shell scripts lets just download some we wrote for this workshop...

```
[vanilla@hopper ~]$ tree workshops
```

```
workshops/
├── intro_workshop
│   ├── code
│   │   ├── calcPiMPI.py
│   │   ├── calcPiSerial.py
│   │   └── vecadd
│   │       ├── Makefile
│   │       ├── vecadd_gpu.cu
│   │       ├── vecadd_mpi_cpu
│   │       ├── vecadd_mpi_cpu.c
│   │       ├── vecaddmpi_cpu.sh
│   │       └── vecadd_mpi_gpu.c
│
└── data
    ├── H2O.gjf
    └── step_sizes.txt
└── slurm
    ├── calc_pi_array.sh
    ├── calc_pi_mpi.sh
    ├── calc_pi_parallel.sh
    ├── calc_pi_serial.sh
    ├── gaussian.sh
    ├── hostname_mpi.sh
    ├── vecadd_hopper.sh
    ├── vecadd_xena.sh
    ├── workshop_example2.sh
    ├── workshop_example3.sh
    └── workshop_example.sh

```

Run tree to see how the workshops directories are organized...

```
[vanilla@hopper ~]$ tree workshops
```

```
workshops/
├── intro_workshop
│   ├── code
│   │   ├── calcPiMPI.py
│   │   ├── calcPiSerial.py
│   │   └── vecadd
│   │       ├── Makefile
│   │       ├── vecadd_gpu.cu
│   │       ├── vecadd_mpi_cpu
│   │       ├── vecadd_mpi_cpu.c
│   │       ├── vecaddmpi_cpu.sh
│   │       └── vecadd_mpi_gpu.c
│
├── data
│   ├── H2O.gjf
│   └── step_sizes.txt
└── slurm
    ├── calc_pi_array.sh
    ├── calc_pi_mpi.sh
    ├── calc_pi_parallel.sh
    ├── calc_pi_serial.sh
    ├── gaussian.sh
    ├── hostname_mpi.sh
    ├── vecadd_hopper.sh
    ├── vecadd_xena.sh
    ├── workshop_example2.sh
    ├── workshop_example3.sh
    └── workshop_example.sh

```

```
README.md
```

Run **tree** to see how the workshops directories are organized...

The workshop files are divided into “code”, “slurm”, and “data” directories.

```
[vanilla@hopper intro_workshop]$ pwd  
/users/vanilla/workshops/intro_workshop  
[vanilla@hopper intro_workshop]$ cat slurm/workshop_example.sh  
#!/bin/bash  
#SBATCH --partition debug  
#SBATCH --ntasks 4  
#SBATCH --time 00:05:00  
#SBATCH --job-name ws_example  
#SBATCH --mail-user your_username@unm.edu  
#SBATCH --mail-type ALL
```

hostname

Let's take a look at the **workshop_example.sh** script in the slurm directory...

```
[vanilla@hopper intro_workshop]$ sbatch slurm/workshop_example.sh  
sbatch: Account not specified in script or ~/.default_slurm_account,  
using latest project  
Submitted batch job 5252  
[vanilla@hopper intro_workshop]$
```

We **submit** our slurm
shell script with the
sbatch command.

```
[vanilla@hopper intro_workshop]$ sbatch slurm/workshop_example.sh  
sbatch: Account not specified in script or ~/.default_slurm_account,  
using latest project  
Submitted batch job 5252  
[vanilla@hopper intro_workshop]$
```

Notice that the only output we get is a job id.

This indicates that the script was successfully sent to the scheduler.

The commands in the script will run as soon as the hardware requested is available.

We submit our slurm shell script with the sbatch command.

Workflow

Head Node

User 1

Program A

Script A

User 2

Program B

Script B

Compute Node 01

Compute Node 02

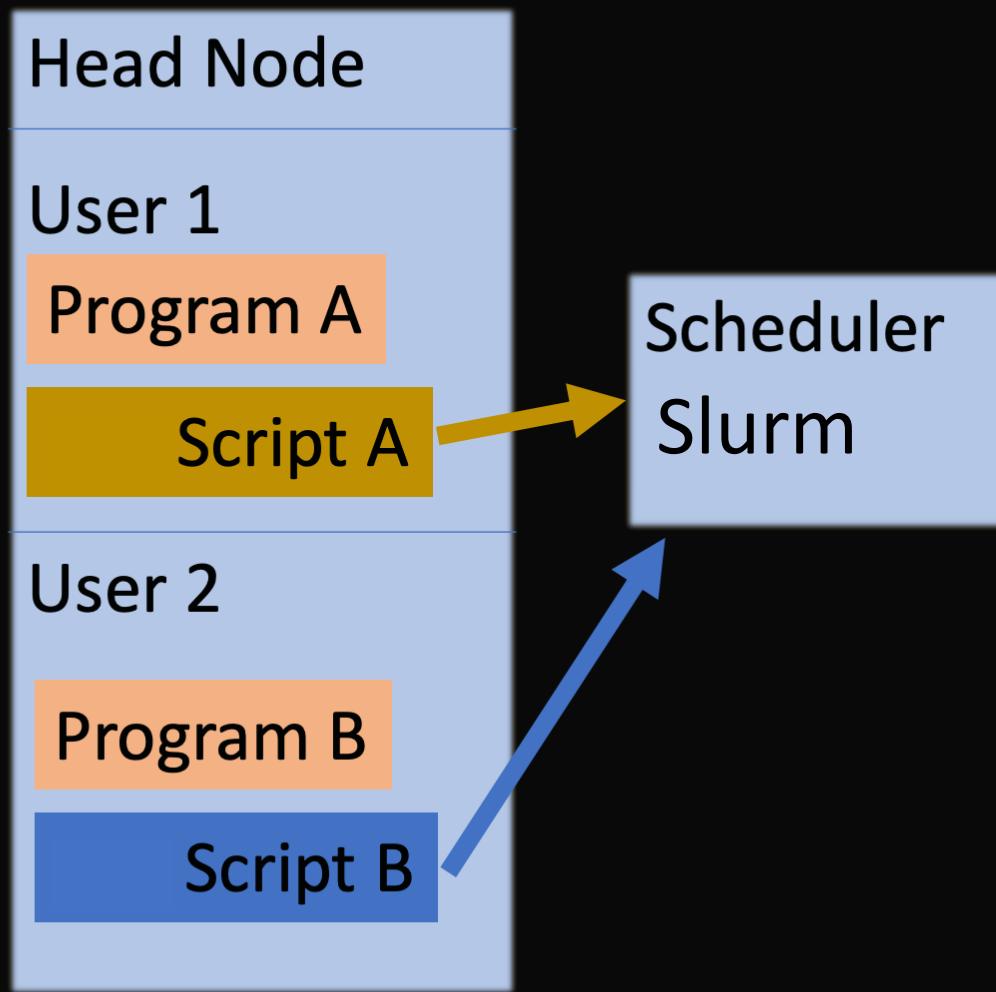
Compute Node 03

Compute Node 04

Compute Node 05

Shared filesystems – All nodes can access the same programs and write output

Workflow



Compute Node 01

Compute Node 02

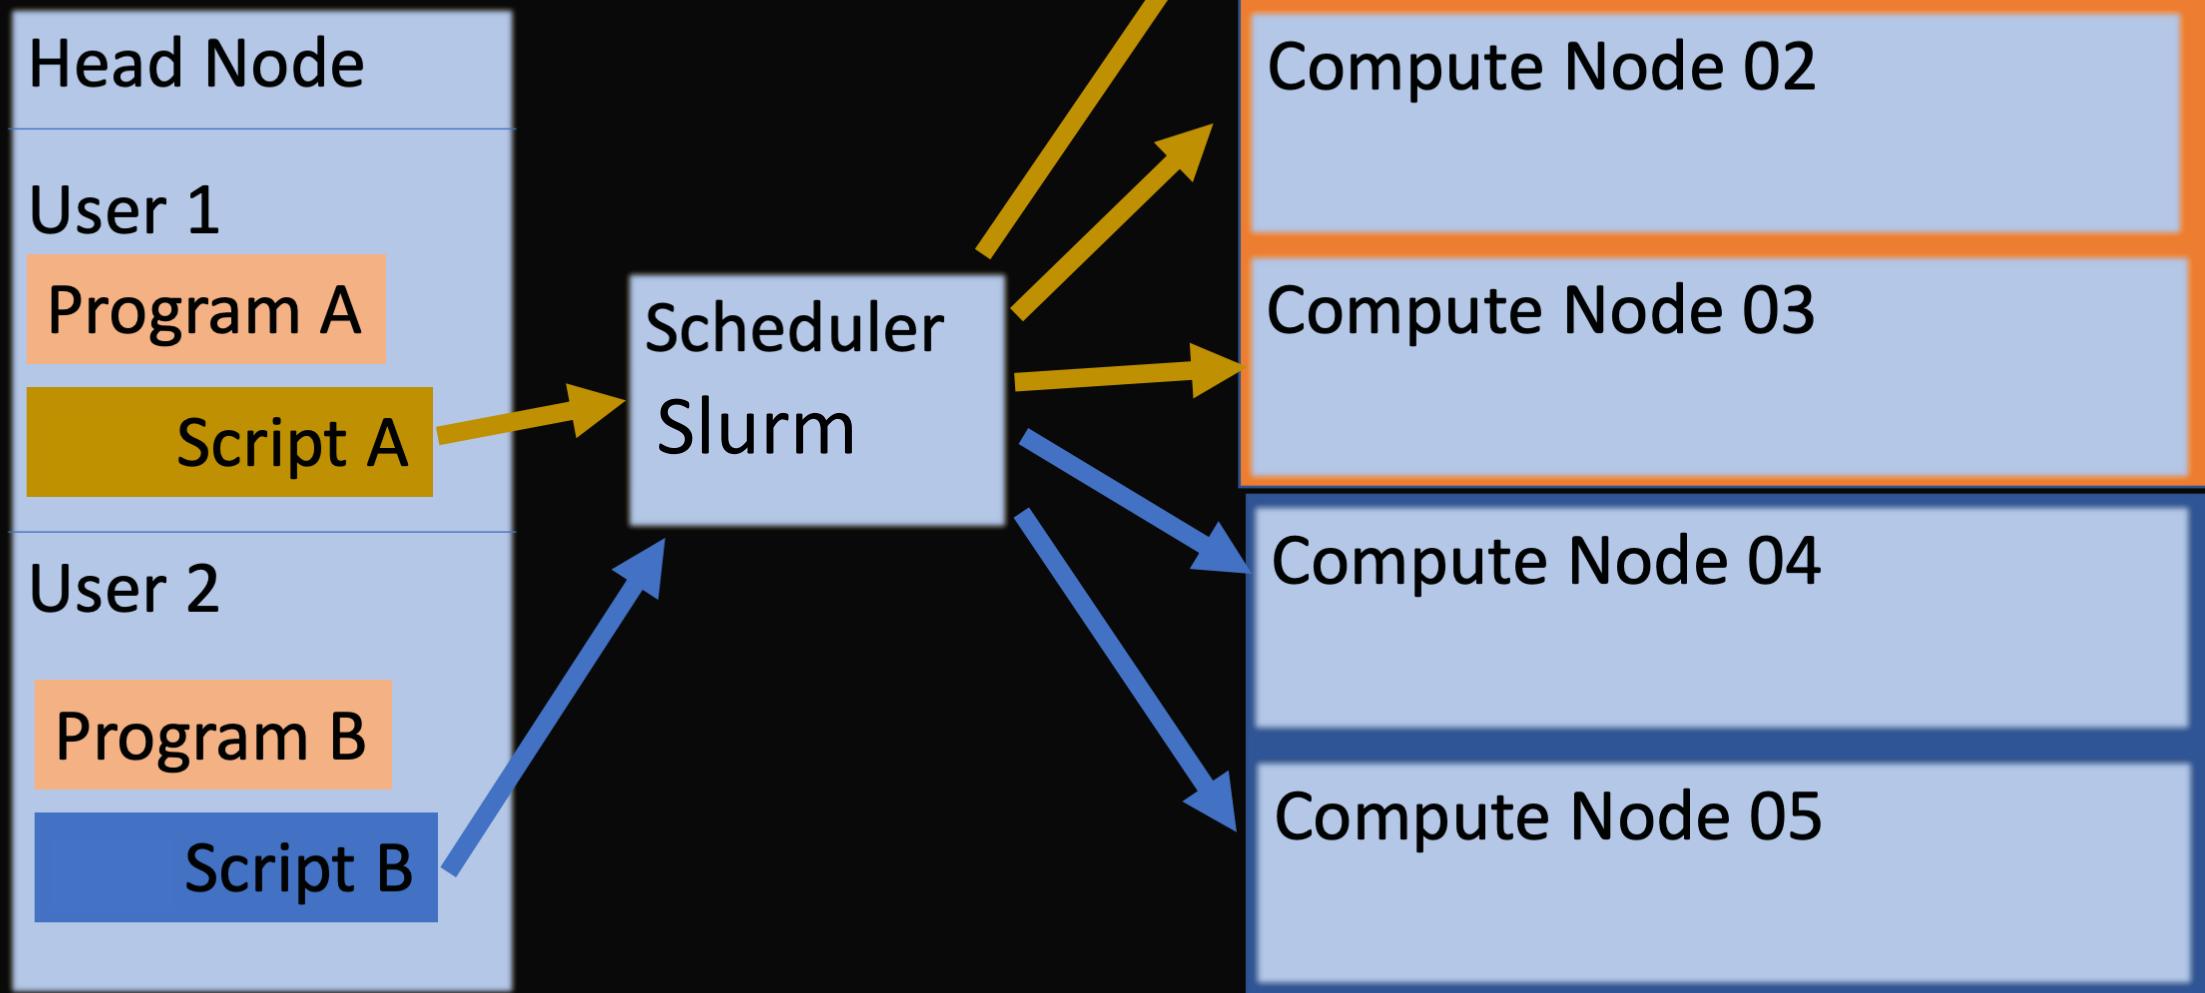
Compute Node 03

Compute Node 04

Compute Node 05

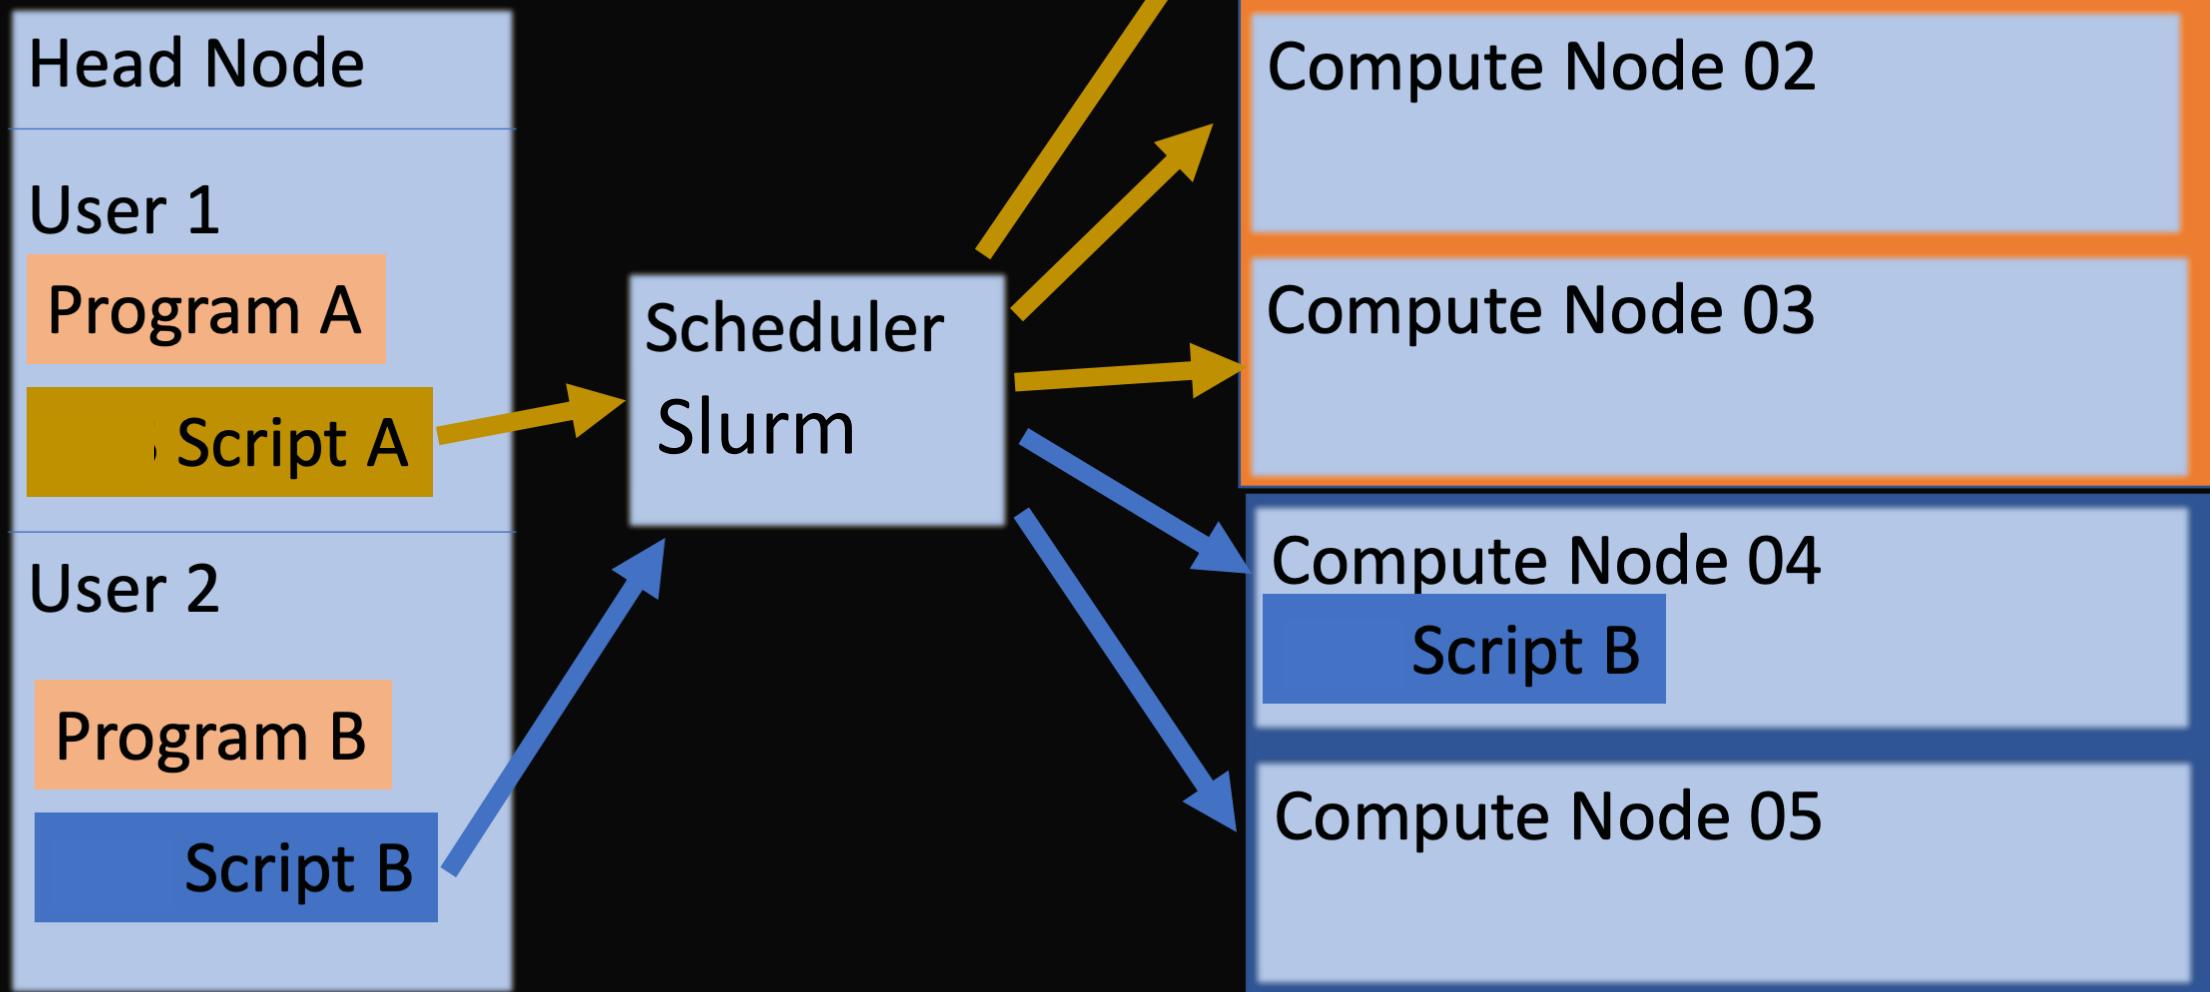
Shared filesystems – All nodes can access the same programs and write output

Workflow



Shared filesystems – All nodes can access the same programs and write output

Workflow



Shared filesystems – All nodes can access the same programs and write output

Workflow

We need something in the script to run the program on all the nodes. E.g. srun.

Head Node

User 1

Program A

Script A

User 2

Program B

Script B

Scheduler
Slurm

Compute Node 01

Script A

Program A

Compute Node 02

Program A

Program A

Compute Node 03

Program A

Program A

Compute Node 04

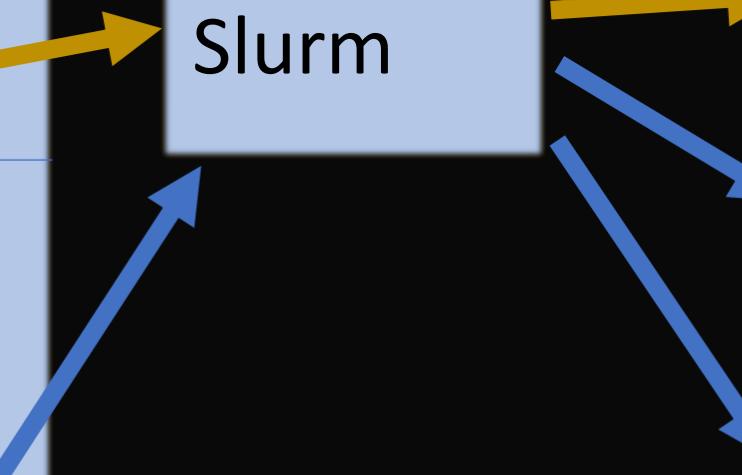
Script B

Program B

Compute Node 05

Program B

Program B



Shared filesystems – All nodes can access the same programs and write output

```
[vanilla@hopper intro_workshop]$ ls  
code  data  pbs  slurm  slurm-5252.out
```

The **hostname** command is very fast so everyone's job should finish in a few seconds.

When it is finished you will have a new file named slurm-{your job id}.out.



```
[vanilla@hopper intro_workshop]$ ls  
code  data  pbs  slurm  slurm-5252.out
```



When it is finished you will have a new file named slurm-{your job id}.out.

```
[vanilla@hopper intro_workshop]$ cat slurm-5252.out  
hopper011
```

```
[vanilla@hopper intro_workshop]$ ls  
code  data  pbs  slurm  slurm-5252.out
```



When it is finished you will have a new file named slurm-{your job id}.out.

```
[vanilla@hopper intro_workshop]$ cat slurm-5252.out  
hopper011
```

Why did it only run the program once instead of 4 times?

```
[vanilla@hopper intro_workshop]$ sbatch slurm/workshop_example1.sh
sbatch: Account not specified in script or ~/.default_slurm_account,
using latest project
Submitted batch job 5252
[vanilla@hopper intro_workshop]$
```

Take a look at the
output file.

```
[vanilla@hopper intro_workshop]$ module load miniconda3
```

```
[vanilla@hopper intro_workshop]$ conda create -n numpy numpy
```

Wait a while – introduce yourselves to your neighbor...

Conda allows you to install software into your home directory. In this case we need the numerical python libraries for calcPiSerial.py

**Let's experiment with a
program that does
slightly more than
print the hostname.**

```
[vanilla@hopper intro_workshop]$source activate numpy  
[vanilla@hopper intro_workshop]$srun --partition debug python  
code/calcPiSerial.py 10  
srun: Using account 2016199 from ~/.default_slurm_account  
You have not been allocated GPUs. To request GPUs, use the -G  
option in your submission script.  
Pi = 3.14242598500109870940, (Diff=0.00083333141130559341)  
(calculated in 0.000005 secs with 10 steps)
```

**Activate the numpy environment and
Run calcPiSerial.py on a compute node.**

**For our example program the more steps it takes the
more accurate it is, but the longer it takes.**

```
[vanilla@hopper intro_workshop]$ cat slurm/calc_pi_serial.sh
#!/bin/bash
#SBATCH --partition debug
#SBATCH --ntasks 1
#SBATCH --time 00:05:00
#SBATCH --job-name calc_pi_serial
#SBATCH --mail-user your_username@unm.edu
#SBATCH --mail-type ALL

module load miniconda3
source activate numpy

cd $SLURM_SUBMIT_DIR
python code/calcPiSerial.py 1000000000
[vanilla@hopper intro_workshop]$
```

```
[vanilla@hopper intro_workshop]$ sbatch slurm/calc_pi_serial.sh
sbatch: Using account 2016199 from ~/.default_slurm_account
Submitted batch job 5263
vanilla@hopper:~/workshops/intro_workshop$ squeue -me
JOBID PARTITION      NAME      USER ST          TIME   NODES NODELIST(REASON)
5263    debug calc_pi_  vanilla  R          0:44       1 hopper011
```

Edit `slurm/calc_pi_serial.sh`.

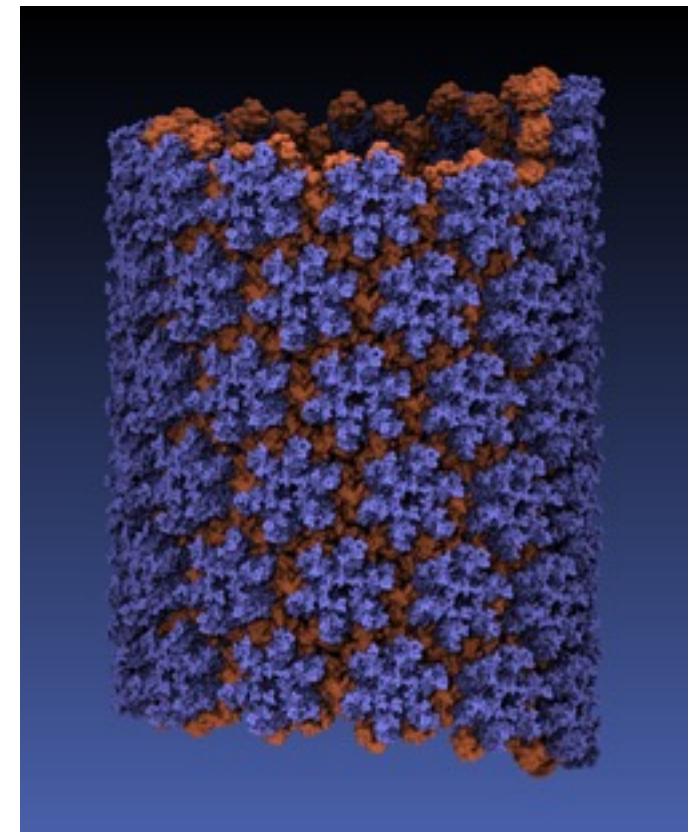
Change the email address to your address and submit the script.

Then enter `squeue --me` to see the job status.

Take a look at the job output.



NAMD and VMD are part of the team winning the **2020 ACM Gordon Bell Special Prize** for high performance computing-based **COVID-19 research**, for the paper **AI-Driven Multiscale Simulations Illuminate Mechanisms of SARS-CoV-2 Spike Dynamics**, presented at Supercomputing 2020, Nov 18, 2020.



```
vanilla@hopper:~ $ cd workshops/namd  
mfricke@hopper:~/workshops/namd $ ls  
namd.slurm ubq_ws_eq.conf ubq_ws.pdb ubq_ws.psf par_all27_prot_lipid.inp
```

namd.slurm - SLURM Batch Script

ubq_ws_eq.conf - NAMD Configuration
Simulation info in TCL

ubq_ws.pdb - Protein Databank File
Metadata and atom locations

ubq_ws.psf - Protein Structure File
Molecular force field Info

par_all27_prot_lipid.inp - CHARMM Parameter File
Protein/Lipid force fields

```
mfricke@hopper:~/workshop/namd $ cat namd.slurm
```

```
#!/bin/bash
#SBATCH --partition debug
#SBATCH --job-name namd-test
#SBATCH --time 00:05:00
#SBATCH --nodes 1
#SBATCH --ntasks-per-node 4
#SBATCH --mail-user your@email.address
#SBATCH --mail-type ALL

module load namd/2.14/verbs

# NAMD Binaries are not SLURM aware - so we use charmrun directly
charmrun $(which namd2) ++auto-provision ubq_ws_eq.conf
```

SLURM Batch Script

```
vanilla@hopper:~/workshops/namd $ squeue --me
```

```
vanilla@hopper:~/workshops/namd $ squeue --me
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
1814848	debug	namd-tes	vanilla	R	0:04	1	hopper011

```
vanilla@hopper:~/workshops/namd $ ssh hopper011
```

```
Last login: Fri Feb 17 10:04:34 2023 from 172.17.0.2
```

```
vanilla@hopper011:~ $ htop
```

Check the status of your job...

Note the hostname of the compute node you were allocated. In the example it is **hopper011**: SSH to that hostname and run the **htop** command.

```
matthew — mfricke@hopper011:~ — ssh hopper — 113x27

1[|||||||||||||||||100.0%] 17[ 0.0%]
2[ 0.0%] 18[ 0.0%]
3[|||||||||||||100.0%] 19[ 0.0%]
4[ 0.0%] 20[ 0.0%]
5[ 0.0%] 21[ 0.0%]
6[ 0.0%] 22[ 0.0%]
7[|||||||||||||100.0%] 23[ 0.0%]
8[ 0.0%] 24[ 0.0%]
9[|||||||||||||99.3%] 25[ 0.0%]
10[ 0.0%] 26[ 0.0%]
11[ 0.0%] 27[ 0.0%]
12[ 0.0%] 28[ 0.0%]
13[ 0.0%] 29[ 0.0%]
14[ 0.0%] 30[ 0.0%]
15[ 0.0%] 31[ 0.0%]
16[ 0.0%] 32[ 0.0%]
Mem[|||||||13.3G/92.9G] Tasks: 44, 88 thr, 428 kthr; 5 running
Swp[ 0K/85.2G] Load average: 3.07 1.35 0.72
Uptime: 36 days, 07:43:13

PID USER PRI NI VIRT RES SHR S CPU%△MEM% TIME+ Command
 1 root 20  0 232M 11524 8592 S 0.0 0.0 0:18.18 /sbin/init
 587 root 20  0 696M 19304 4740 S 0.0 0.0 4:23.93 /warewulf/bin/wwclient
 590 root 20  0 696M 19304 4740 S 0.0 0.0 0:03.35 /warewulf/bin/wwclient
 591 root 20  0 696M 19304 4740 S 0.0 0.0 0:08.69 /warewulf/bin/wwclient
F1Help F2Setup F3Search F4Filter F5Tree F6SortBy F7Nice -F8Nice +F9Kill F10Quit
```

```
vanilla@hopper:~/workshops/namd $ ls  
namd.slurm  
par_all27_prot_lipid.inp  
slurm-1814849.out  
ubq_ws_eq.conf  
ubq_ws_eq.coor  
ubq_ws_eq.coor.BAK  
ubq_ws_eq.eq.dcd  
ubq_ws_eq.eq.dcd.BAK  
ubq_ws_eq.restart.coor  
ubq_ws_eq.restart.vel  
ubq_ws_eq.restart.xsc  
ubq_ws_eq.vel  
ubq_ws_eq.vel.BAK  
ubq_ws_eq.xsc  
ubq_ws_eq.xsc.BAK  
ubq_ws.pdb  
ubq_ws.psf  
vanilla@hopper:~/workshops/namd $
```

Check your output after the job completes...

ubq_ws_eq.restart.vel - atomic velocities

ubq_ws_eq.restart.coor - atomic coordinates

.restart. - checkpoint files

```
vanilla@hopper:~/workshops/namd $ ssh xena
Last login: Fri Feb 17 10:08:12 2023 from hopper.alliance.unm.edu
<snip>
```

Xena is our GPU and large memory cluster.
NAMD can use GPU acceleration.

```
vanilla@xena:~/workshops/namd $ module spider namd
```

```
-----  
namd/2.14:  
-----
```

```
Versions:
```

```
namd/2.14/cuda-multicore  
namd/2.14/cuda-verbs
```

Find the NAMD module names on Xena with the “spider” command

```
vanilla@hopper:~ $ qgrok
queues      free    busy    offline   jobs   nodes   CPUs   GPUs   CPUs/node   GPUs/node   Memory/node   time_limit   CPU
_limit
----  -----  -----  -----  -----  -----  -----  -----  -----  -----  -----  -----
-- 
singleGPU  0       23     0       13     23     368    23     16     1        61G      2-00:00:00  192
dualGPU    1       2      1       1      4      64     8      16     2        61G      2-00:00:00  192
bigmem-1TB 1       1      0       1      2      128    0      64     0        1006G    2-00:00:00  192
bigmem-3TB 2       0      0       0      2      128    0      64     0        3.0T     2-00:00:00  192
debug      2       0      0       0      2      32     2      16     1        61G      4:00:00    8
systems    1       0      0       0      1      16     1      16     1        61G      2-00:00:00  192
totals:    6       26     1       15     33     720    33
```

Find the available partitions with “qgrok”

```
mfricke@hopper:~/workshops/namd$ cat namd_xena.slurm
#!/bin/bash
#SBATCH --partition debug
#SBATCH --gpus 1 # Request 1 GPU
#SBATCH --job-name namd-test
#SBATCH --time 00:05:00
#SBATCH --nodes 1
#SBATCH --ntasks-per-node 4
#SBATCH --mail-user your@email.address
#SBATCH --mail-type ALL

module load namd/2.14/cuda-verbs

# NAMD Binaries are not SLURM aware - so we use charmrun directly
charmrun $(which namd2) ++auto-provision ubq_ws_eq.conf
```

Now we can write a new SLURM script for Xena and just modify the module name and request a GPU

```
vanilla@xena:~/workshops/namd [master ?]$ sbatch namd_xena.slurm
sbatch: Using account 2016199 from ~/.default_slurm_account
Submitted batch job 333532
```

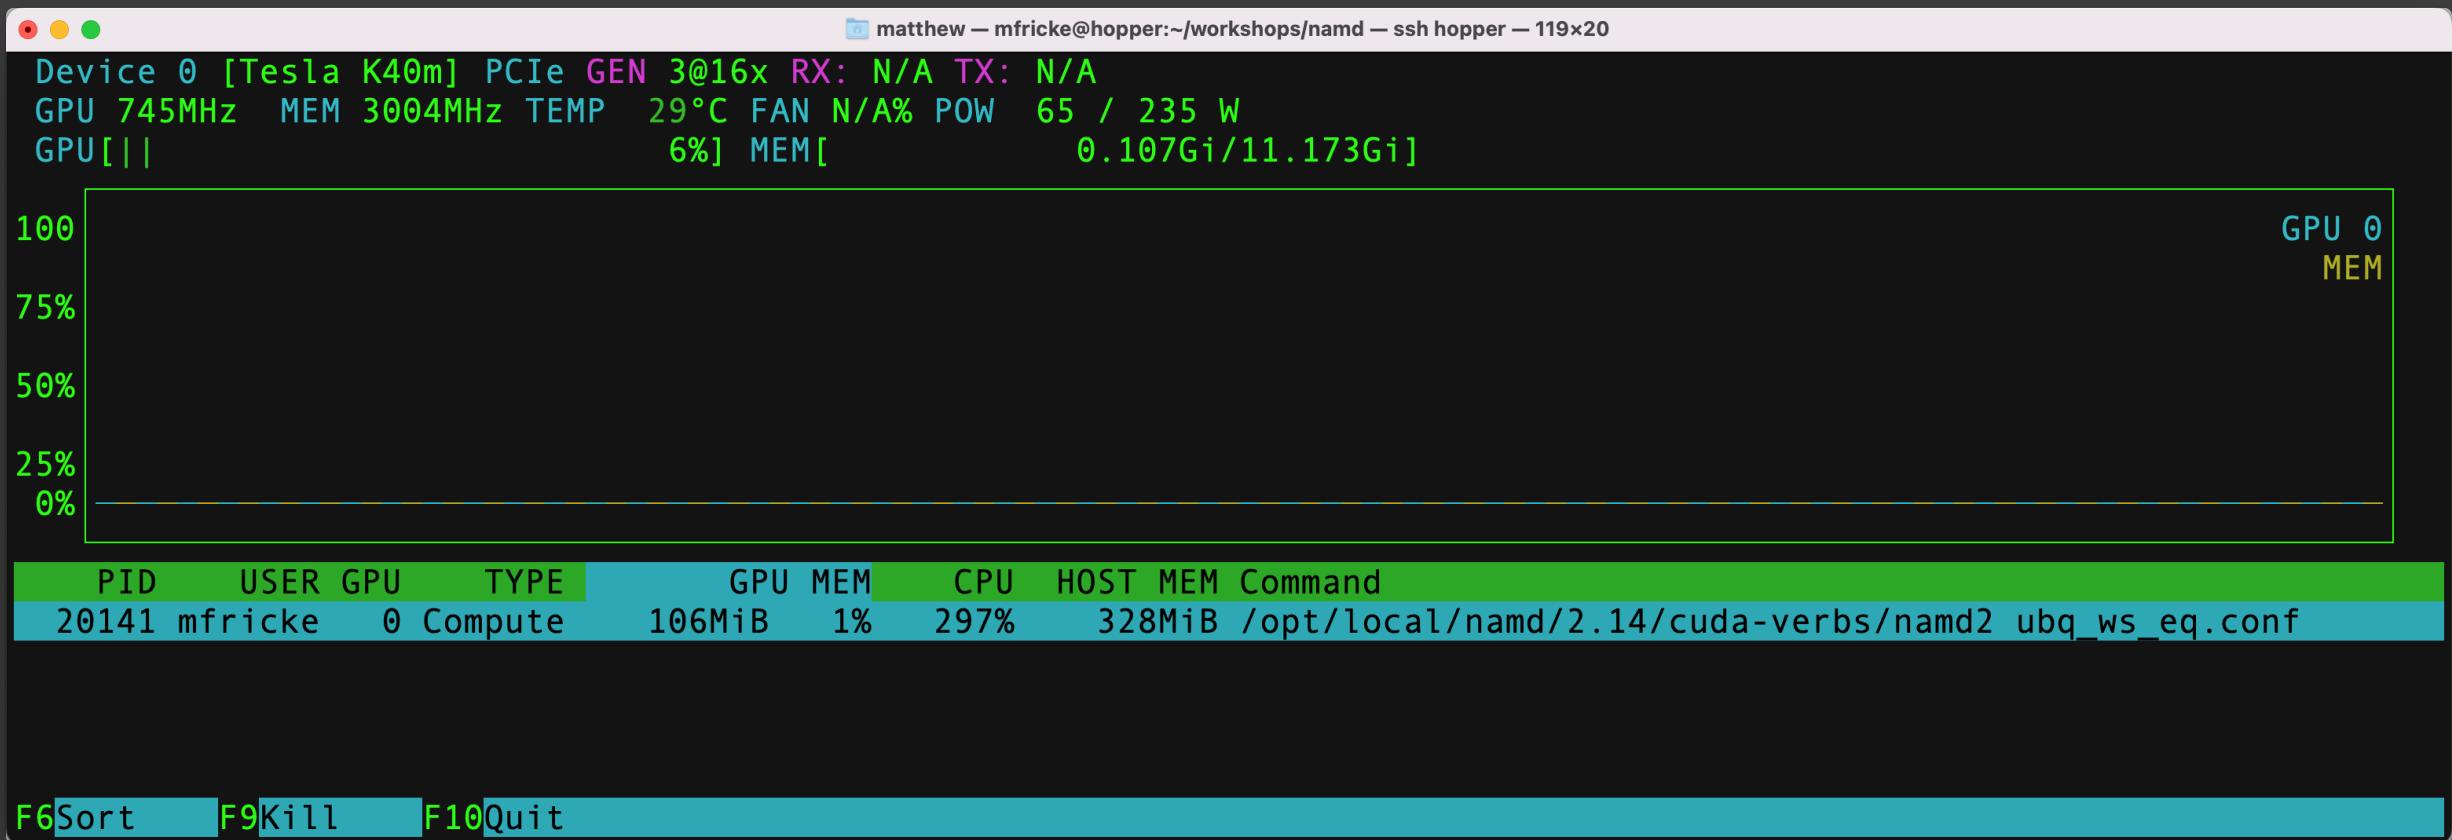
```
vanilla@xena:~/workshops/namd [master ?]$ squeue --me
JOBID PARTITION NAME      USER      ST      TIME   NODES NODELIST(REASON)
333532 debug    namd-test vanilla    R      0:04      1 xena-test
```

```
vanilla@xena:~/workshops/namd [master ?]$ ssh xena-test
Last login: Fri Feb 17 10:19:37 2023 from xena
vanilla@xena-test:~ $ module load nvtop
vanilla@xena-test:~ $ nvtop
```

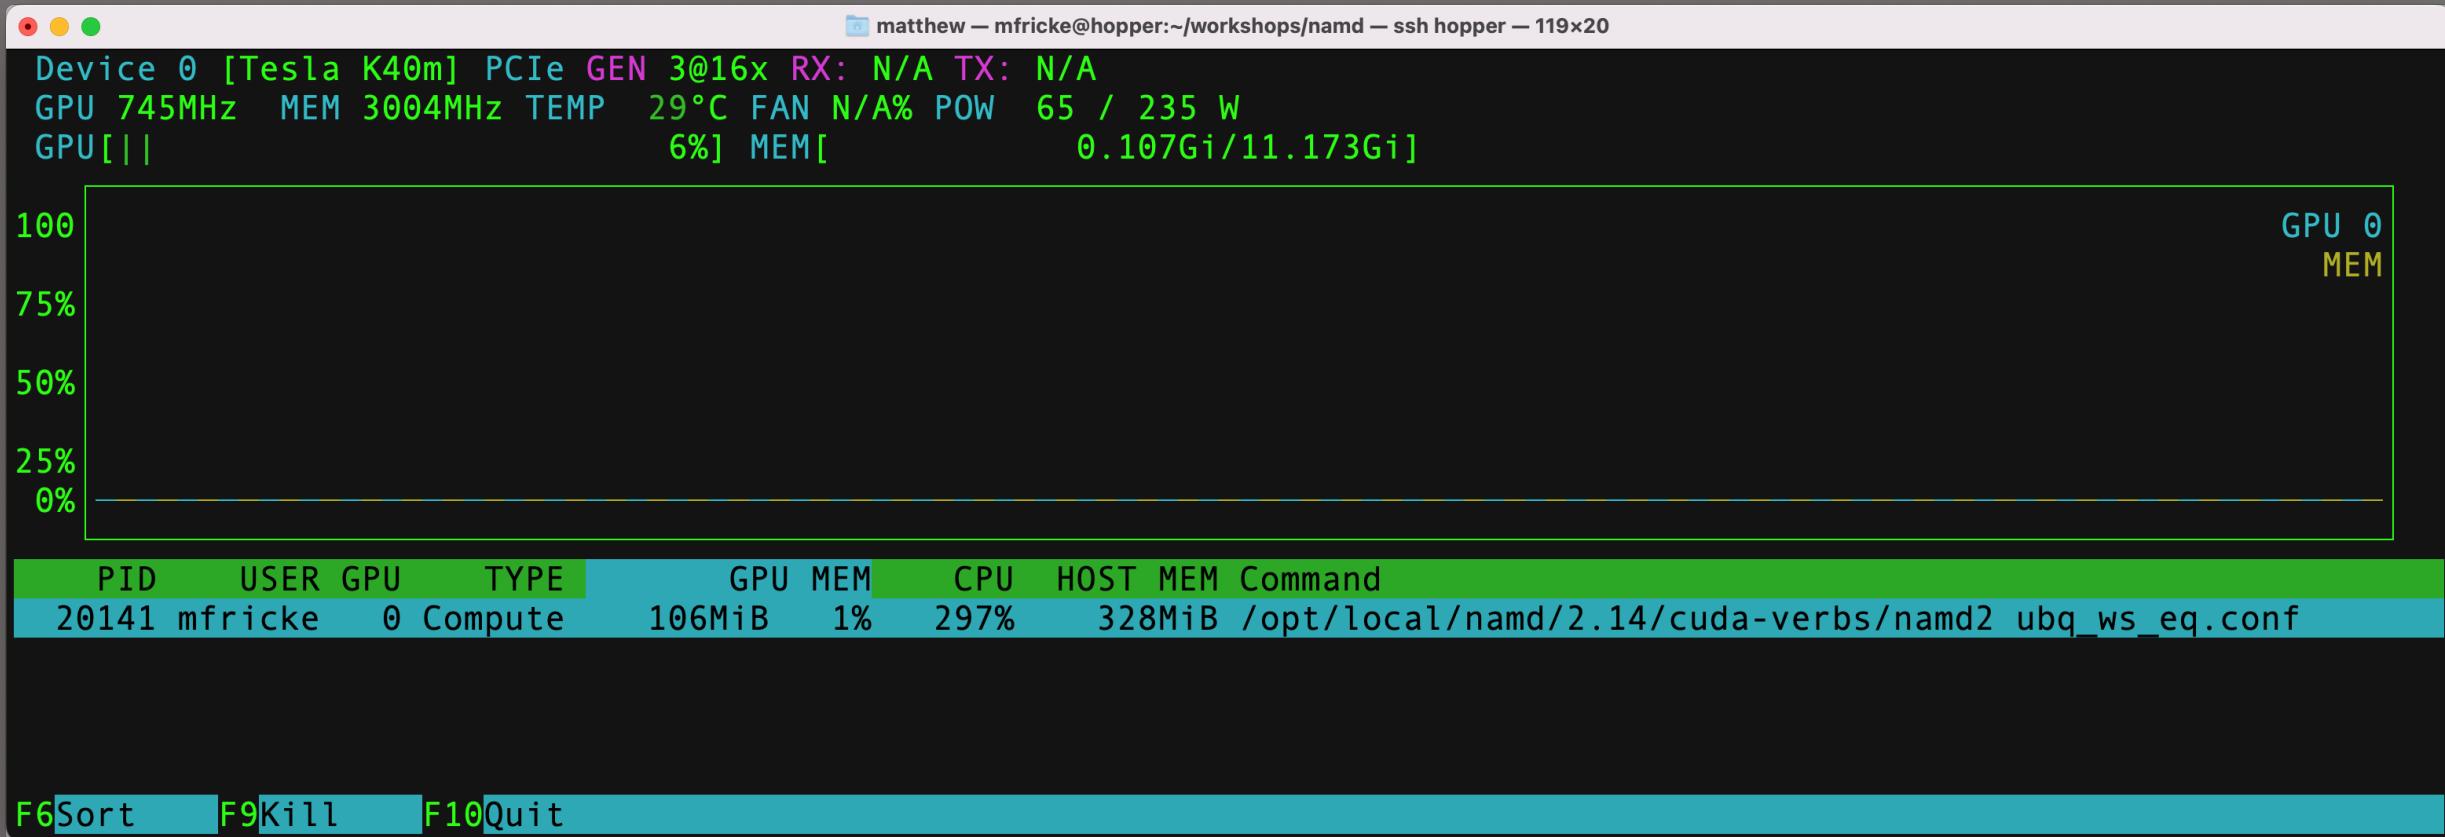
Check the status of your job...

Note the hostname of the compute node you were allocated. In the example it is **xena-test**:

SSH to that hostname and run the **nvtop** command.



```
Device 0 [Tesla K40m] PCIe GEN 3@16x RX: N/A TX: N/A
GPU 745MHz MEM 3004MHz TEMP 29°C FAN N/A% POW 65 / 235 W
GPU[|||] 6%] MEM[ 0.107Gi/11.173Gi]
```



You have learned

- how to run programs using the SLURM scheduler
- the difference between interactive and batch jobs
- how to check the status of your jobs
- how to select debug vs general SLURM partitions
- how to ask for the hardware resources you need
- you ran the NAMD molecular dynamics code using SLURM

